

6.874 Deep Learning in the Life Sciences

## Lecture 6

# **Generative Models: GANs, VAEs, Learning Representations**

Prof. Manolis Kellis

Slides Credit: Ben Lengerich, Relja Arandjelovic,  
Fei-Fei Li, Justin Johnson, Serena Yeung

# Interpretable Deep Learning

## 1. Intro to Interpretability

- 1a. Interpretability definition:** Convert implicit NN information to human-interpretable information
- 1b. Motivation:** Verify model works as intended; debug classifier; make discoveries; Right to explanation
- 1c. Ante-hoc** (train interpretable model) vs. **Post-hoc** (interpret complex model; degree of “locality”)

## 2. Interpreting Deep Neural Networks

- 2a. Interpreting Models** (macroscopic, understand internals) vs. **decisions** (microscopic, practical applications)
- 2b. Interpreting Models:** Weight visualization, Surrogate model, Activation maximization, Example-based
- 2c. Interpreting Decisions:**
  - Example-based
  - Attribution Methods: why are gradients noisy?
  - Gradient-based Attribution: SmoothGrad, Interior Gradient
  - Backprop-based Attribution: Deconvolution, Guided Backpropagation

## 3. Evaluating Attribution Methods

- 3a. Qualitative: Coherence:** Attributions should highlight discriminative features / objects of interest
- 3b. Qualitative: Class Sensitivity:** Attributions should be sensitive to class labels
- 3c. Quantitative: Sensitivity:** Removing feature with high attribution → large decrease in class probability
- 3d. Quantitative: ROAR & KAR.** Low class prob cuz image unseen → remove pixels, retrain, measure acc. drop

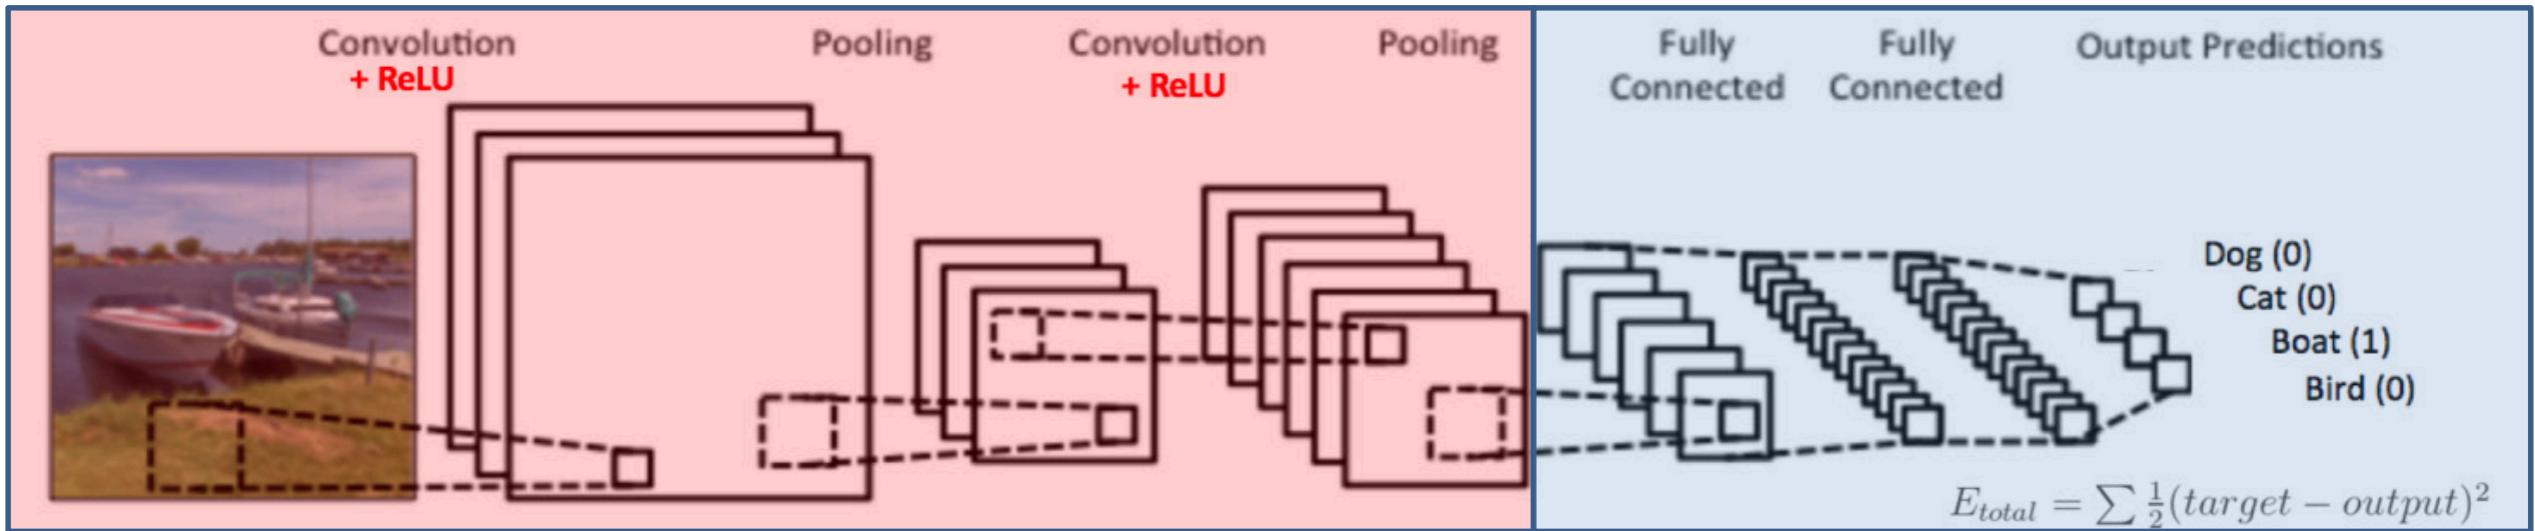
# Previously

---

- We've seen ways to model  $X \rightarrow Y$  for a fixed dataset
- But that's not what the world looks like:
  - No  $Y$
  - Limited samples
  - Many related datasets
  - Changes over time
- Can we learn the rules which govern how the real world varies?

# Learning Representations

# Key idea: Representation learning



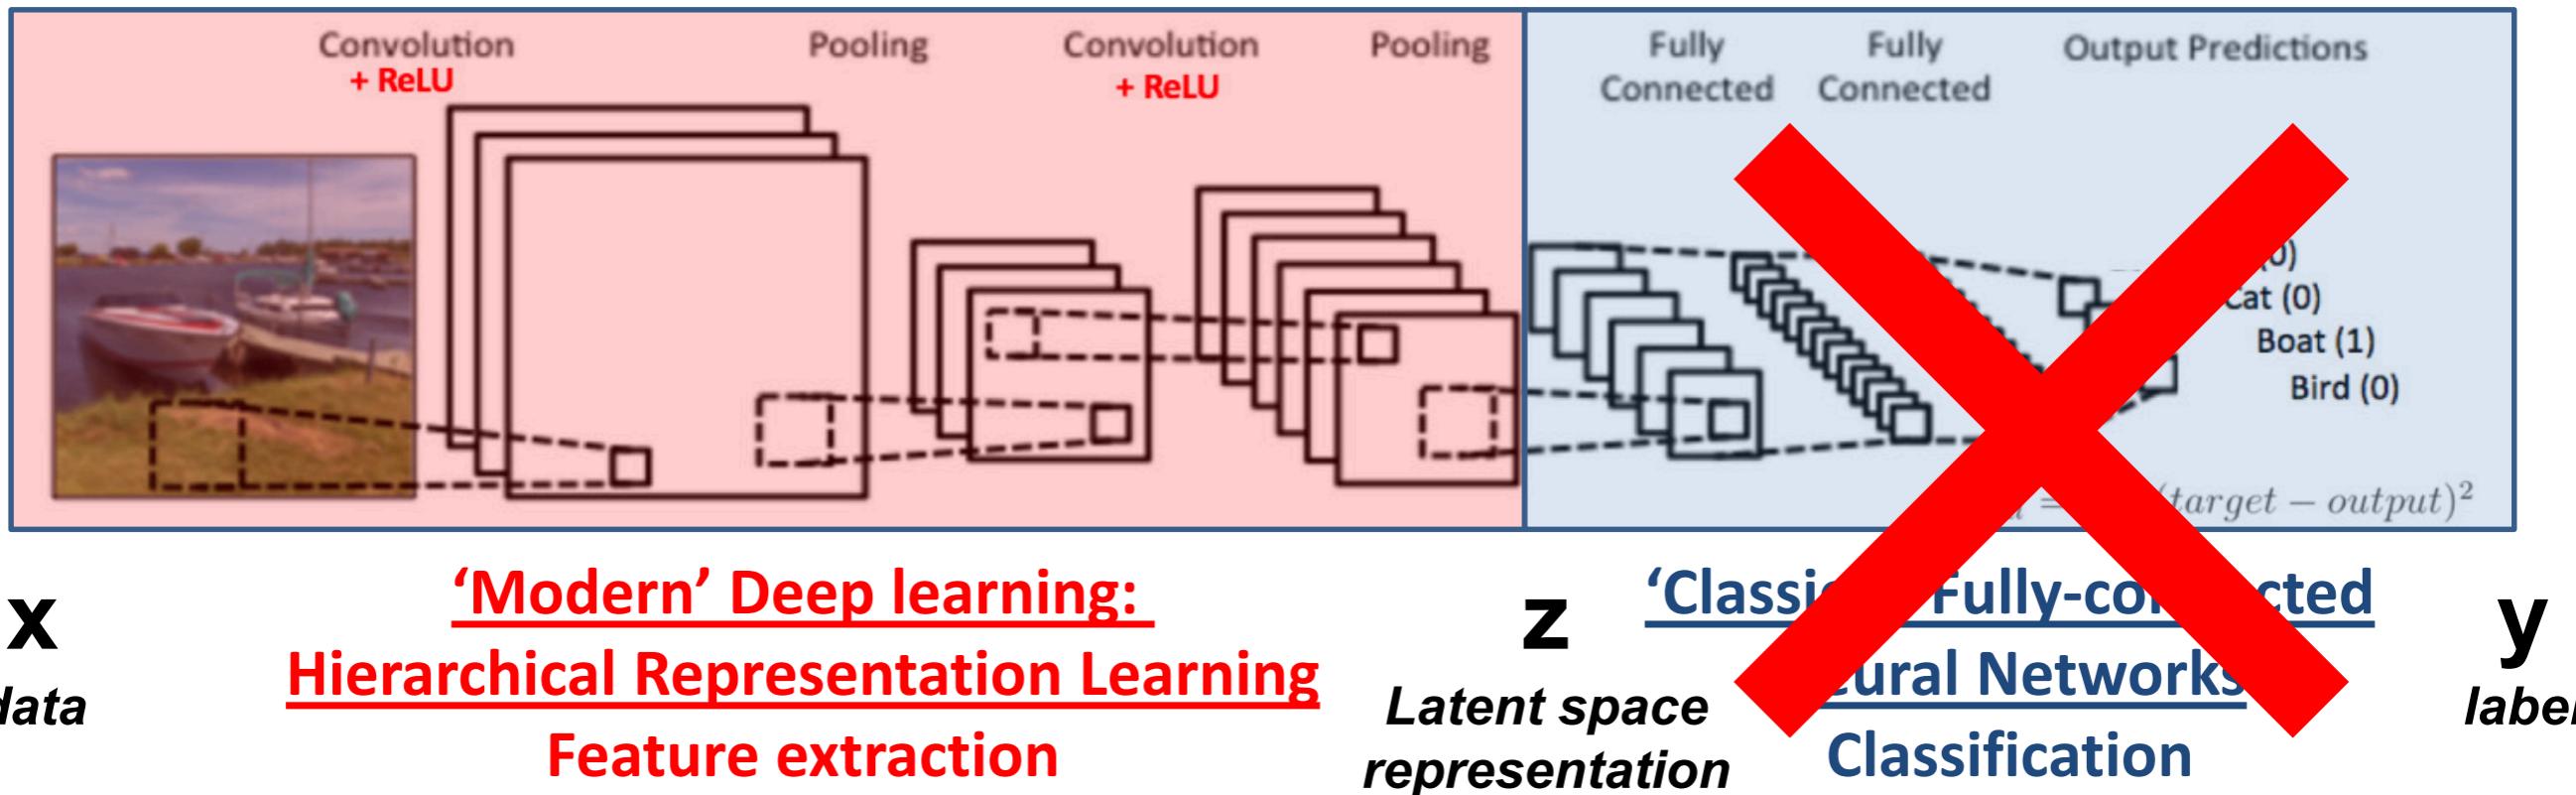
'Modern' Deep learning:  
Hierarchical Representation Learning  
Feature extraction

'Classical' Fully-connected  
Neural Networks  
Classification

In deep learning, the two tasks are coupled:

- the **classification task** “drives” the **feature extraction**
- **Extremely powerful and general paradigm**
  - **Be creative!** The field is still at its infancy!
  - New application domains (e.g. beyond images) can have **structure** that current architectures **do not capture/exploit**
  - Genomics/biology/neuroscience can help drive development of **new architectures**

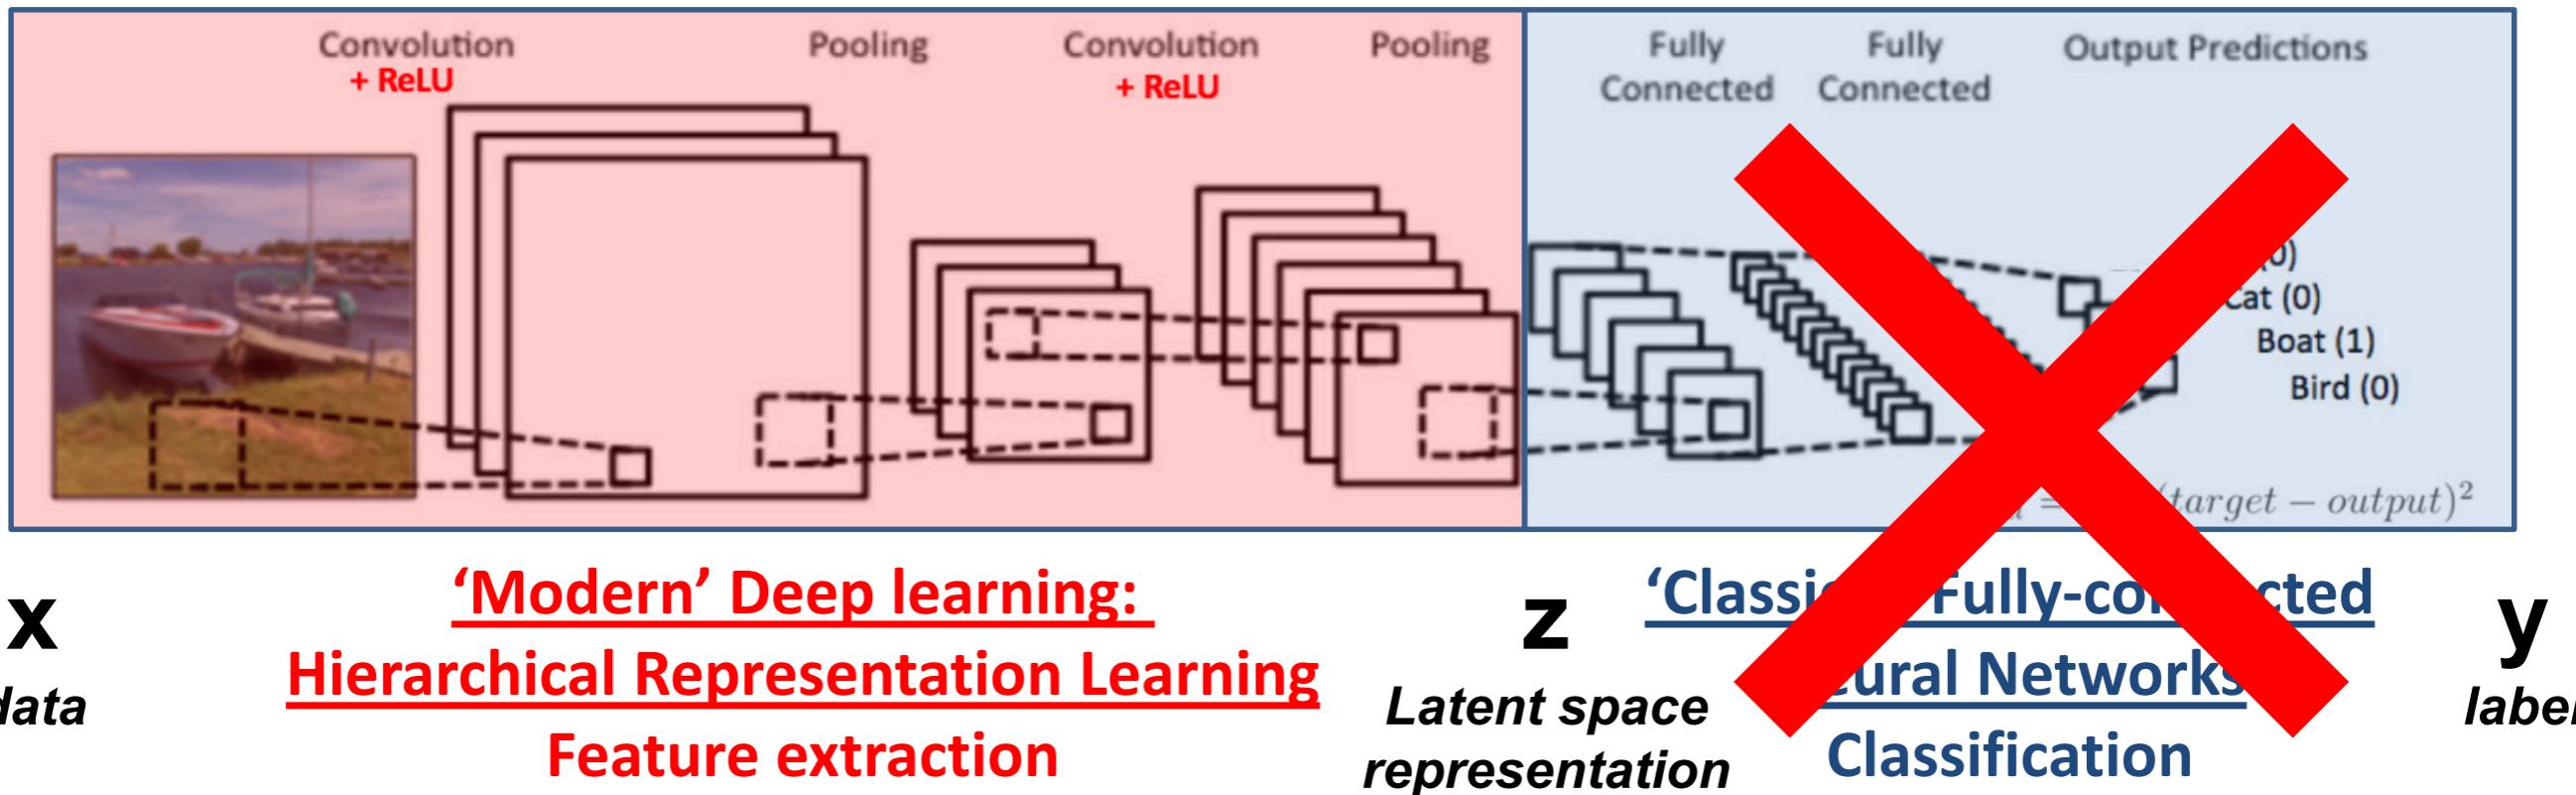
# Representation learning without annotations?



**Many ideas are possible (and yours could be even better!):**

1. Predict the future: RNNs, Video
2. Pretext tasks: predict self, before/after, missing patch, correct rotation, colorization, up-sampling, multimodal
3. Compression: Autoencoder (predict self, through clamp), representation **z**
4. Capture parameter distribution (variance): Variational Auto-Encoders
5. Make latent space parameters **z** meaningful, orthogonal, explicit, tuneable
6. Train using a second network: GANs - Improve quality of output images
7. The Sky is the Limit

# Representation learning without annotations?



**Many ideas are possible (and yours could be even better!):**

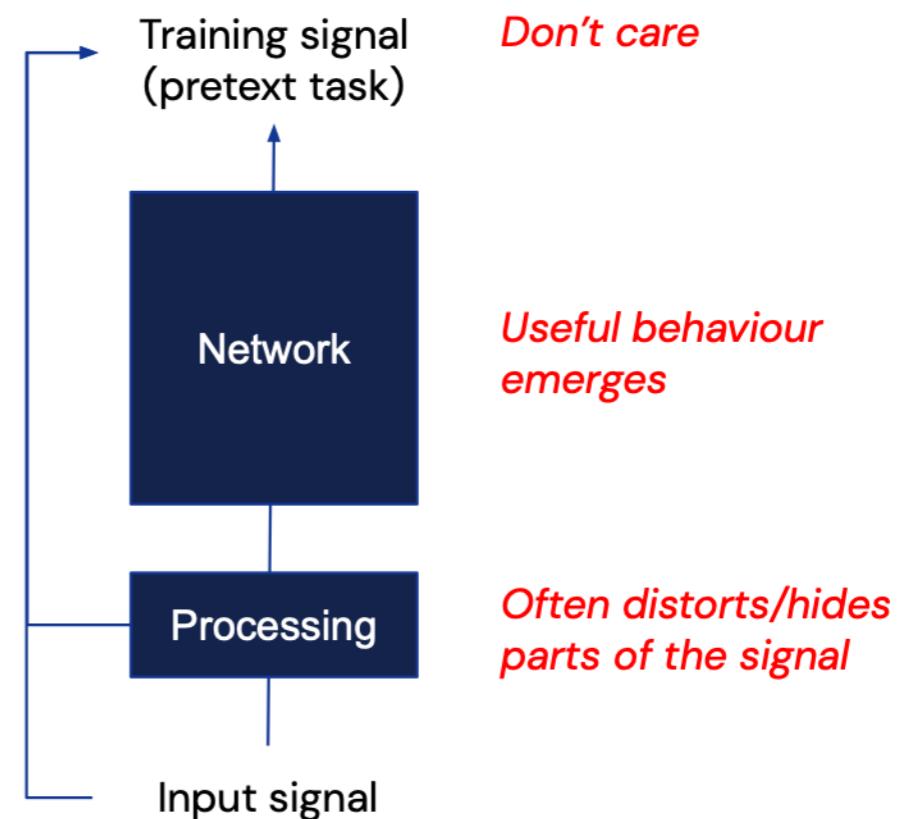
1. Predict the future: RNNs, Video
2. Compression: Autoencoder (predict self, through clamp), representation **z**
3. Pretext tasks: predict self, before/after, missing patch, correct rotation, colorization, up-sampling, multimodal
4. Capture parameter distribution (variance): Variational Auto-Encoders
5. Make latent space parameters **z** meaningful, orthogonal, explicit, tuneable
6. Train using a second network: GANs - Improve quality of output images
7. The Sky is the Limit

# Pretext Tasks

# A tour of pretext tasks

## Self-supervised learning

- **Goal:** Learn good representations
- **Means:** Construct a pretext task
  - Don't care about the pretext task itself
  - Only important it enables learning



*Don't care*

*Useful behaviour emerges*

*Often distorts/hides parts of the signal*

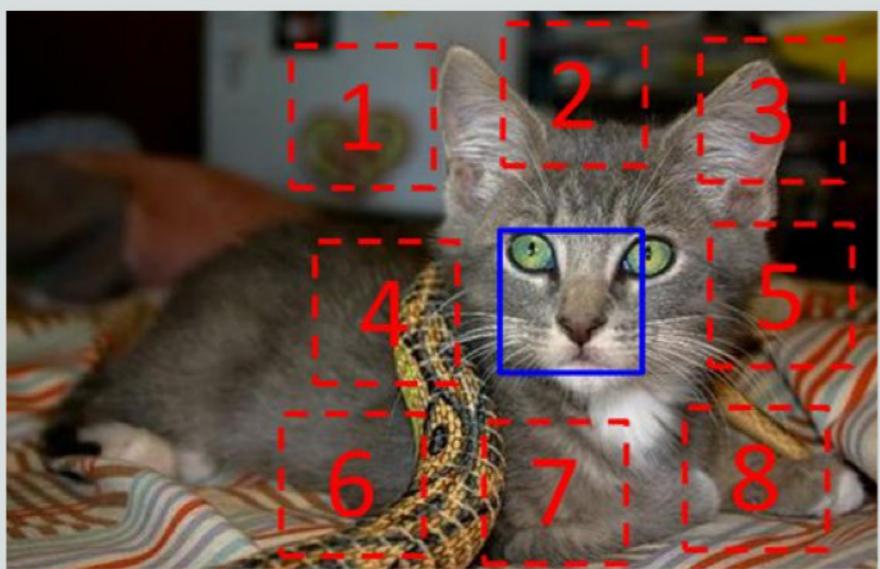
## Rough pretext task classification

- Inferring structure
- Transformation prediction
- Reconstruction
- Exploiting time
- Multimodal
- Instance classification

## Disclaimer

- Rough classification of tasks, some fit multiple categories
- Trying to cover many but inevitably missing many works
- Often have to pick one of multiple concurrent similar methods
- If A comes before B in this presentation, it doesn't mean A did it first

# Inferring structure



Slide Credit: Relja Arandjelovic

## Context prediction

Can you guess the spatial configuration for the two pairs of patches?

Question 1:



Question 2:



## Context prediction

Can you guess the spatial configuration for the two pairs of patches? **Much easier if you recognize the object!**

Question 1:



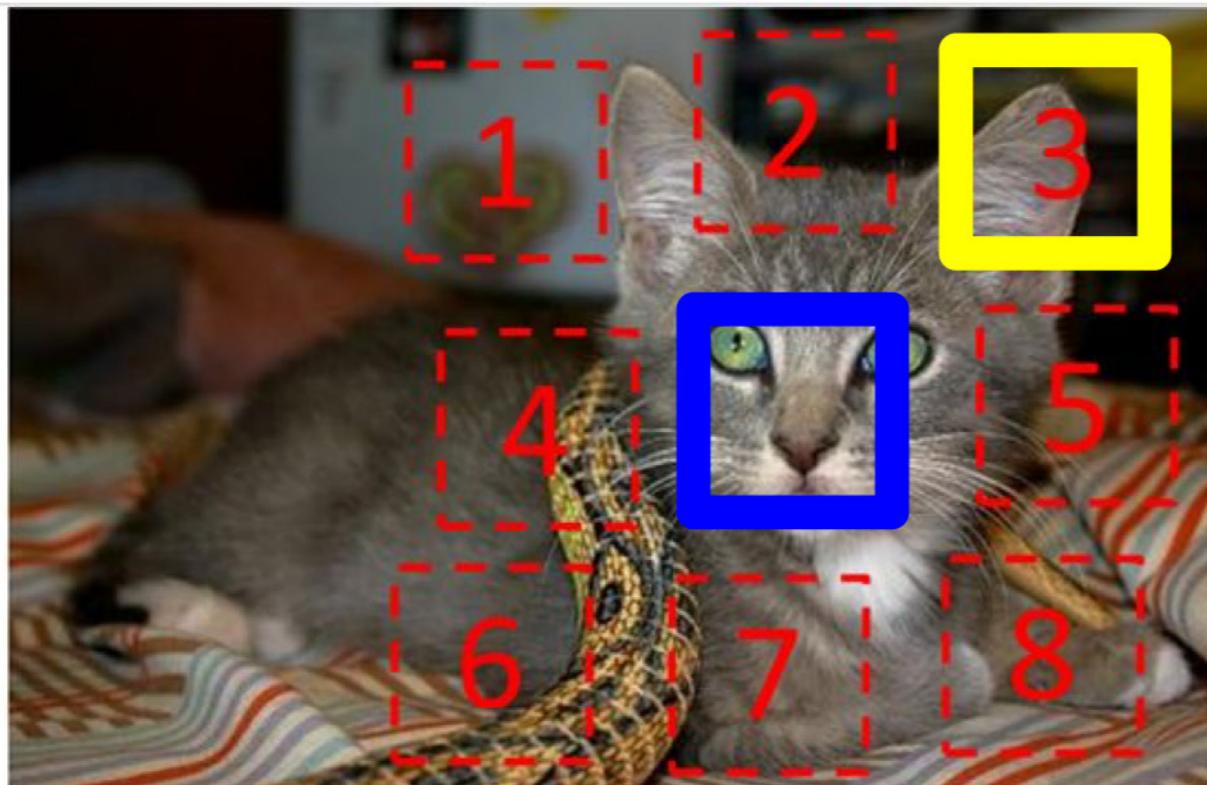
Question 2:



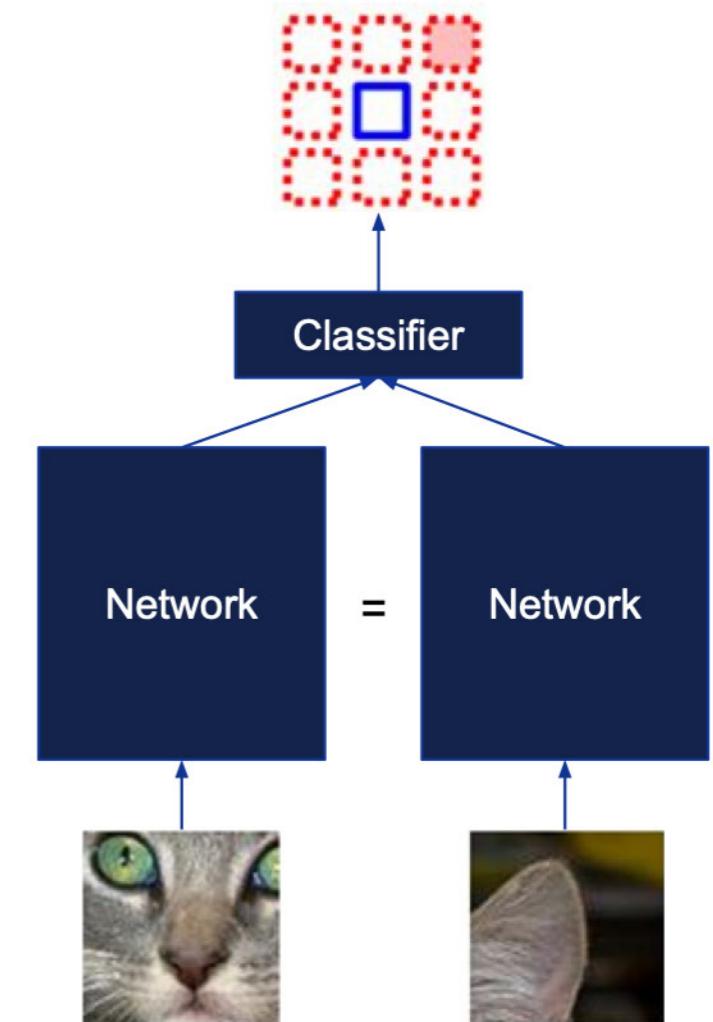
Intuition

- The network should learn to recognize object parts and their spatial relations

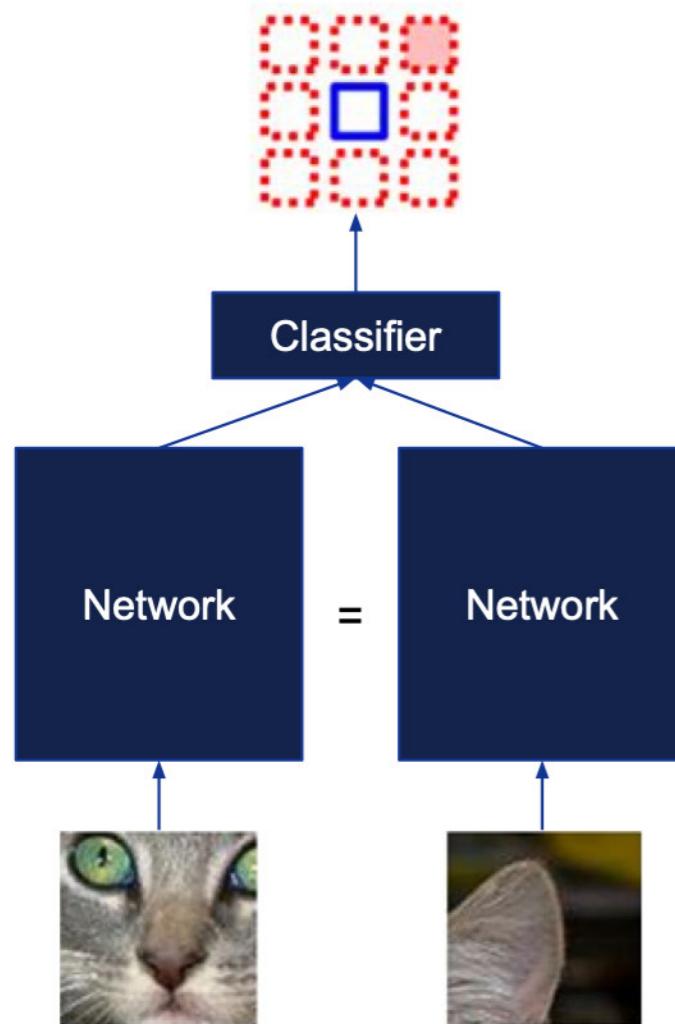
## Context prediction



$$X = (\text{[Patch from Box 3]}, \text{[Patch from Box 4]}); Y = 3$$



# Context prediction



## Pros

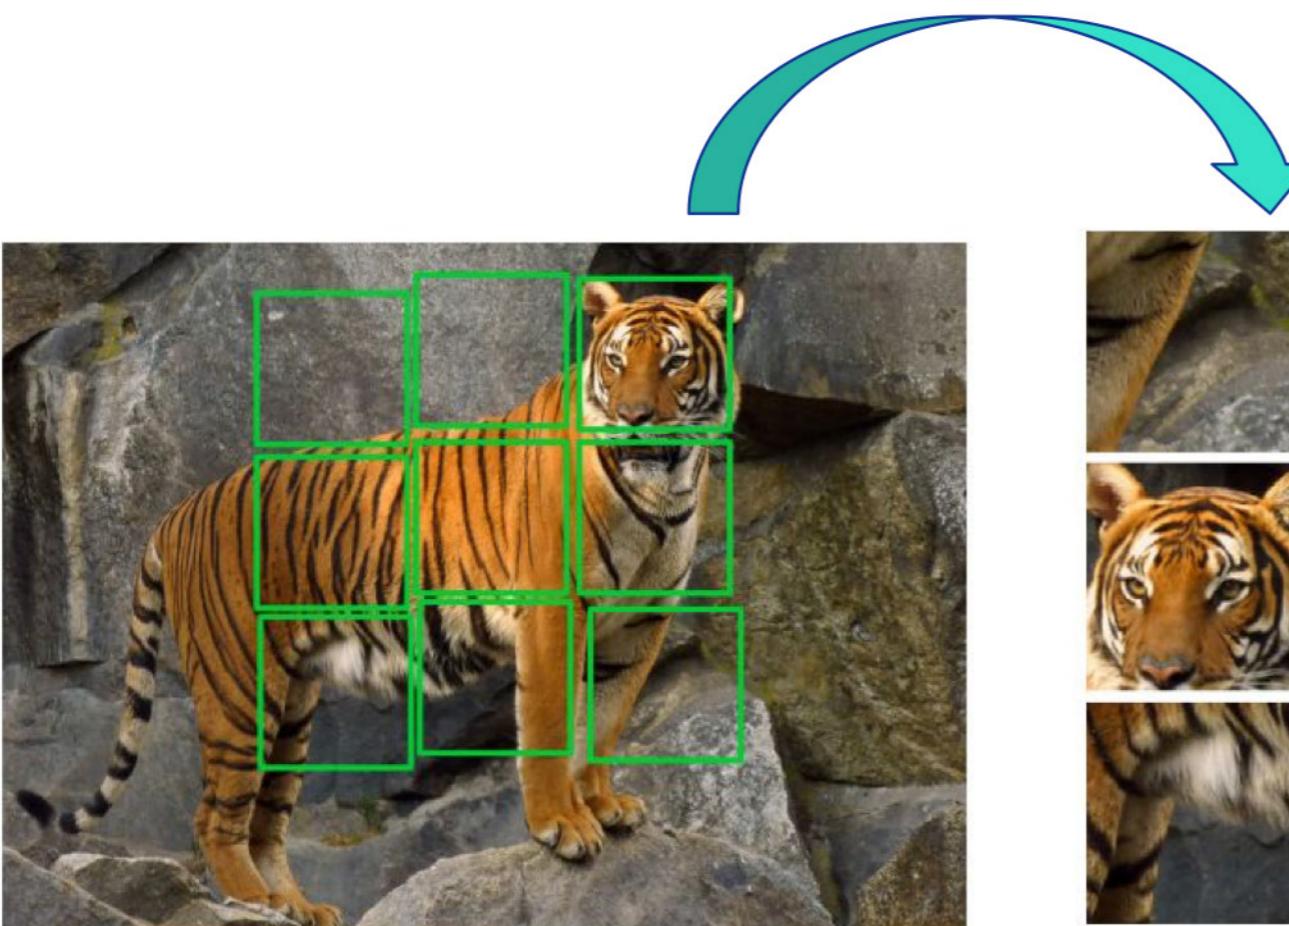
- (arguably) The first self-supervised method
- Intuitive task that should enable learning about object parts

## Cons

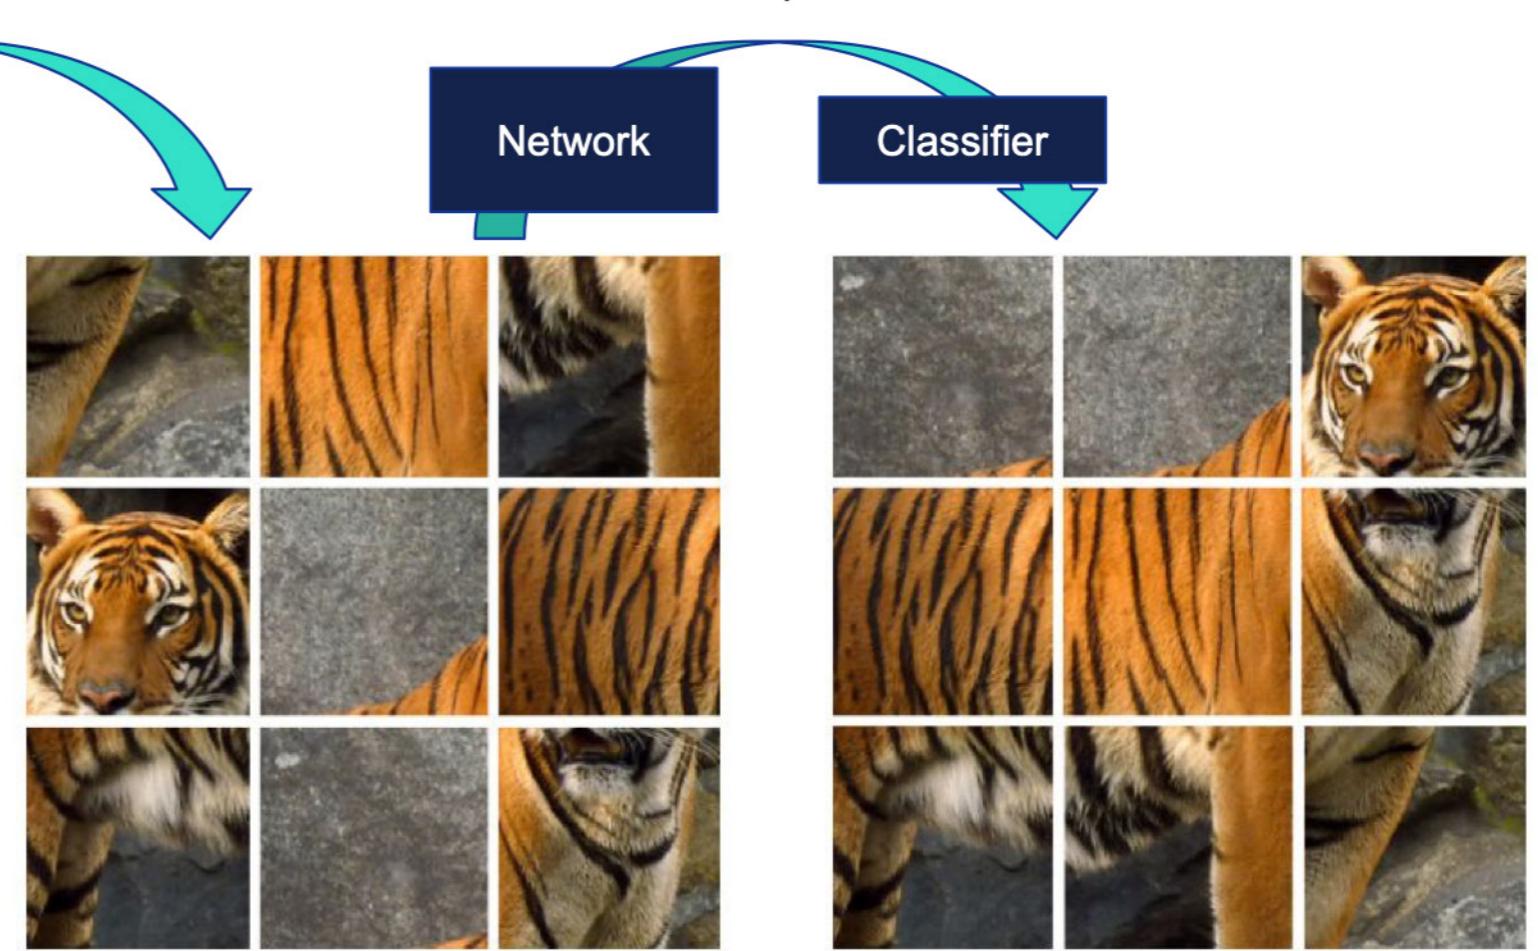
- Assumes training images are photographed with canonical orientations (and canonical orientations exist)
- Training on patches, but trying to learn image representations
- Networks can “cheat” so special care is needed [discussed later]
  - Further gap between train and eval
- Not fine-grained enough due to no negatives from other images
  - e.g. no reason to distinguish cat from dog eyes
- Small output space – 8 cases (positions) to distinguish?

## Jigsaw puzzles

Divide image into patches and permute them

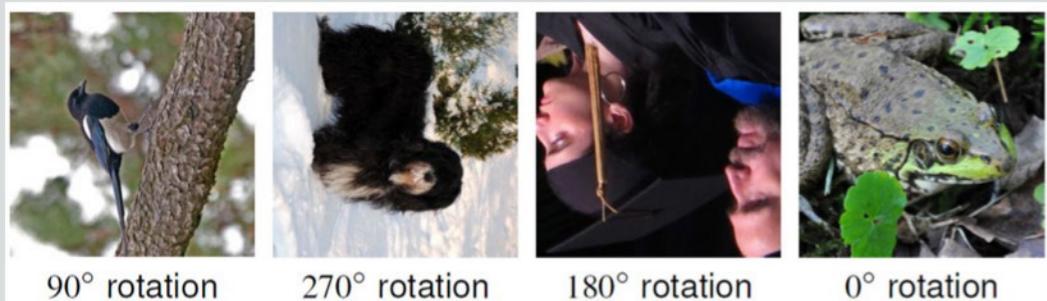


Predict the permutation



Pros & Cons: Same as for context prediction apart from being harder

# Transformation prediction

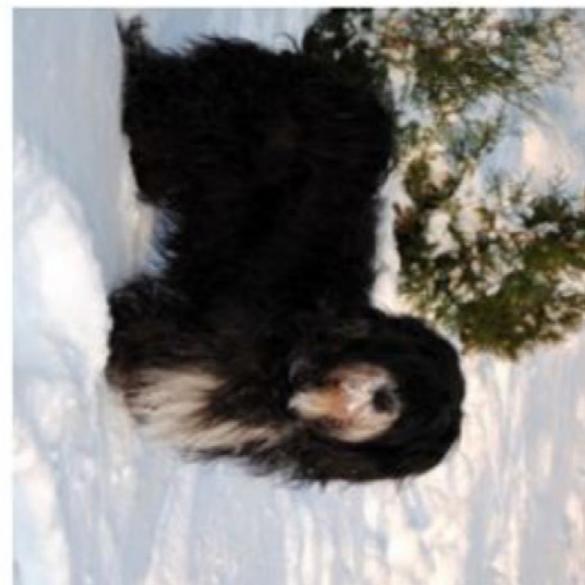
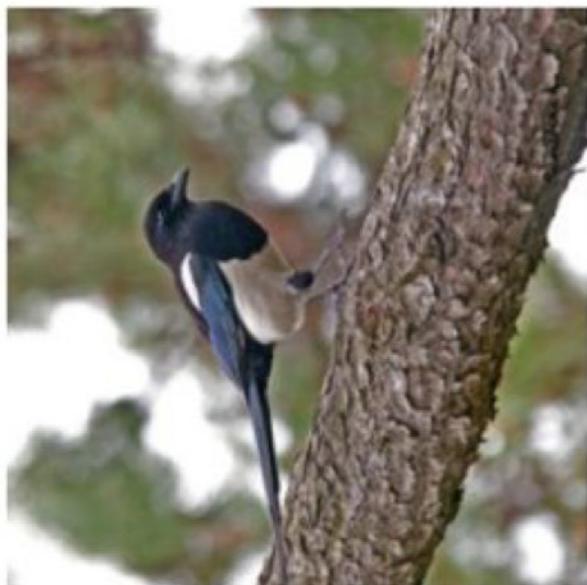


Slide Credit: Relja Arandjelovic

[“Unsupervised representation learning by predicting image rotations”, Gidaris et al. 18]

## Rotation prediction

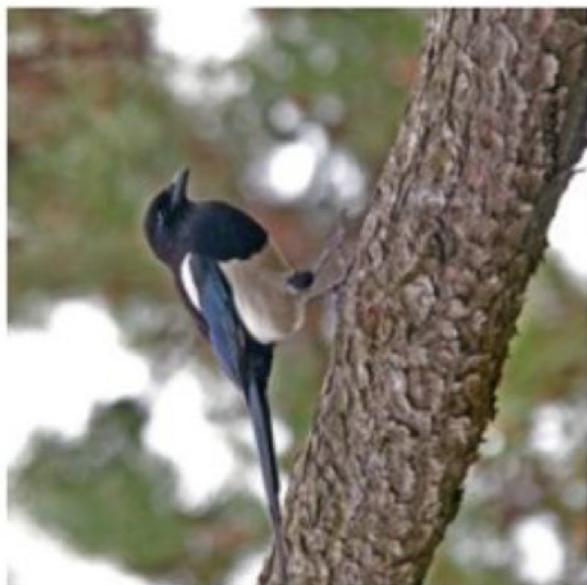
Can you guess how much rotated is applied?



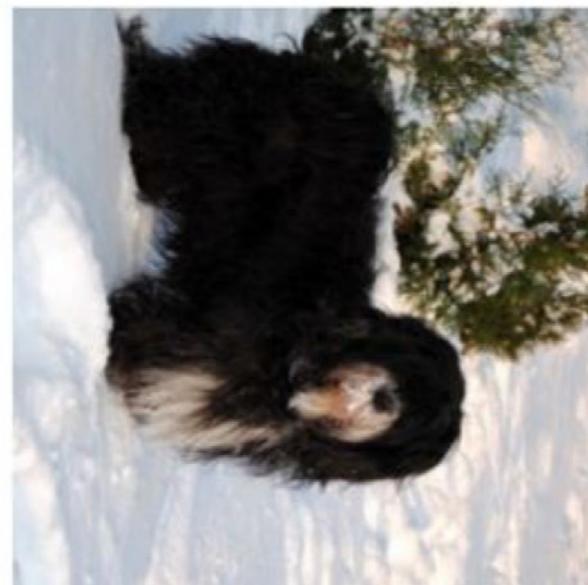
[“Unsupervised representation learning by predicting image rotations”, Gidaris et al. 18]

## Rotation prediction

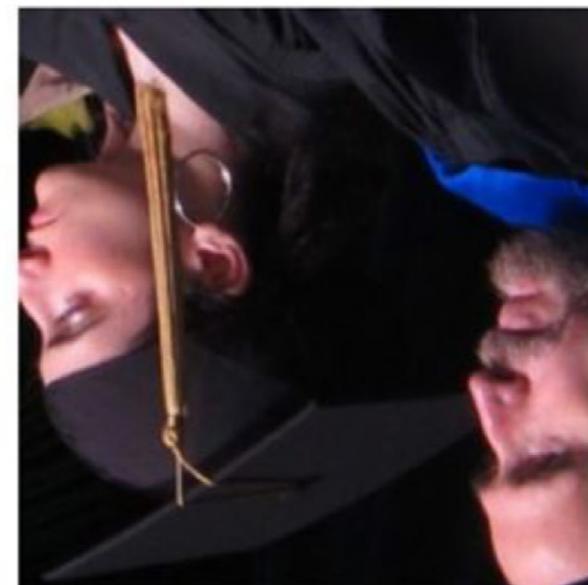
Can you guess how much rotated is applied? **Much easier if you recognize the content!**



90° rotation



270° rotation

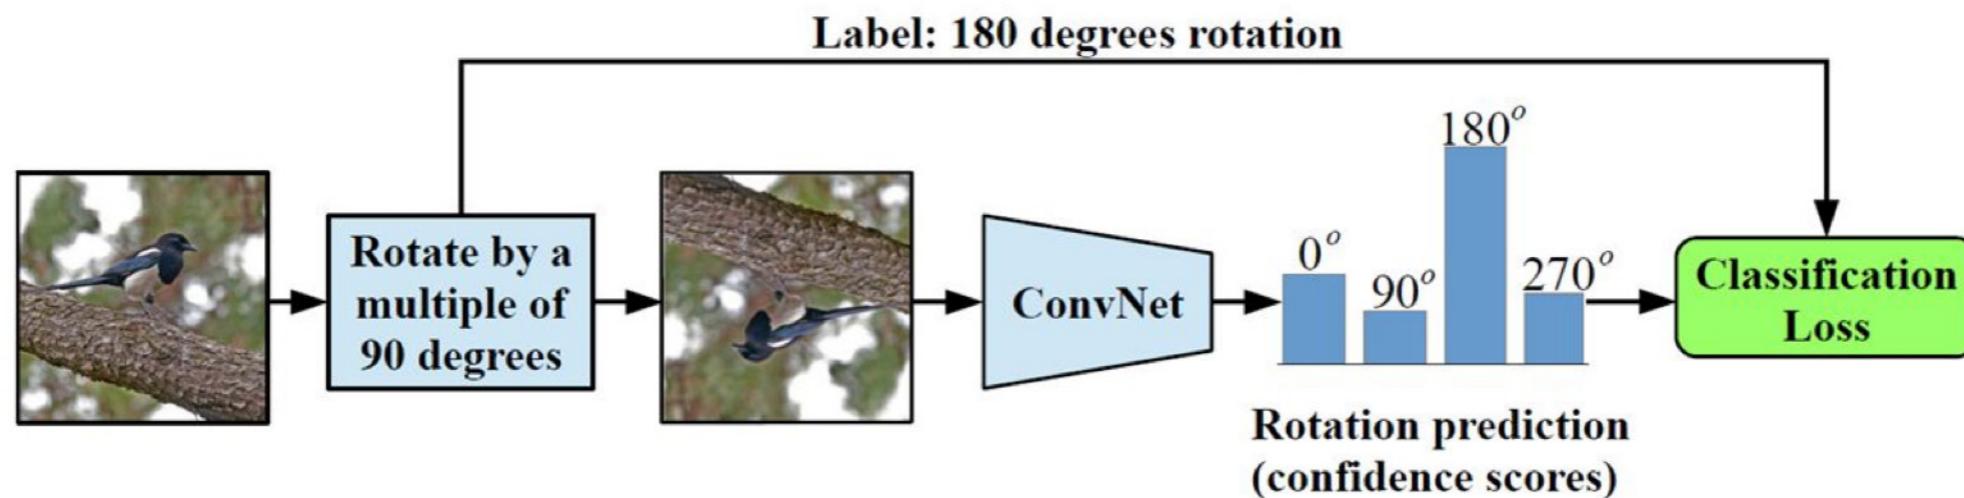


180° rotation



0° rotation

## Rotation prediction



### Pros

- Very simple to implement and use, while being quite effective

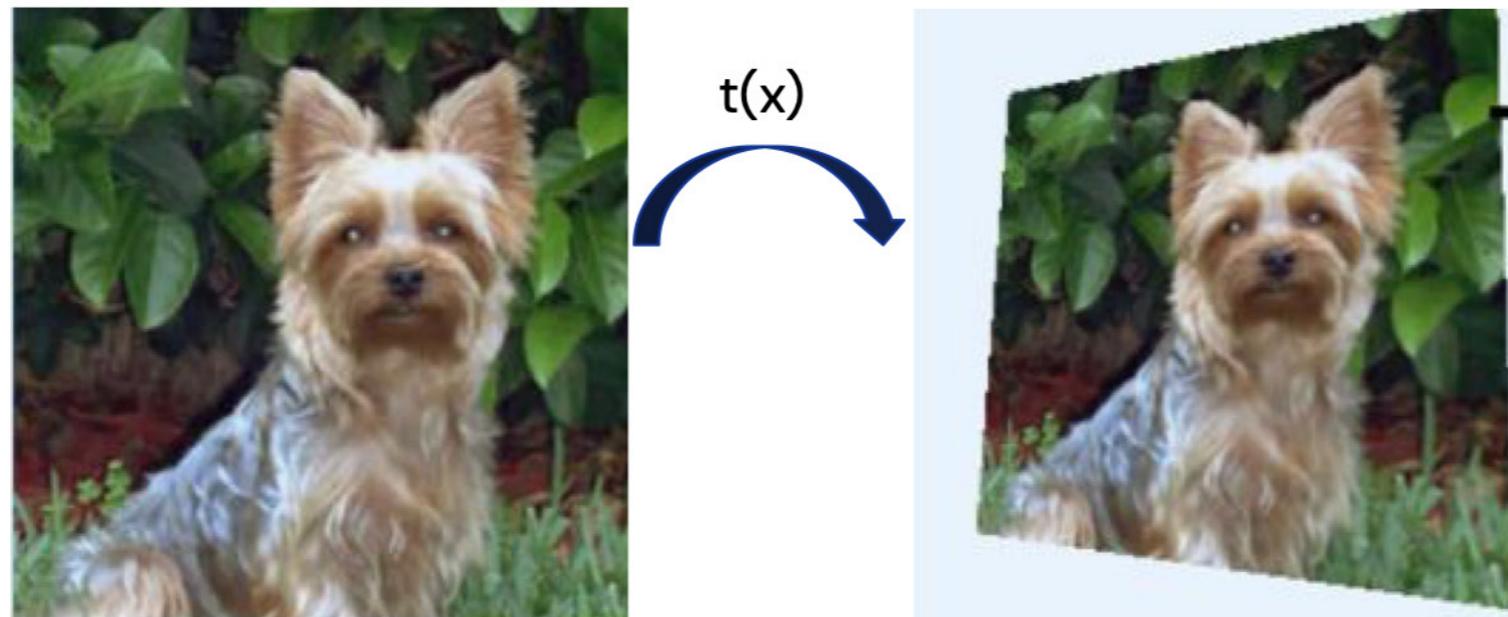
### Cons

- Assumes training images are photographed with canonical orientations (and canonical orientations exist)
- Train–eval gap: no rotated images at eval
- Not fine-grained enough due to no negatives from other images
  - e.g. no reason to distinguish cat from dog
- Small output space – 4 cases (rotations) to distinguish [not trivial to increase; see later]
- Some domains are trivial e.g. StreetView ⇒ just recognize sky

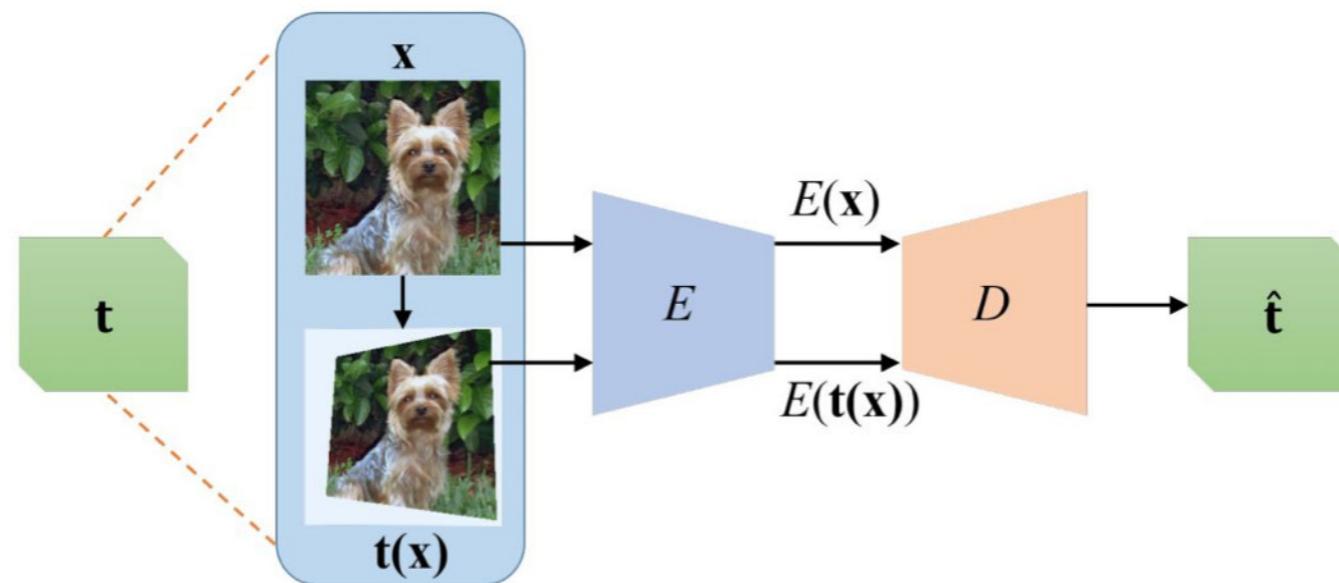
[“AET vs. AED: Unsupervised representation learning by auto-encoding transformations rather than data”, Zhang et al. 19]

## Relative transformation prediction

Estimate the transformation between two images. **Requires good features**



## Relative transformation prediction



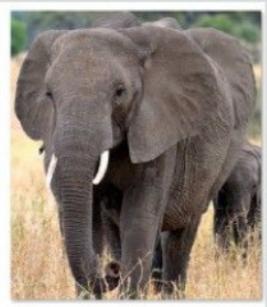
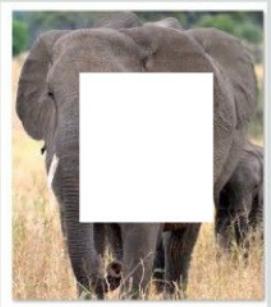
### Pros

- In line with classical computer vision, e.g. SIFT was developed for matching

### Cons

- Train–eval gap: no transformed images at eval
- Not fine-grained enough due to no negatives from other images
  - e.g. no reason to distinguish cat from dog
- Questionable importance of semantics vs low-level features (assuming we want semantics)
  - Features are potentially not invariant to transformations

# Reconstruction



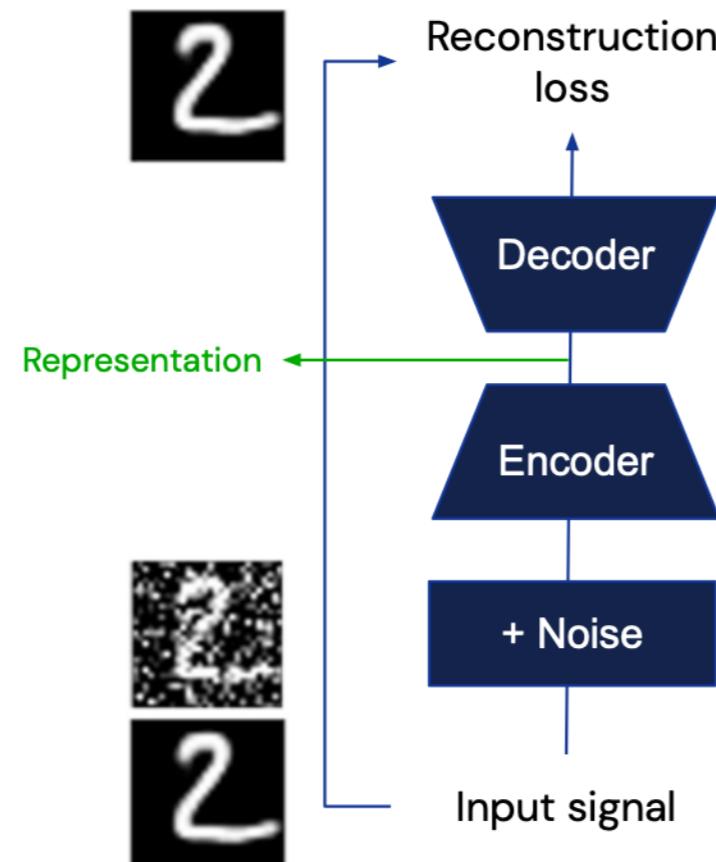
Slide Credit: Relja Arandjelovic

[“Extracting and composing robust features with denoising autoencoders”, Vincent et al. 08]

# Denoising autoencoders

What is the noise and what is the signal?

Recognizing the digit helps!



## Pros

- Simple classical method
- Apart from representations, we get a denoiser for free

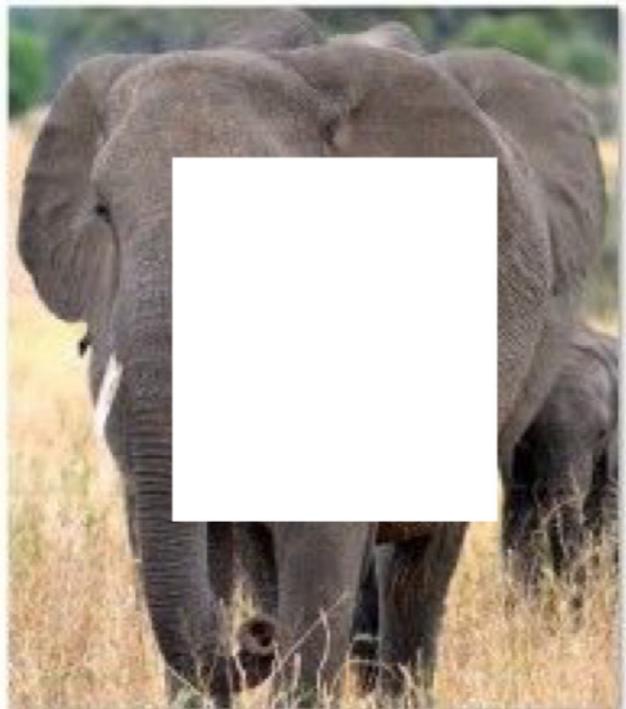
## Cons

- Train–eval gap: training on noisy data
- Too easy, no need for semantics – low level cues are sufficient

[“Context encoders: Feature learning by inpainting”, Pathak et al. 16]

## Context encoders

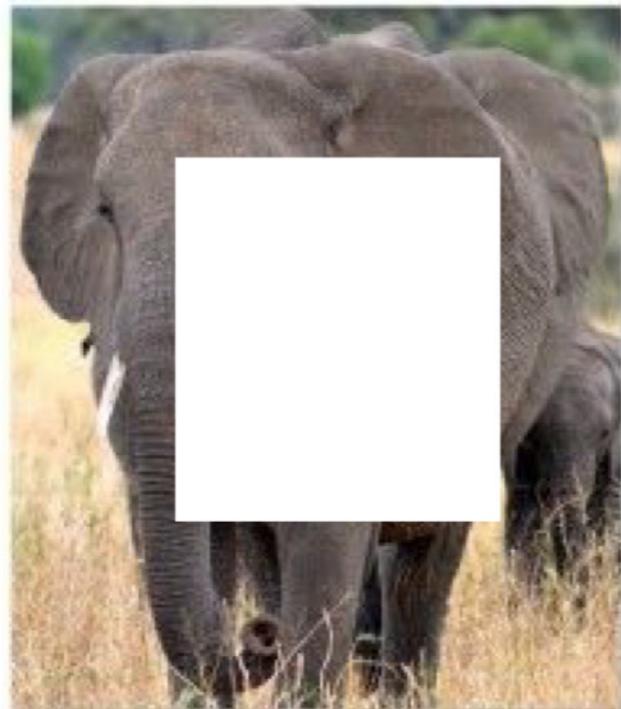
What goes in the middle?



[“Context encoders: Feature learning by inpainting”, Pathak et al. 16]

## Context encoders

What goes in the middle? **Much easier if you recognize the objects!**



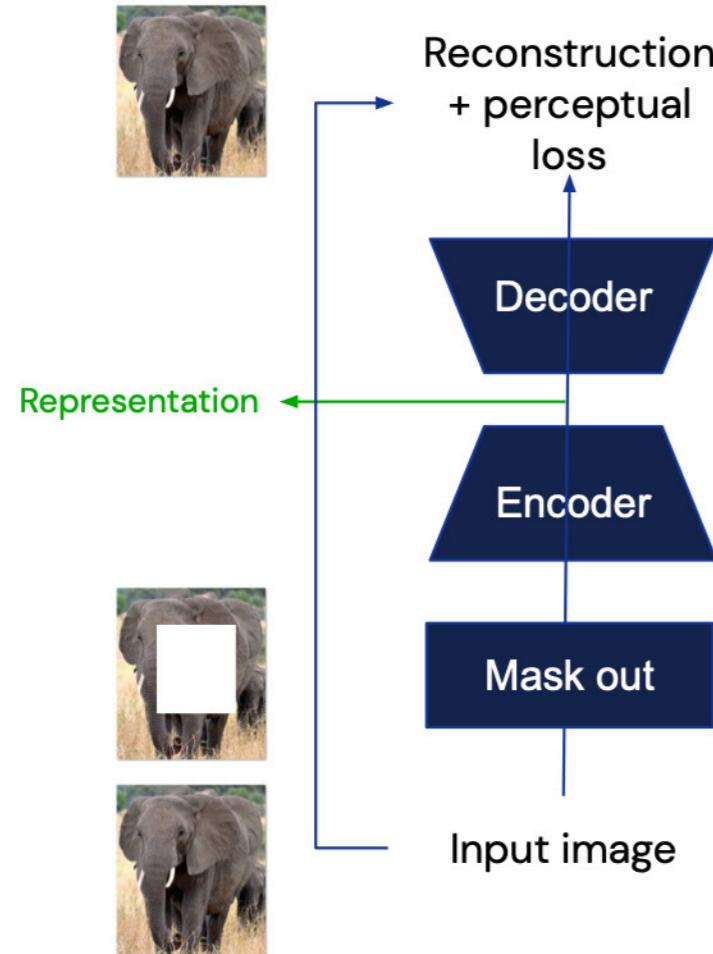
Natural language processing (e.g. word2vec, BERT)

All [MASK] have tusks. ⇒ All **elephants** have tusks.

[“Distributed representations of words and phrases and their compositionality”, Mikolov et al. 13]

[“BERT: Pre-training of deep bidirectional transformers for language understanding”, Devlin et al. 18]

# Context encoders



## Pros

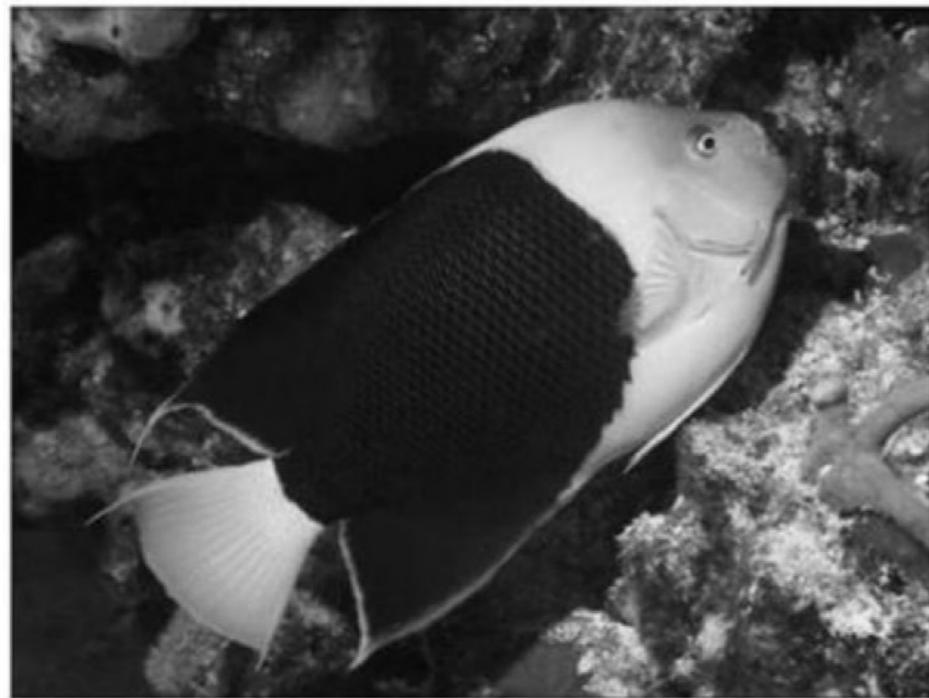
- Requires preservation of fine-grained information

## Cons

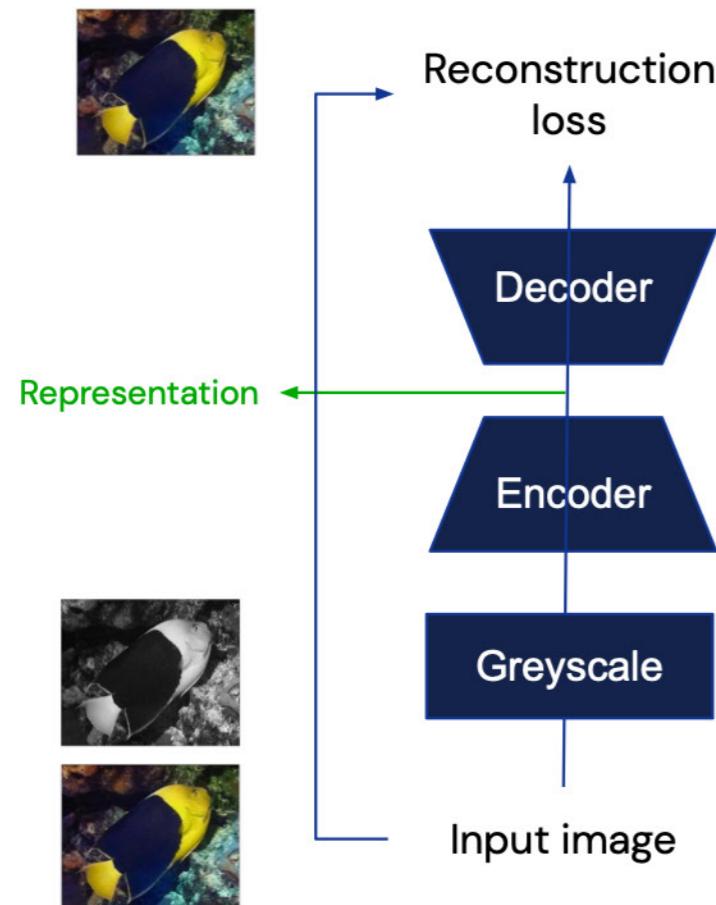
- Train-eval gap: no masking at eval
- Reconstruction is too hard and ambiguous
- Lots of effort spent on “useless” details: exact colour, good boundary, etc

## Colorization

What is the colour of every pixel? Hard if you don't recognize the object!



## Context encoders



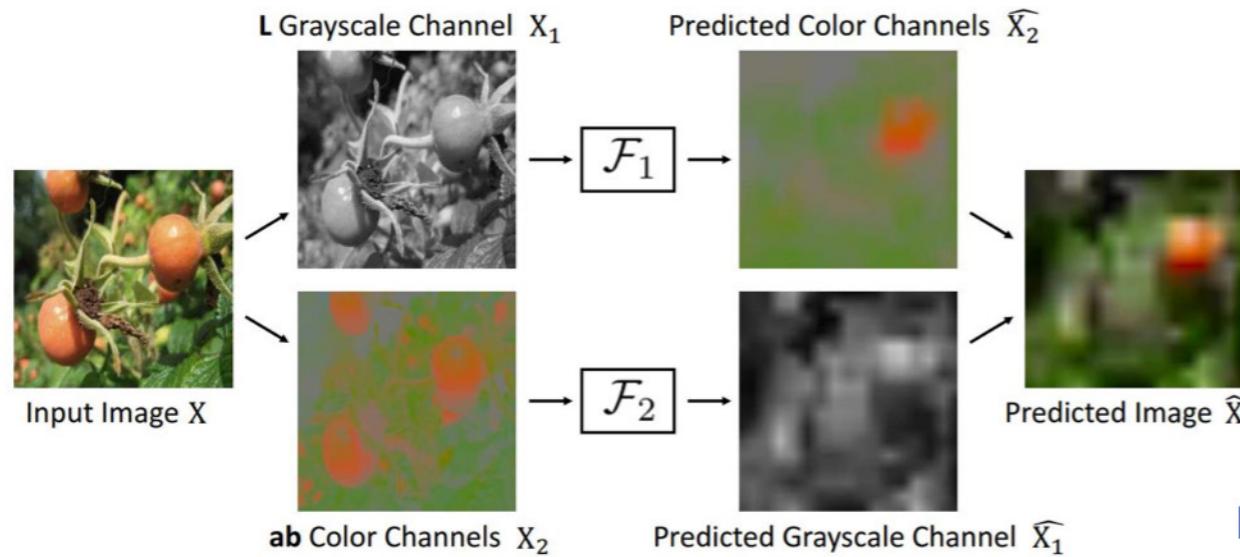
### Pros

- Requires preservation of fine-grained information

### Cons

- Reconstruction is too hard and ambiguous
- Lots of effort spent on “useless” details: exact colour, good boundary, etc
- Forced to evaluate on greyscale images, losing information

## Context encoders $\Rightarrow$ Split-brain encoders



### Pros

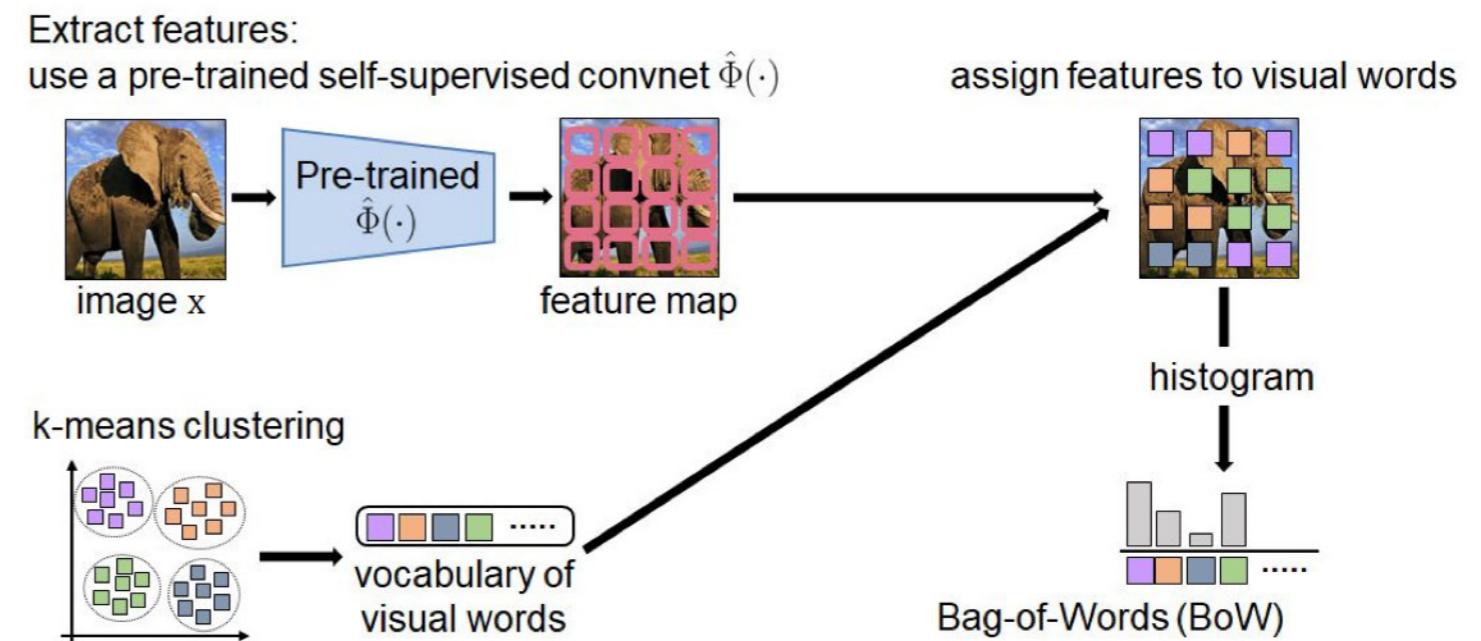
- Requires preservation of fine-grained information

### Cons

- Reconstruction is too hard and ambiguous
- Lots of effort spent on “useless” details: exact colour, good boundary, etc
- ~~Forced to evaluate on grayscale images, losing information~~
- Processes different chunks of the input independently

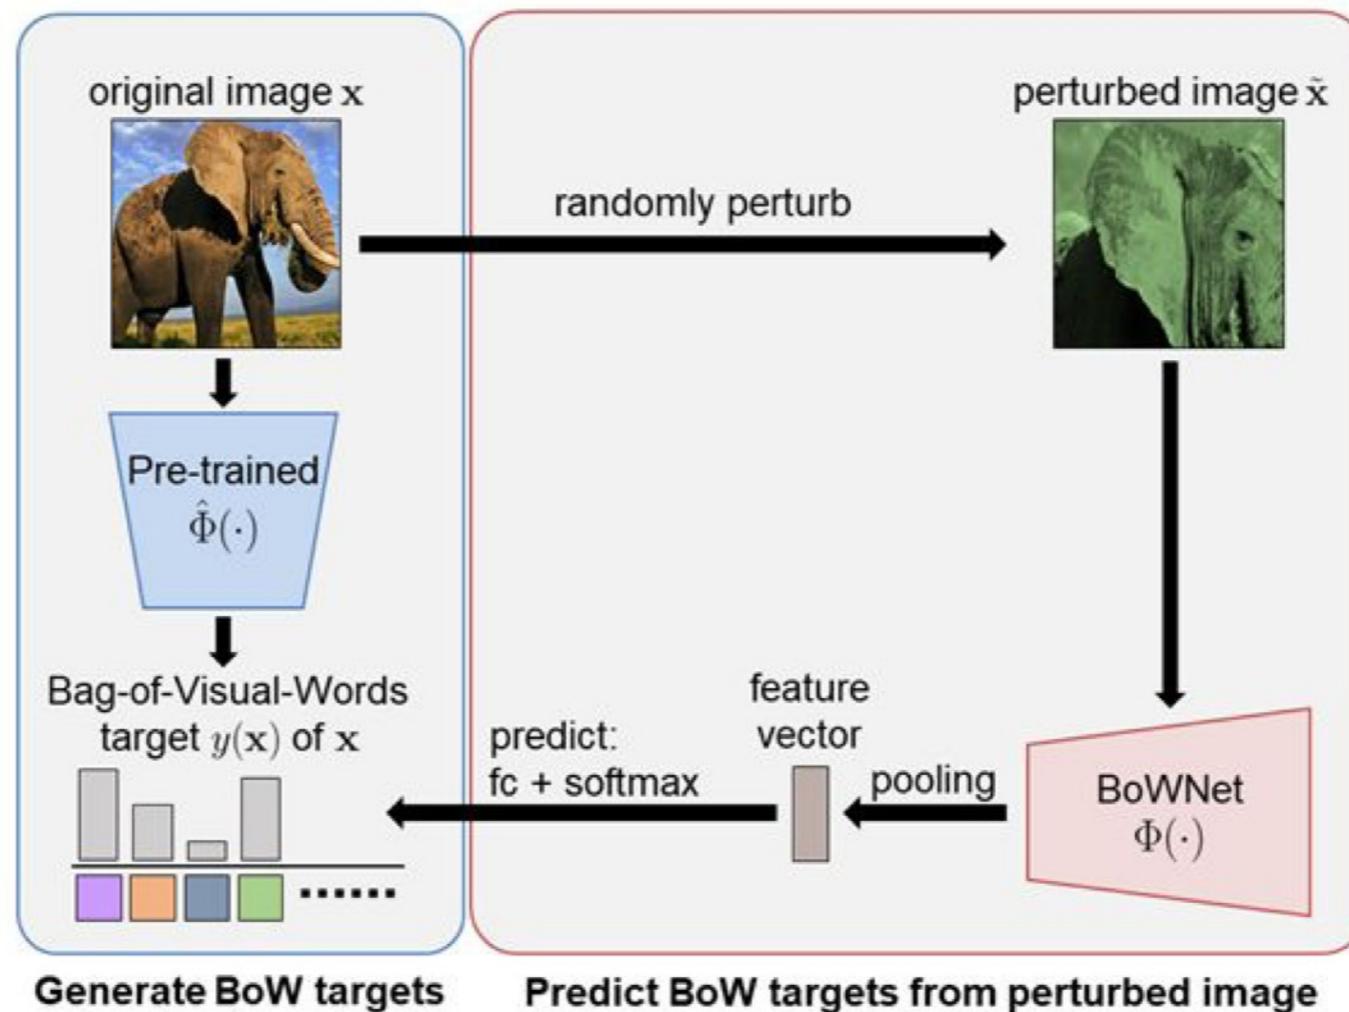
# Predicting bag-of-words

## Bag-of-words reminder



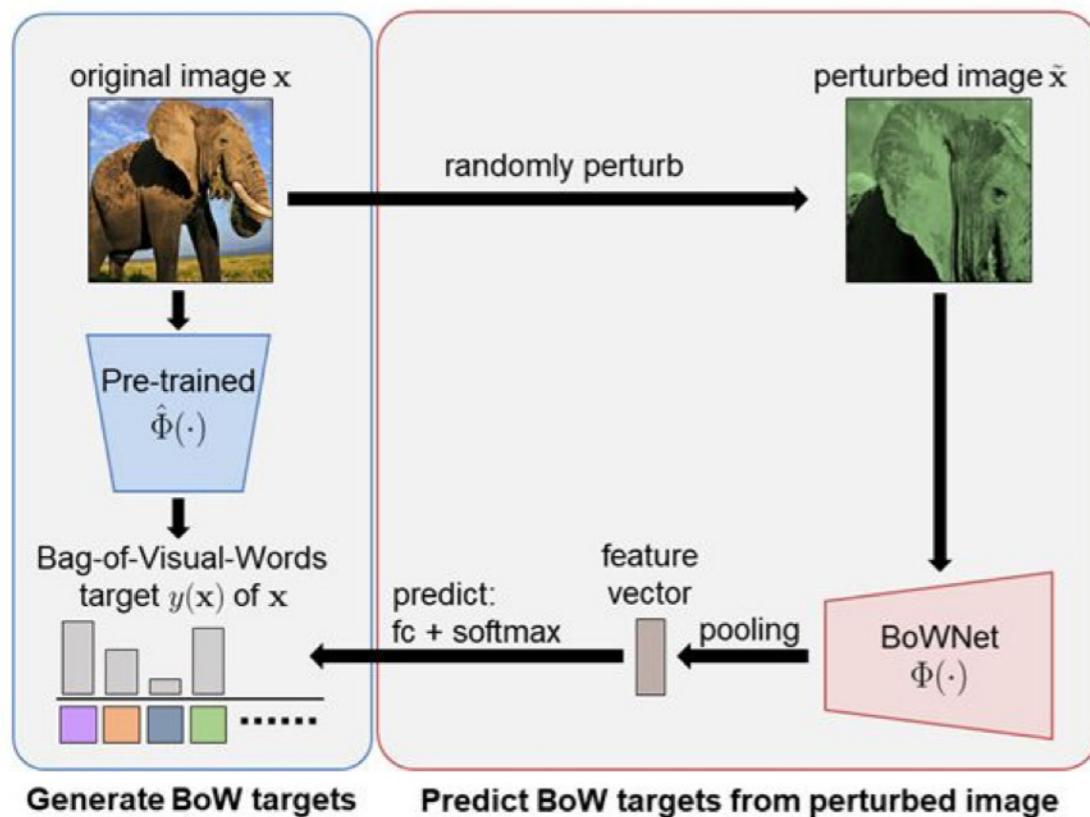
- Loses low-level details
- Encodes mid/high-level concepts

## Predicting bag-of-words



Inspired by NLP: targets = discrete concepts (words)

# Predicting bag-of-words



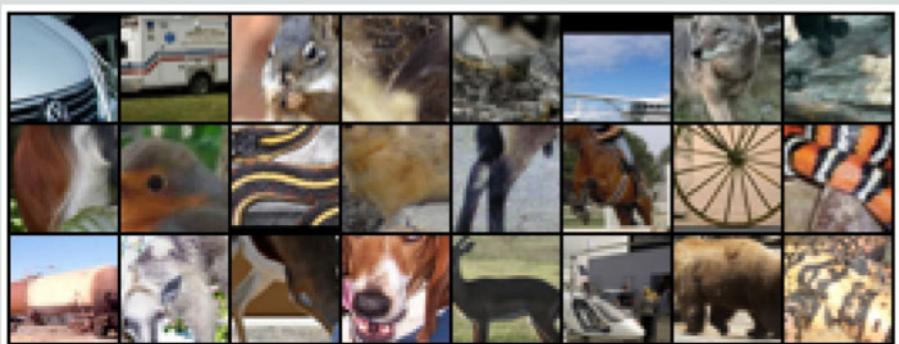
## Pros

- Representations are invariant to desired transformations
- Learn contextual reasoning skills
  - Infer words of missing image regions

## Cons

- Requires bootstrapping from another network
  - e.g. hard to learn more fine-grained features
- Pitfalls of BoW
  - (partial) loss of spatial information
  - SpatialBoW not improving

# Instance classification

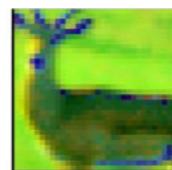


Slide Credit: Relja Arandjelovic

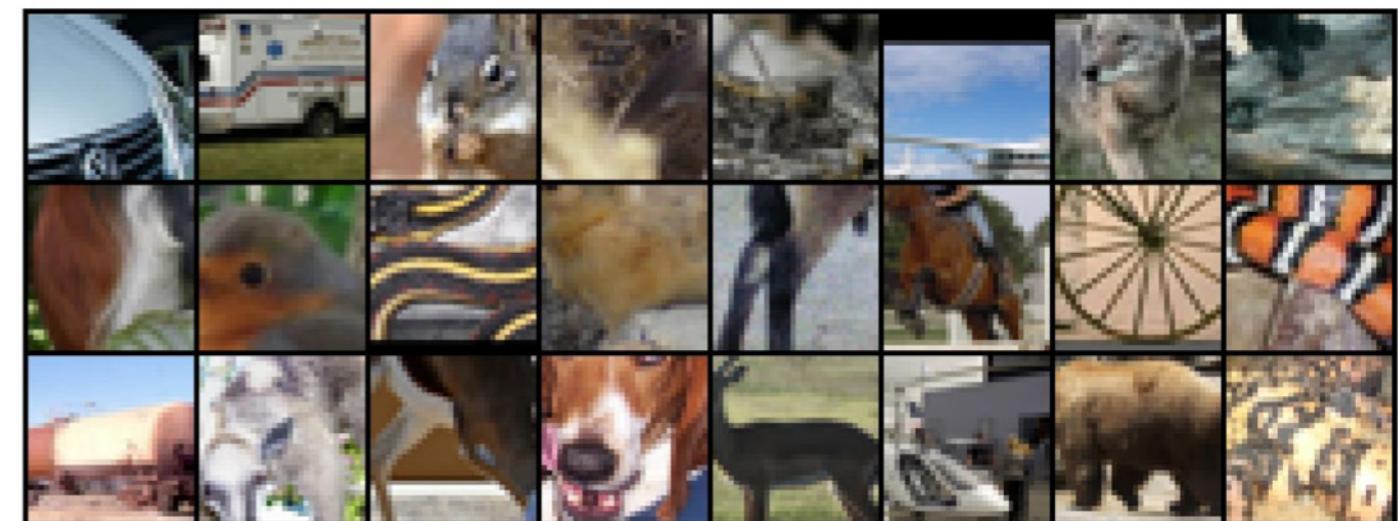
[“Discriminative unsupervised feature learning with exemplar convolutional neural networks”, Dosovitskiy et al.14]

## Exemplar ConvNets

This

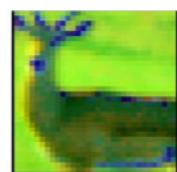


is a distorted crop extracted from an image, which of these crops has the same source image?

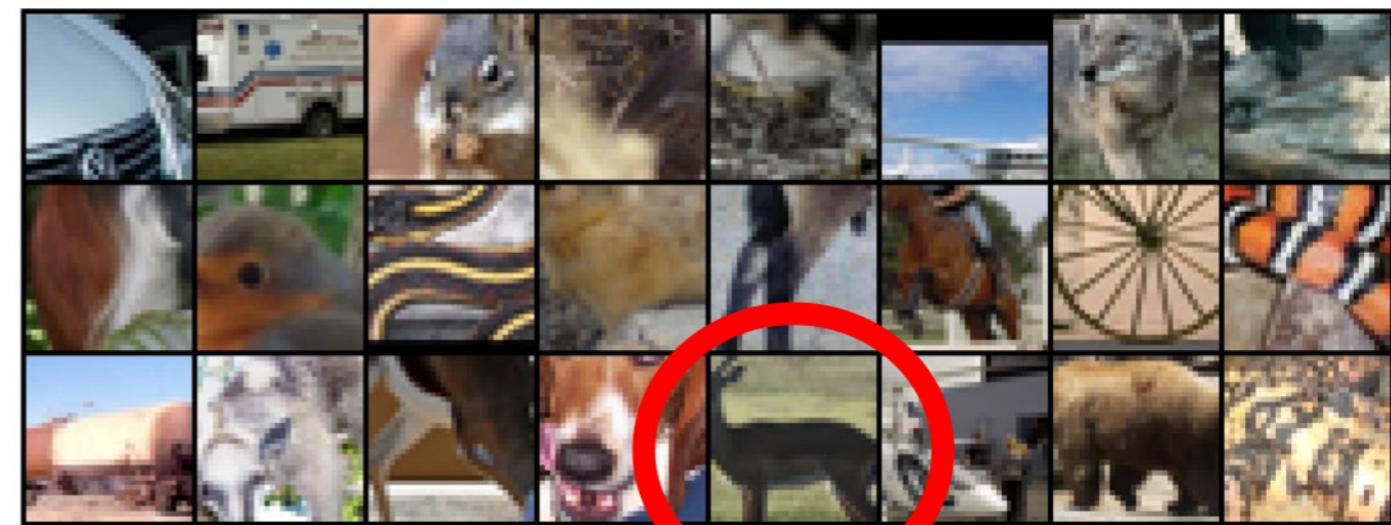


## Exemplar ConvNets

This

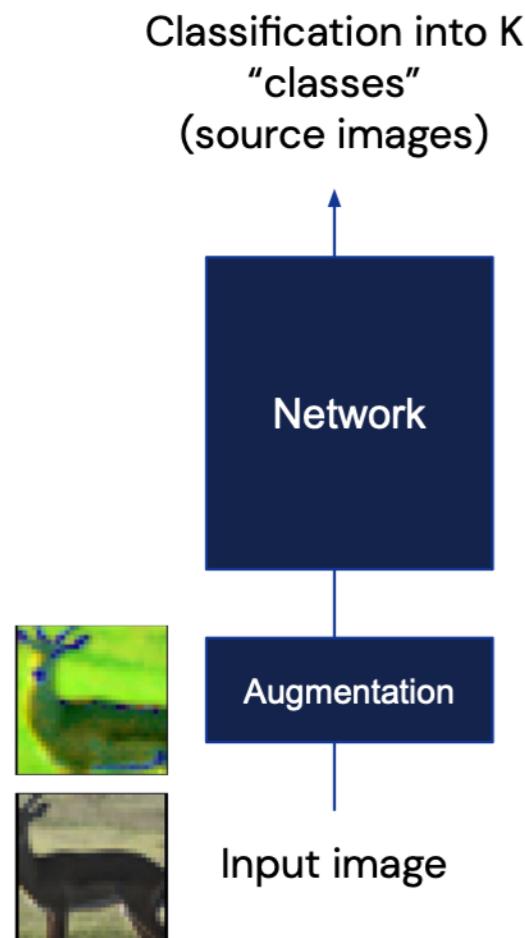


is a distorted crop extracted from an image, which of these crops has the same source image?



Easy if robust to the desired transformations (geometry and colour)

# Exemplar ConvNets



## Pros

- Representations are invariant to desired transformations
- Requires preservation of fine-grained information

## Cons

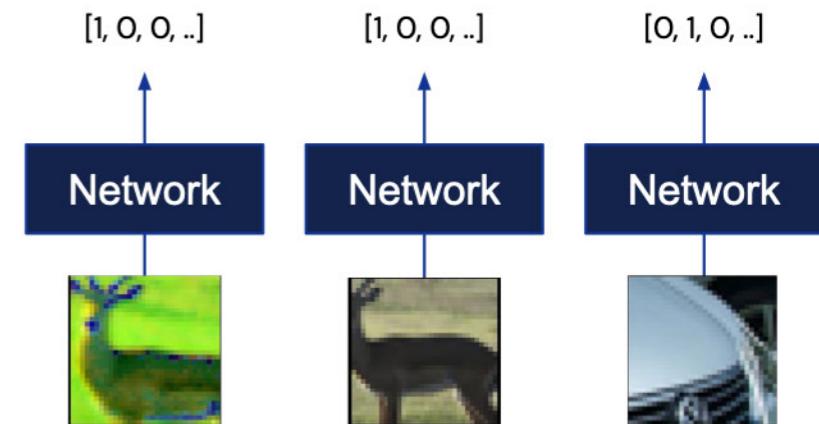
- Choosing the augmentations is important
- Exemplar based: images of the same class or instance are negatives
  - Nothing prevents it from focusing on the background
- Original formulation is not scalable (number of “classes” = dataset size)

# Exemplar ConvNets via metric learning

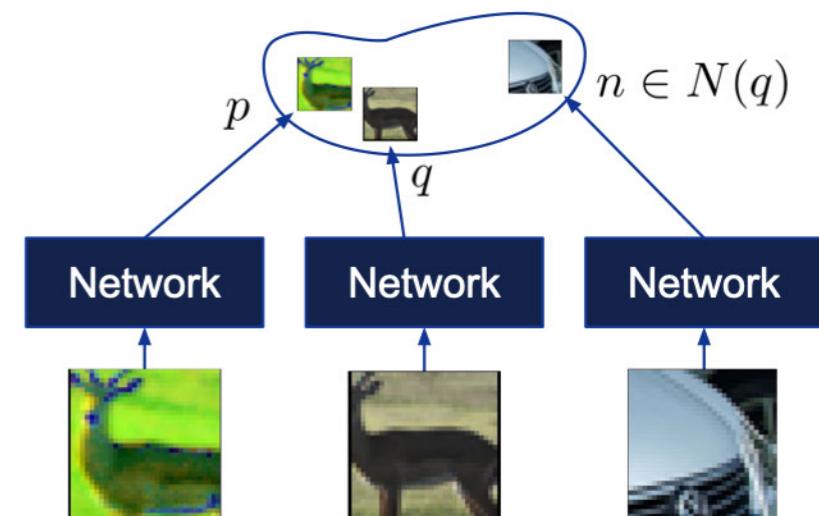
Exemplar ConvNets are not scalable (number of “classes” = number of training images)

- Reformulate in terms of metric learning
- Traditional losses such as contrastive or triplet [“Multi-task self-supervised visual learning”, Doersch and Zisserman 17], [“HowTo100M: Learning a text-video embedding by watching hundred million narrated video clips”, Miech et al. 19]
- Recently popular: InfoNCE [“Representation Learning with Contrastive Predictive Coding”, van den Oord et al. 18]
  - Used by many recent methods: CPC, AMDIM, SimCLR, MoCo, ..

## Classification



## Metric learning



# Noise Contrastive Estimation

InfoNCE loss (a specific popular version)

- For **query**, **positive** and **negative**:

$$-\log \frac{\exp(q^T p)}{\exp(q^T p) + \sum_{n \in N(q)} \exp(q^T n)}$$

- Like a ranking loss:  $(q,p)$  should be close,  $(q,n)$  should be far
- An implementation

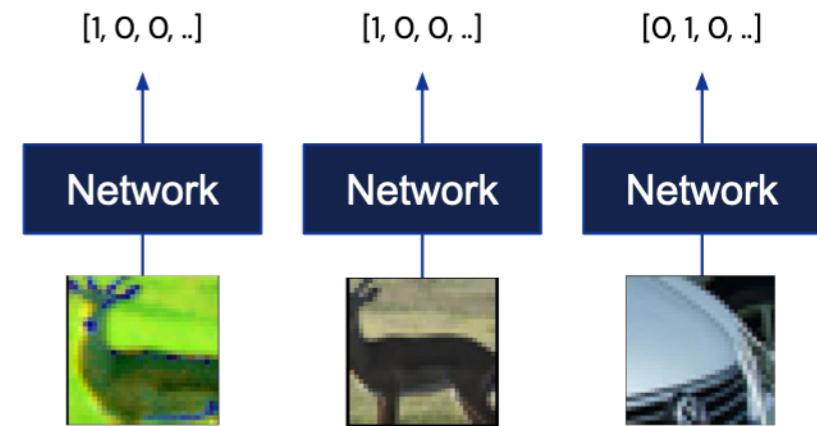
$$\text{logits} = [q^T p, q^T n_1, q^T n_2, \dots] = q^T [p, n_1, n_2, \dots]$$

$$\text{labels} = [1, 0, 0, \dots]$$

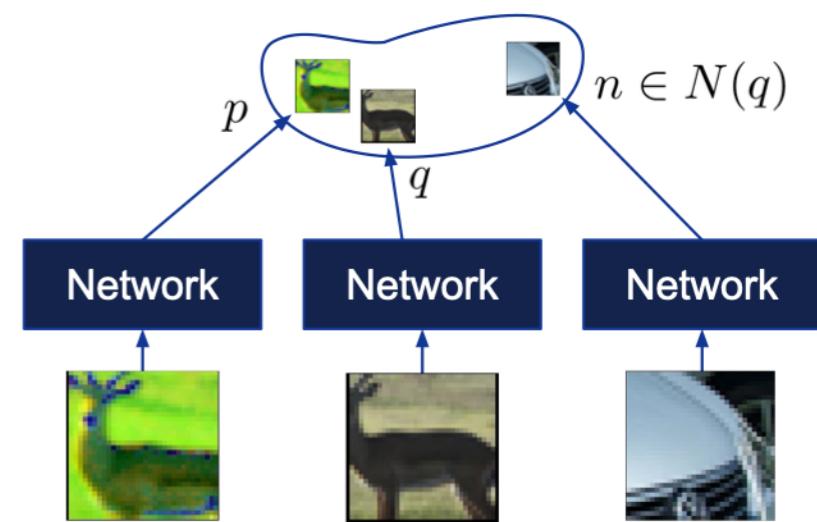
$$\text{InfoNCE} = \text{cross\_entropy}(\text{softmax}(\text{logits}), \text{labels})$$

- Squint and see classification loss
  - Replace  $[p, n_1, n_2, \dots]$  with  $[w_p, w_{n_1}, w_{n_2}, \dots]$
  - Like classification with weight=exemplars
- More details and perspectives in the next part

## Classification



## Metric learning

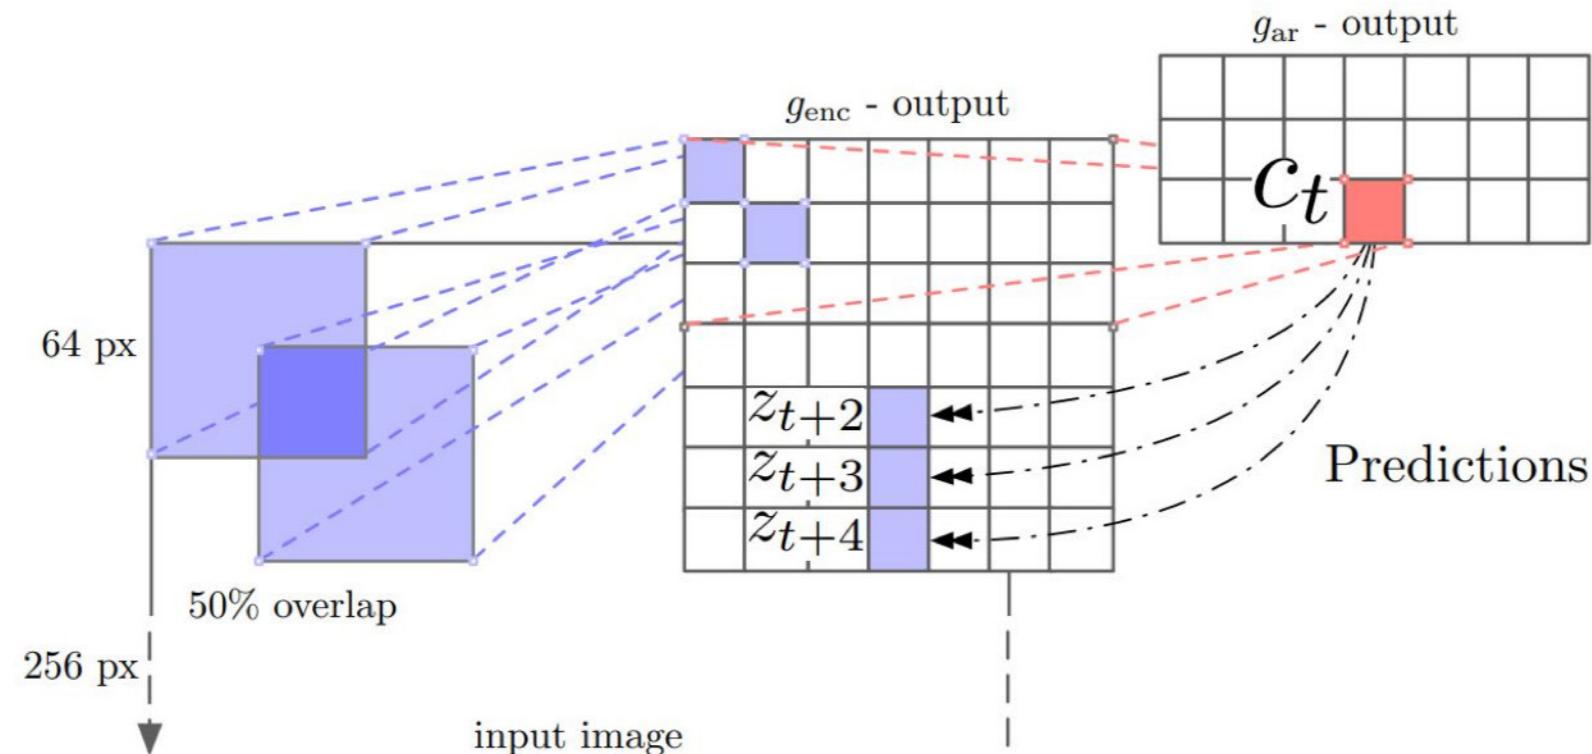


[“Representation Learning with Contrastive Predictive Coding”, van den Oord et al. 18]  
[“Data-efficient image recognition with contrastive prediction coding”, Hénaff et al. 19]

# Contrastive predictive coding (CPC)

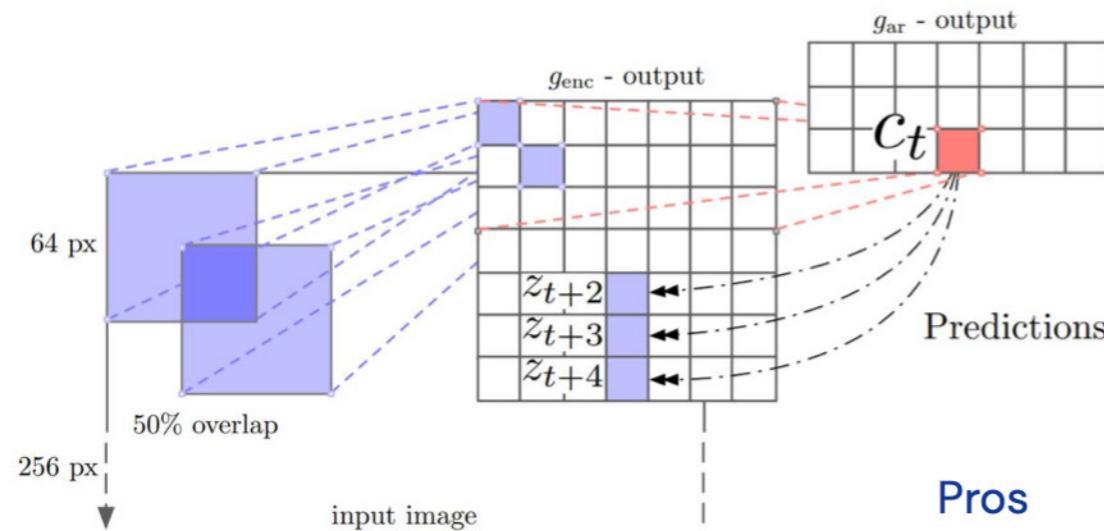
Roughly: Context Prediction + Exemplar ConvNets

- From a patch, predict representations of other patches below it
- Use InfoNCE loss to contrast the (predictions, correct, negatives)
  - Negatives: other patches from the same image and other images



[“Representation Learning with Contrastive Predictive Coding”, van den Oord et al. 18]  
[“Data-efficient image recognition with contrastive prediction coding”, Hénaff et al. 19]

# Contrastive predictive coding (CPC)



## Pros

- Generic framework easily applied to images, video, audio, NLP, ..
- Exemplar: Requires preservation of fine-grained information
- Context prediction: Should enable learning about object parts

## Cons

- Exemplar based: images of the same class or instance are negatives
- Train-eval gap: training on patches, evaluating on images
- Assumes training images are photographed with canonical orientations (and canonical orientations exist)
- Somewhat slow training due to dividing into patches

# Exploiting time



Slide Credit: Relja Arandjelovic

[“Learning features by watching objects move”, Pathak et al. 16]

## Watching objects move

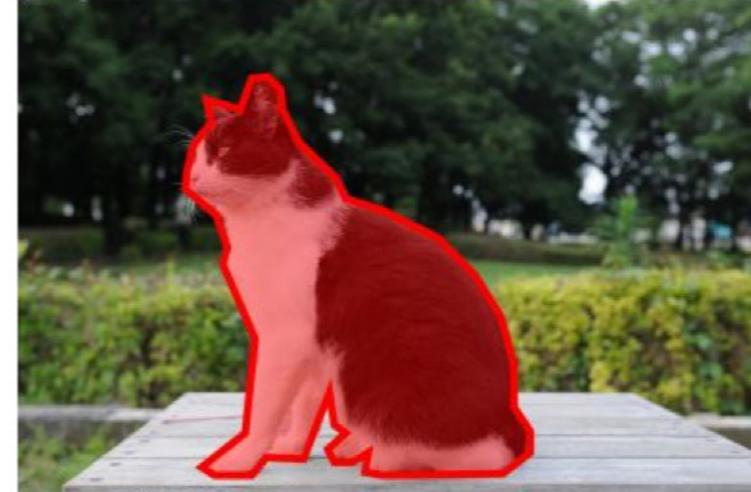
Which pixels will move?



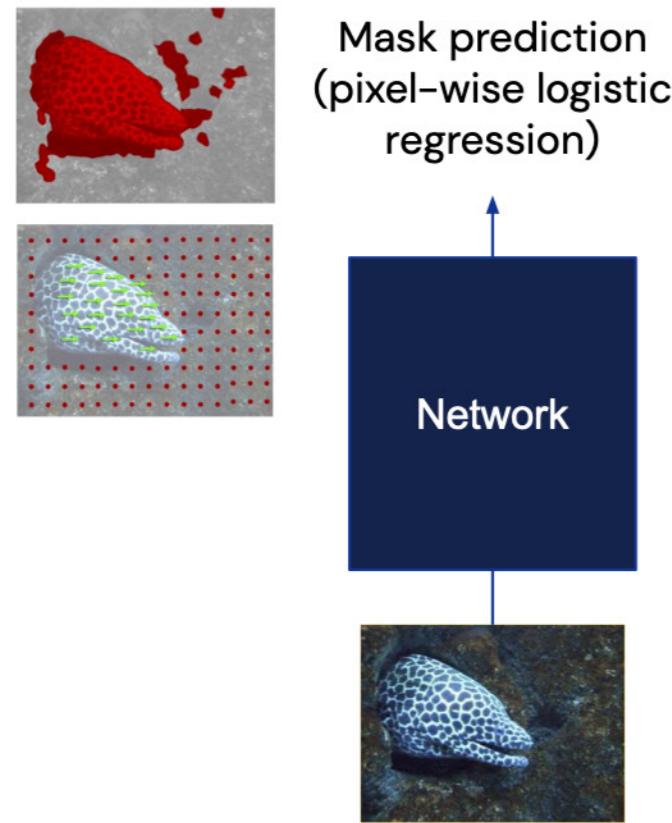
[“Learning features by watching objects move”, Pathak et al. 16]

## Watching objects move

Which pixels will move? **Easy if we can segment objects!**



## Watching objects move



### Pros

- Emerging behaviour: segmentation
- No train–eval gap

### Cons

- “Blind spots”: stationary objects
- Potential focus on large salient objects
- Depends on an external motion segmentation algorithm
- Cannot be extended to temporal nets (pretext task would be trivial)

[“Tracking emerges by colorizing videos”, Vondrick et al. 18]

## Tracking by colorization

Given an earlier frame, colourize the new one.



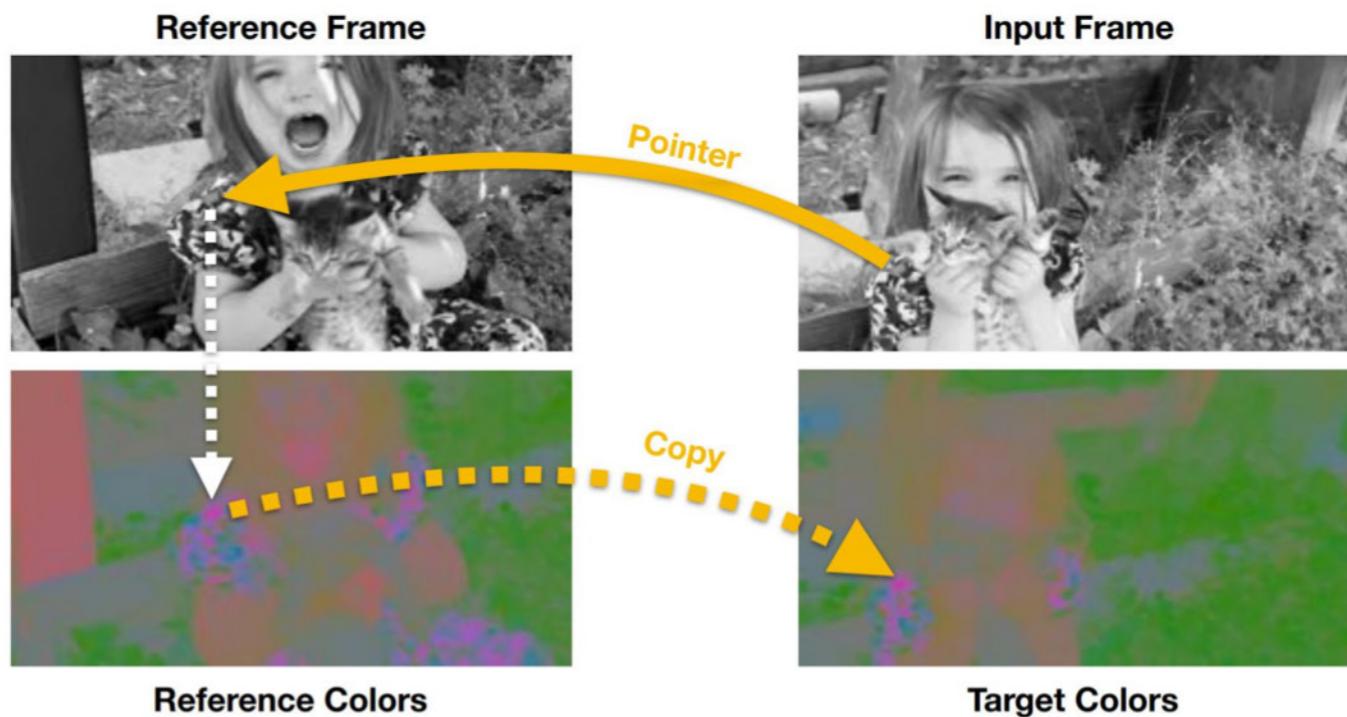
[“Tracking emerges by colorizing videos”, Vondrick et al. 18]

## Tracking by colorization

Given an earlier frame, colourize the new one. **Easy if everything can be tracked!**



## Tracking by colorization



### Pros

- Emerging behaviour: tracking, matching, optical flow, segmentation

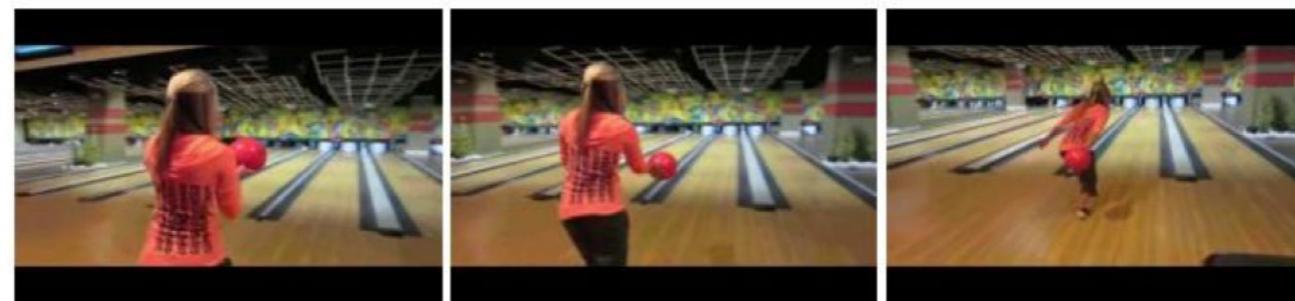
### Cons

- Low level cues are effective – less emphasis on semantics
- Forced to evaluate on greyscale frames, losing information

[“Shuffle and learn: Unsupervised learning using temporal order verification”, Misra et al. 16]

## Temporal ordering

Is this sequence of frames correctly ordered?



[“Shuffle and learn: Unsupervised learning using temporal order verification”, Misra et al. 16]

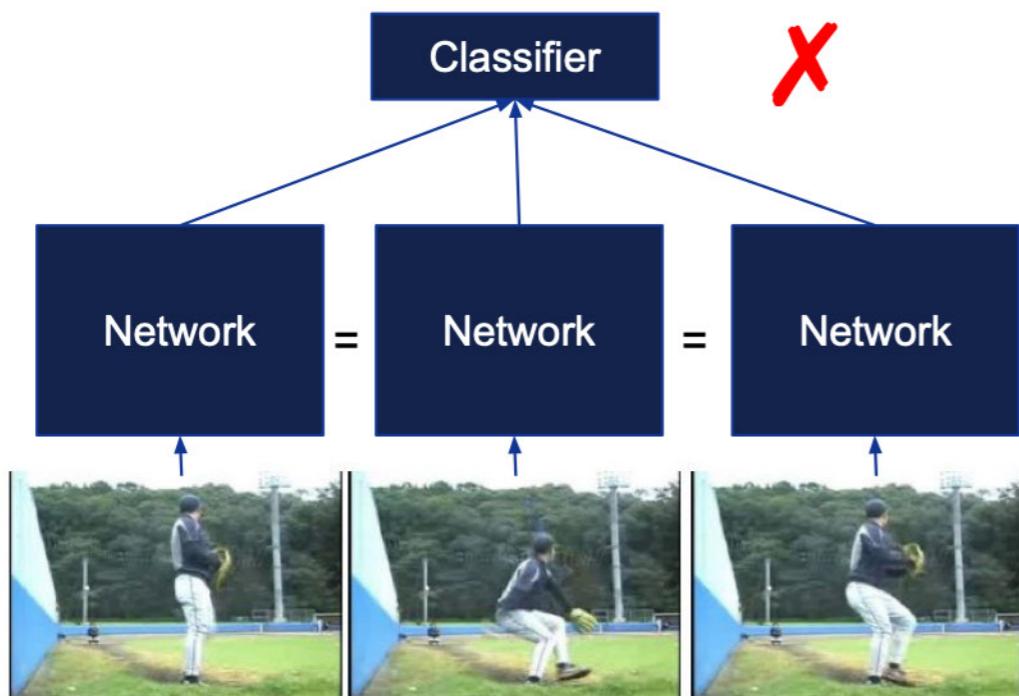
## Temporal ordering

Is this sequence of frames correctly ordered? Easy if we recognize the action and human pose!



[“Shuffle and learn: Unsupervised learning using temporal order verification”, Misra et al. 16]

## Temporal ordering



### Pros

- No train-eval gap
- Learns to recognize human pose

### Cons

- Mostly focuses on human pose – not always sufficient
- Questionable if it can be extended to temporal nets (potentially task becomes too easy)

### Extensions

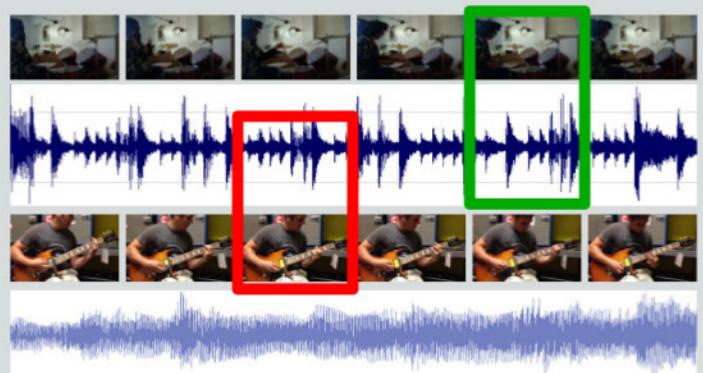
- N frames with one randomly placed – find it

[“Self-supervised video representation learning with odd-one-out networks”, Fernando et al. 16]

- Ranking loss: embeddings should be similar for frames close in time and dissimilar for far away frames

[“Time-contrastive networks: Self-supervised learning from video”, Sermanet et al. 17]

# Multimodal



Slide Credit: Relja Arandjelovic

## Audio-visual correspondence

Does the sound go with the image?

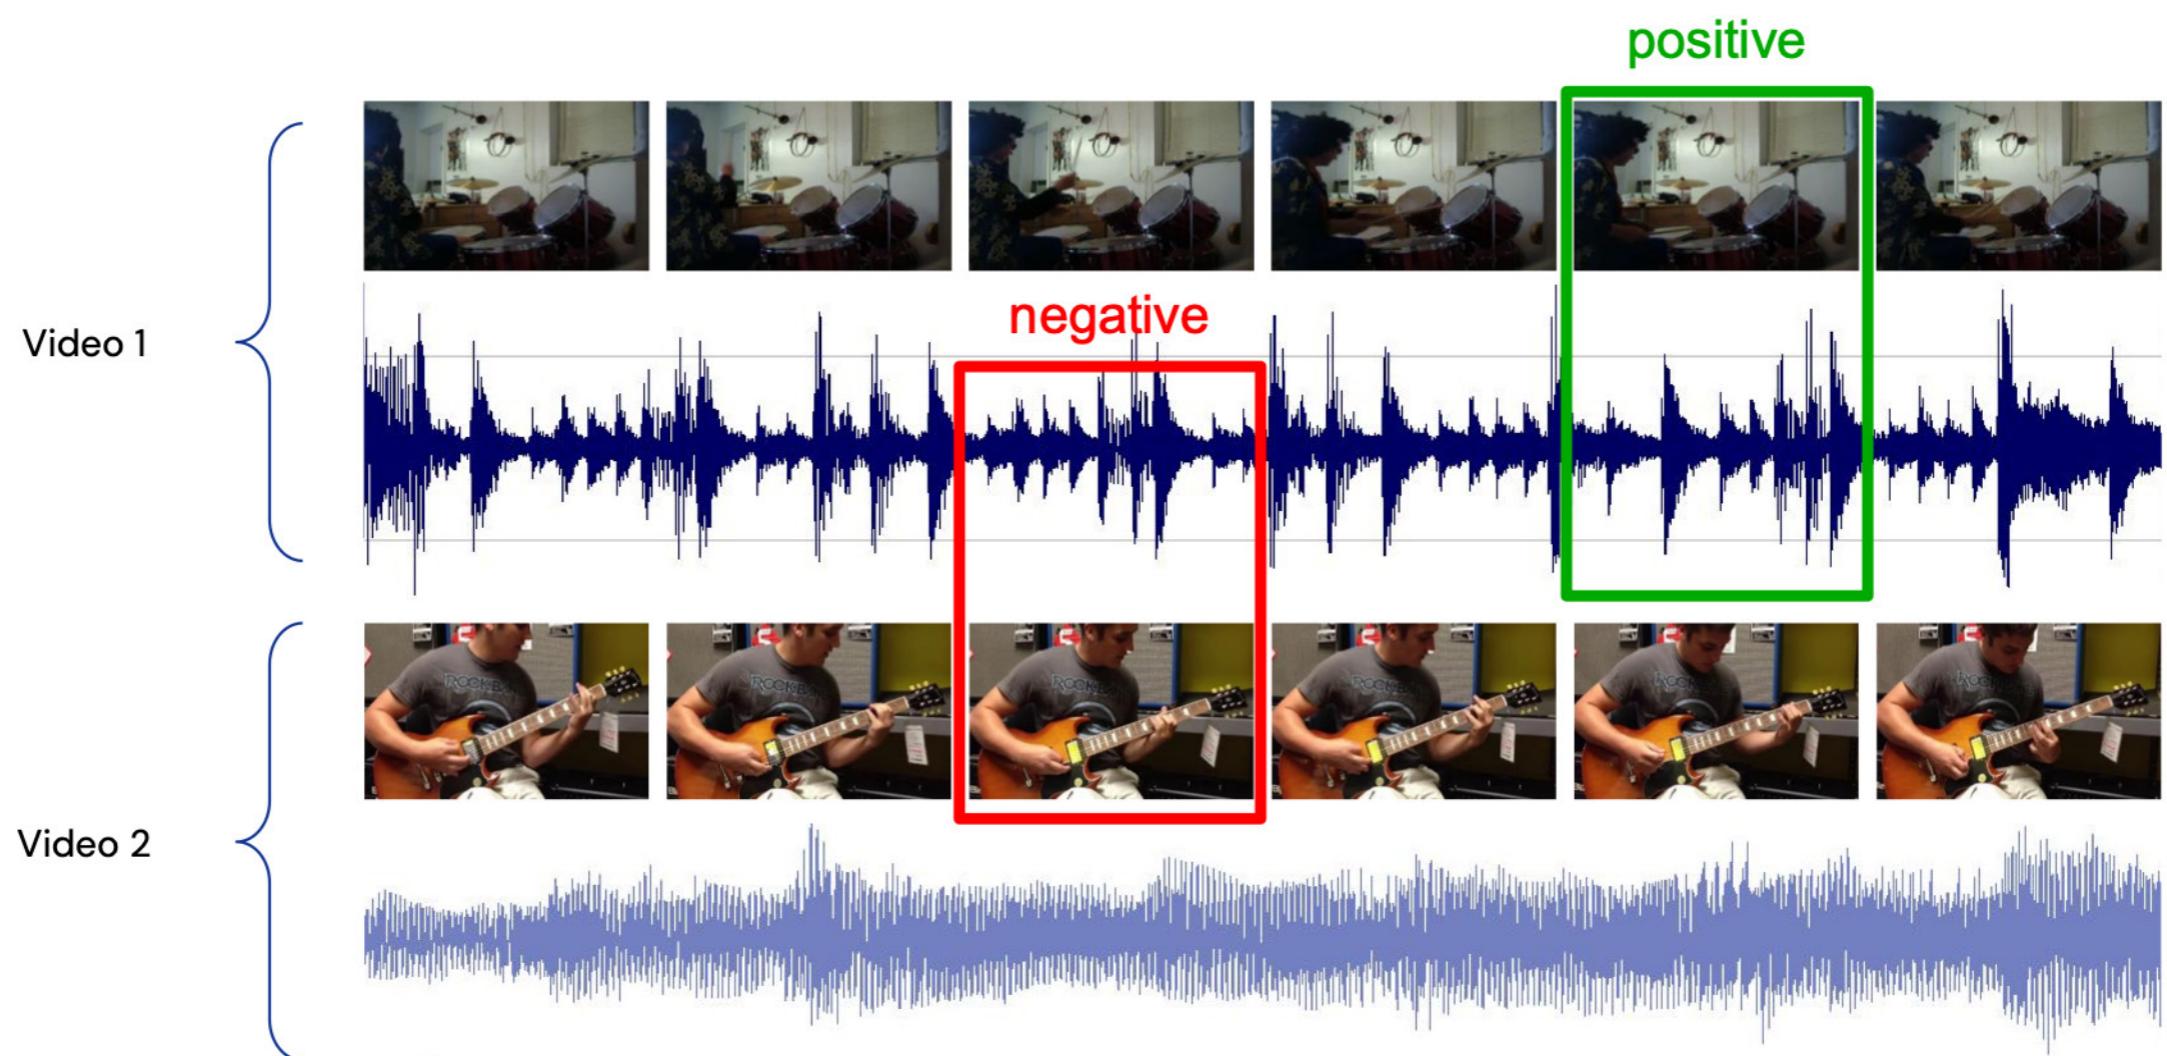


## Audio-visual correspondence

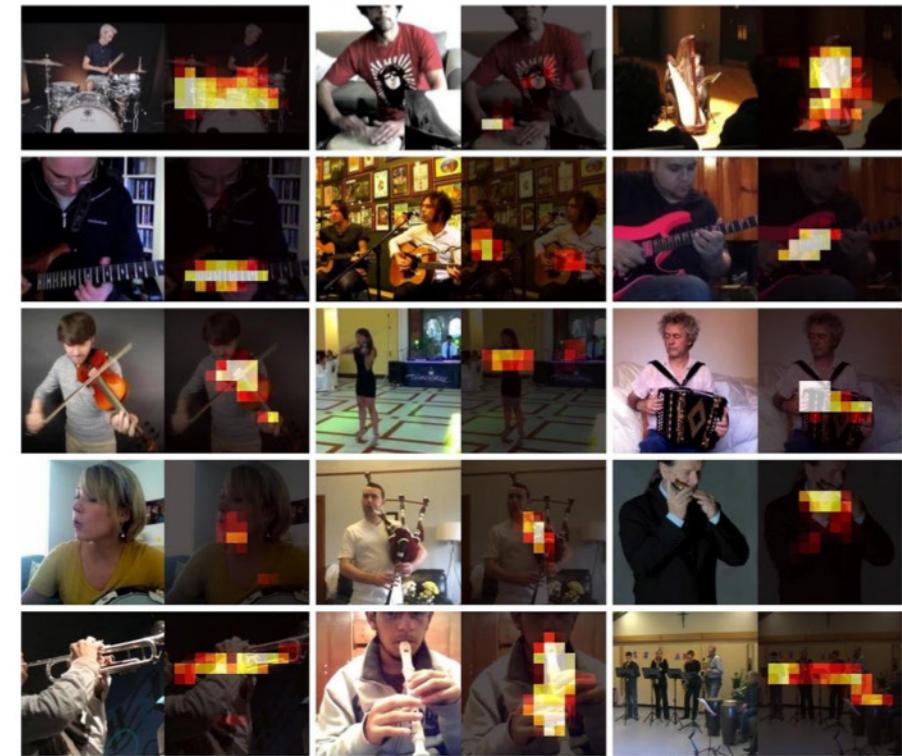
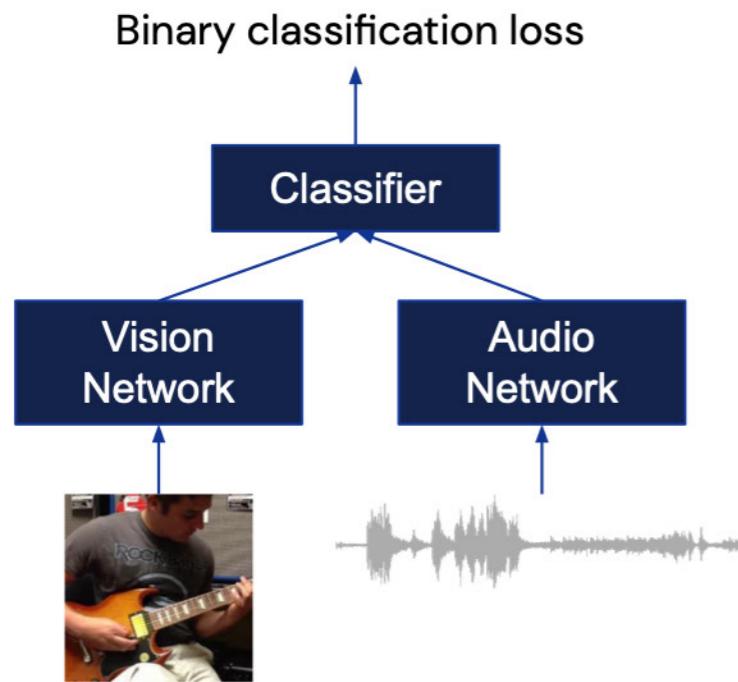
Does the sound go with the image? Easy if we recognize what is happening in both the frame and the audio



## Audio-visual correspondence



# Audio-visual correspondence



[“Objects that sound”, Arandjelović et al. 18]

## Pros

- Natural different views of the training data, no need for augmentations
- No train–eval gap
- Representations in both modalities for free

## Cons

- “Blind spots”: not everything makes a sound
- Exemplar based: videos of the same class or instance are negatives
- Small output space – two cases (corresponds or not)
  - Can be improved by contrastive approaches

## Leveraging narration

Does the narration go with the video?

(Text obtained from automatic speech recognition)



## Leveraging narration

Does the narration go with the video? **Easy if we recognize what is happening in the video and narrations**  
(Text obtained from automatic speech recognition)

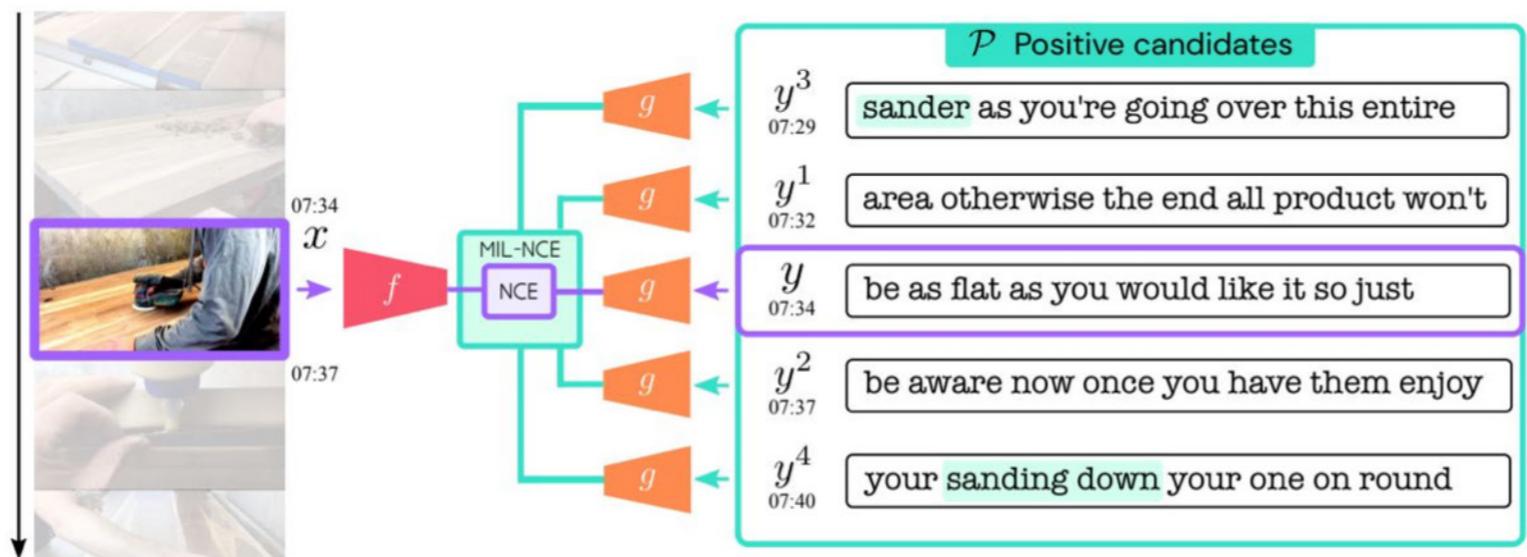


Complication compared to the audio-visual case:

- Narration and visual content are less aligned

# Leveraging narration

Multiple instance learning extension of the NCE loss



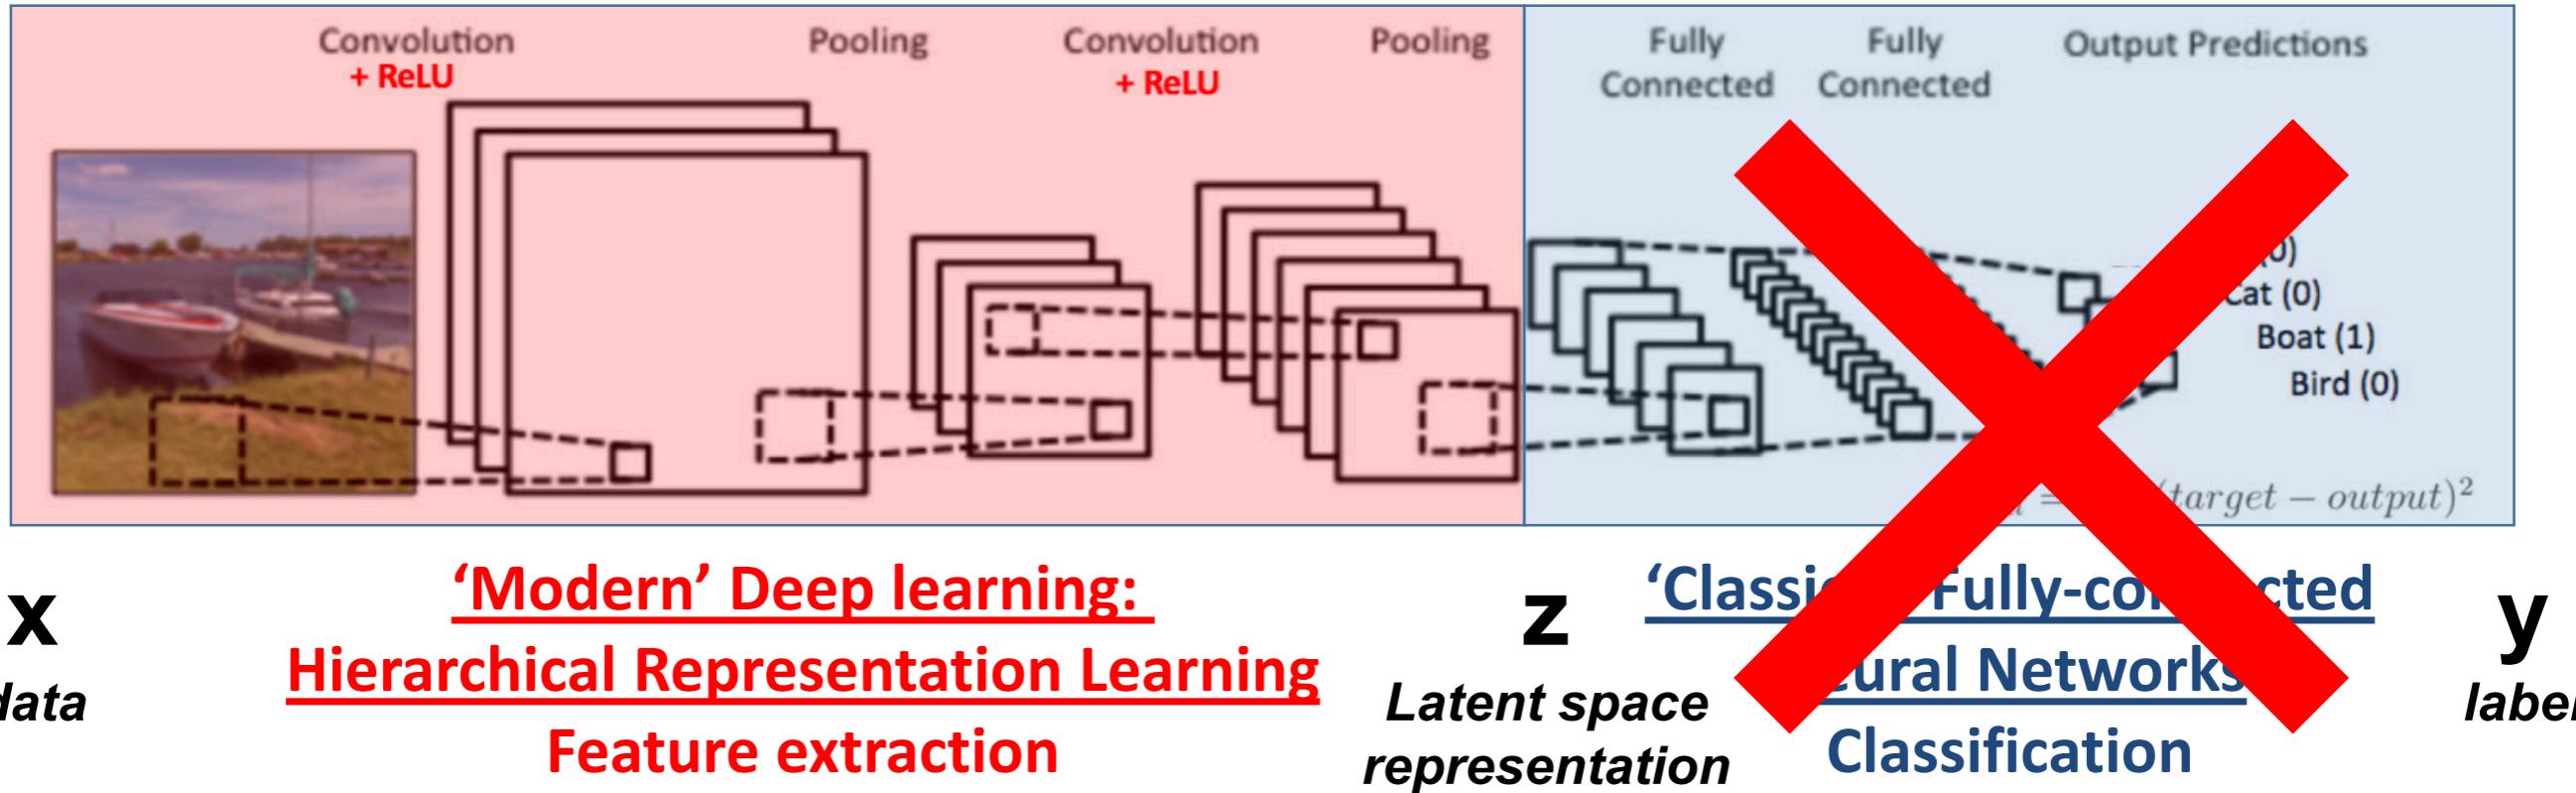
## Pros

- Natural different views of the training data, no need for augmentations
- No train–eval gap
- Representations in both modalities for free

## Cons

- “Blind spots”: not everything is mentioned in narrations
- Exemplar based: videos of the same class or instance are negatives
- Assumes a single language, potentially non-trivial to extend to more

# Representation learning without annotations?



**Many ideas are possible (and yours could be even better!):**

1. Predict the future: RNNs, Video
2. Compression: Autoencoder (predict self, through clamp), representation  $z$
3. Pretext tasks: predict self, before/after, missing patch, correct rotation, colorization, up-sampling, multimodal
4. Capture parameter distribution (variance): Variational Auto-Encoders
5. Make latent space parameters  $z$  meaningful, orthogonal, explicit, tuneable
6. Train using a second network: GANs - Improve quality of output images
7. The Sky is the Limit

# Auto-Encoders

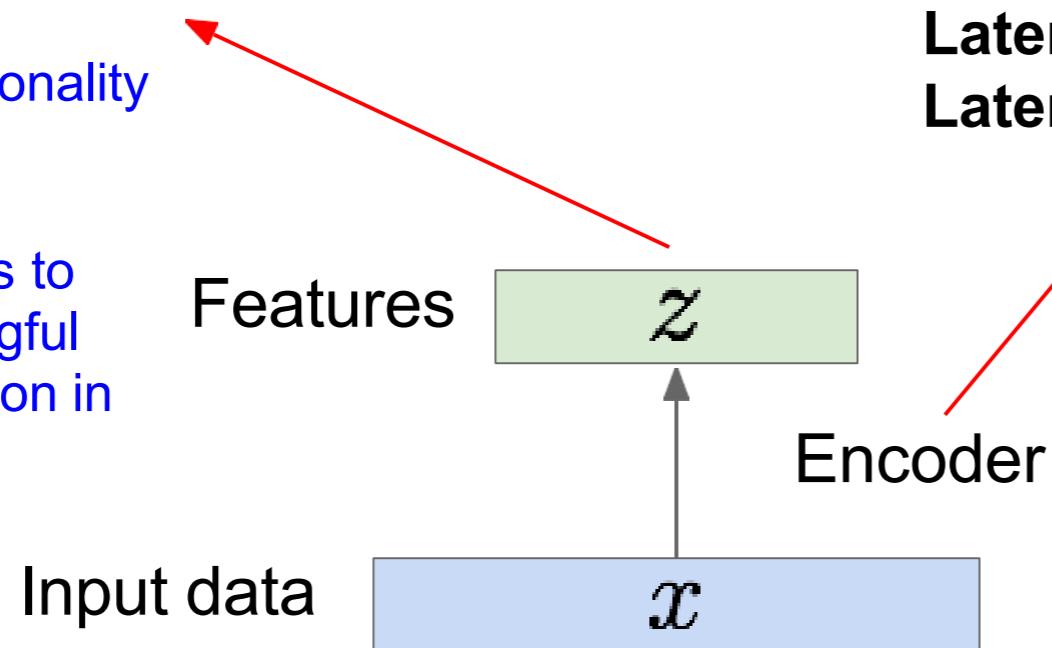
# Some background first: Autoencoders

Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data

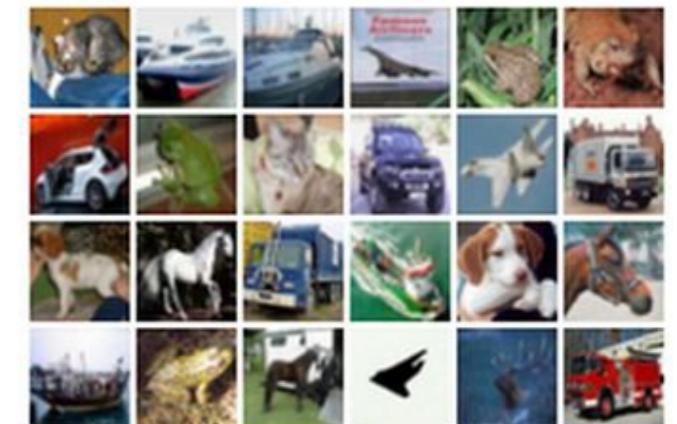
**z** usually smaller than **x**  
(dimensionality reduction)

## Q: Why dimensionality reduction?

A: Want features to capture meaningful factors of variation in data



**Originally:** Linear +  
nonlinearity (sigmoid)  
**Later:** Deep, fully-connected  
**Later:** ReLU CNN

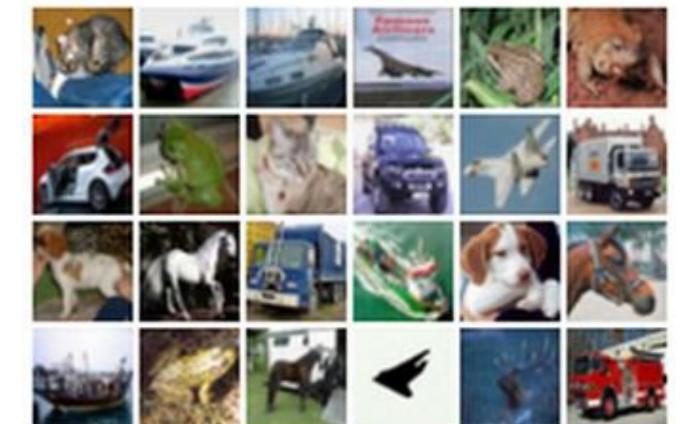
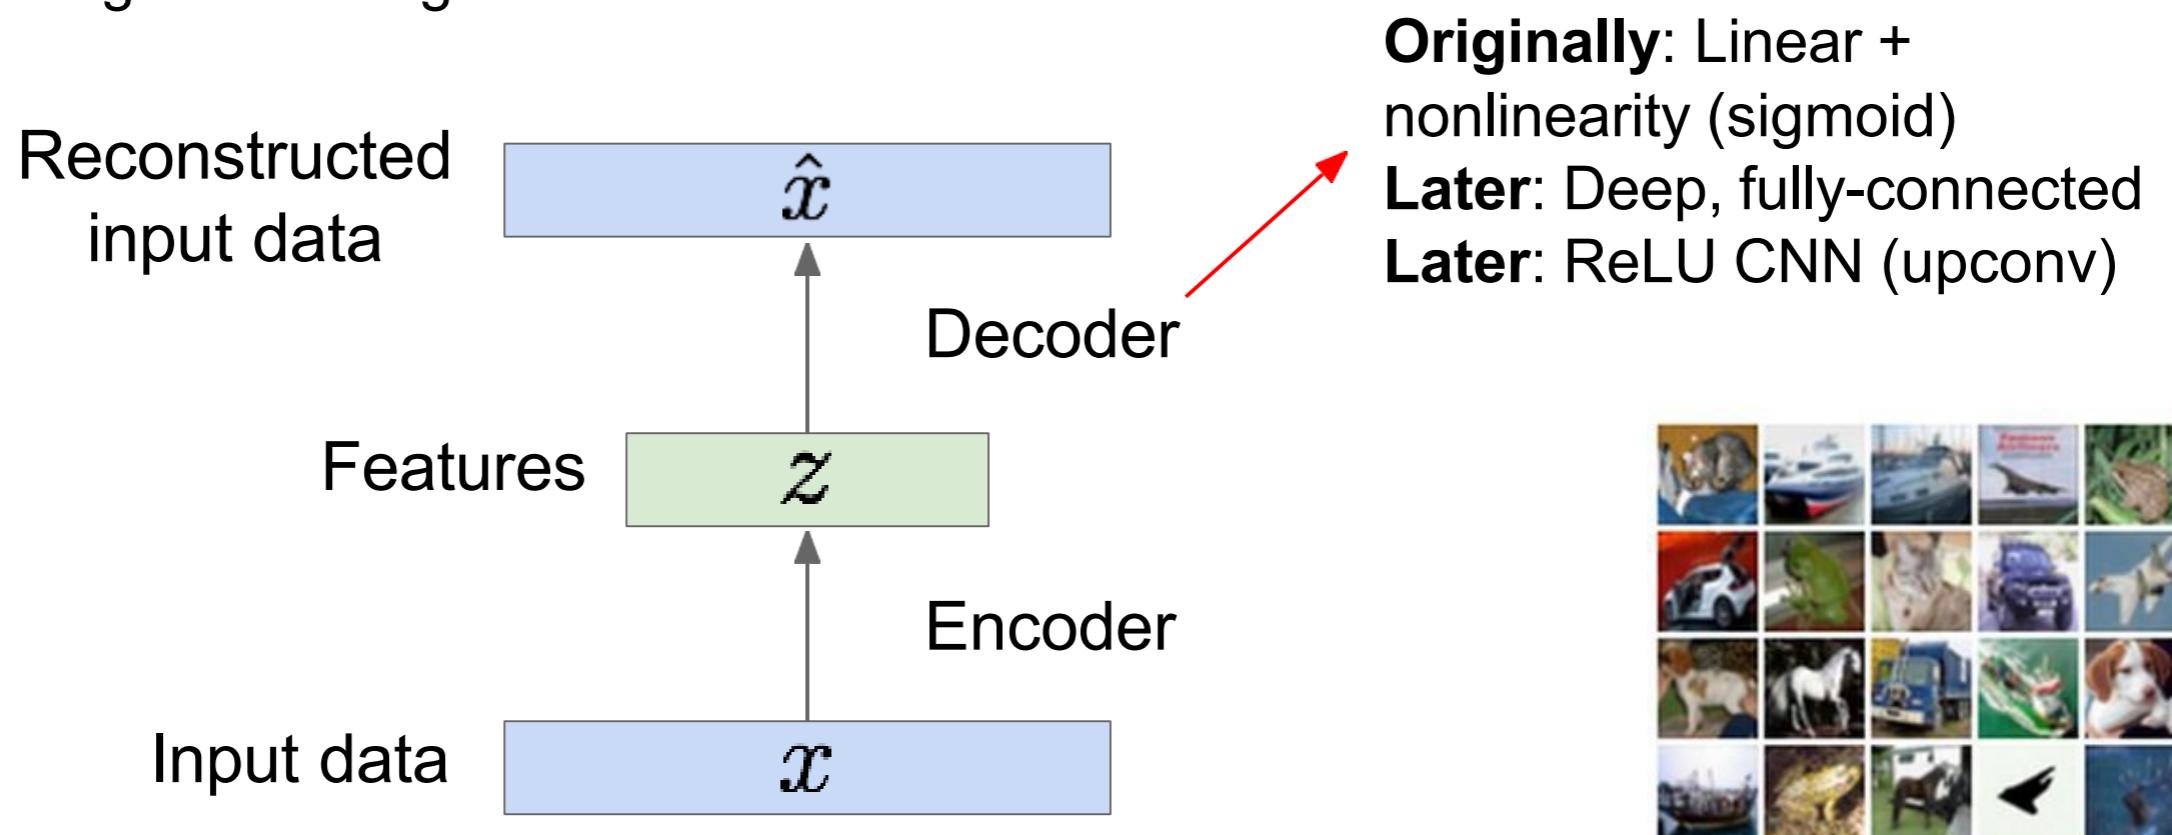


# Some background first: Autoencoders

How to learn this feature representation?

Train such that features can be used to reconstruct original data

“Autoencoding” - encoding itself



# Some background first: Autoencoders

Train such that features can be used to reconstruct original data

Doesn't use labels!

L2 Loss function:

$$\|x - \hat{x}\|^2$$

Reconstructed input data

Features

Input data

$\hat{x}$

Decoder

$z$

Encoder

$x$

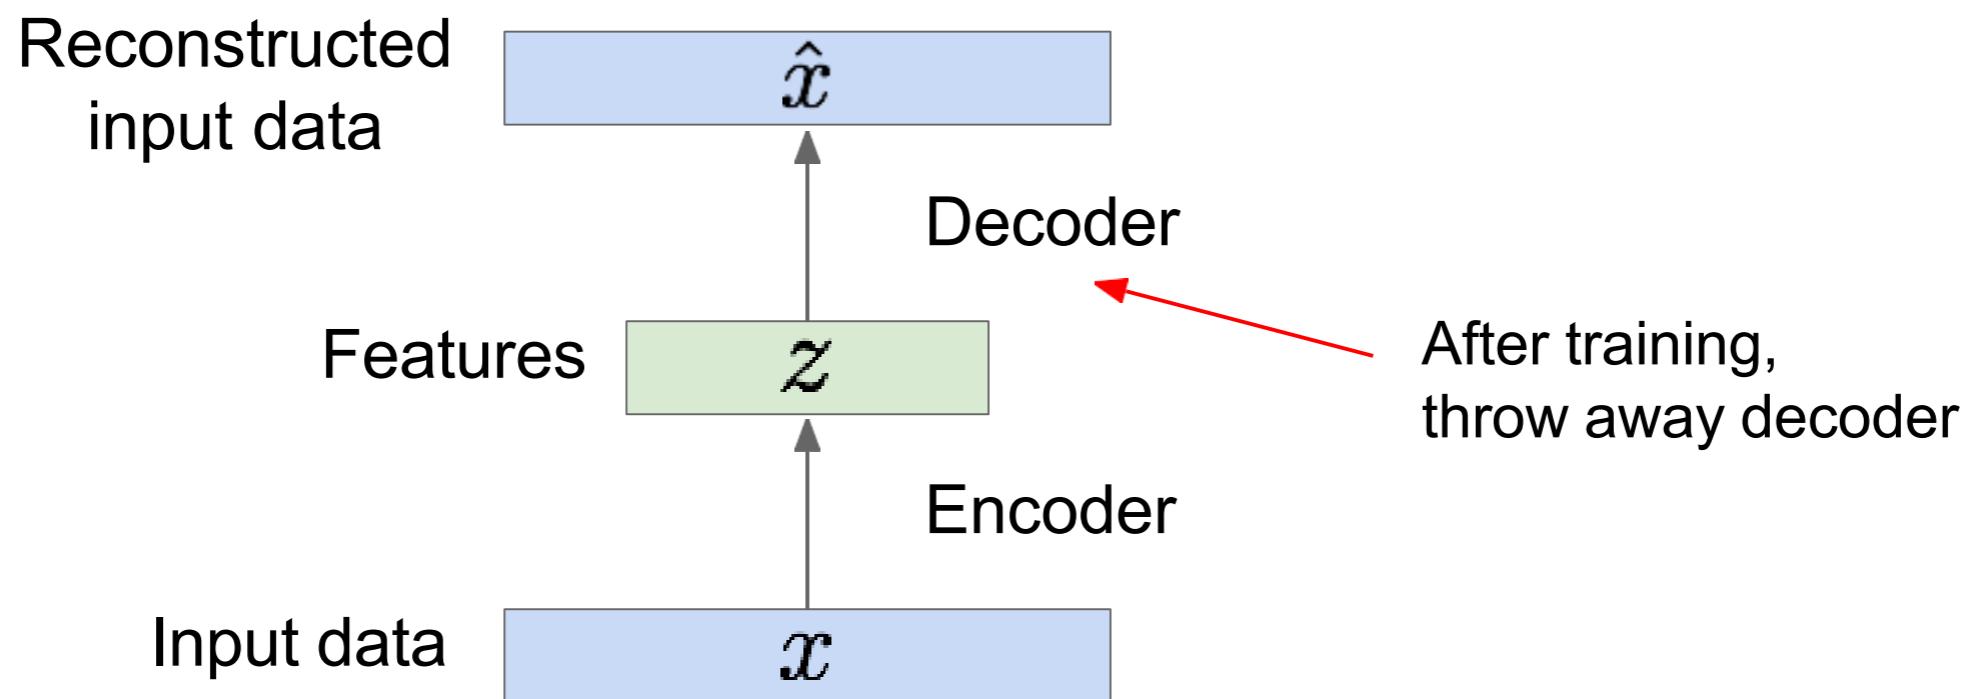
Reconstructed data



Encoder: 4-layer conv  
Decoder: 4-layer upconv

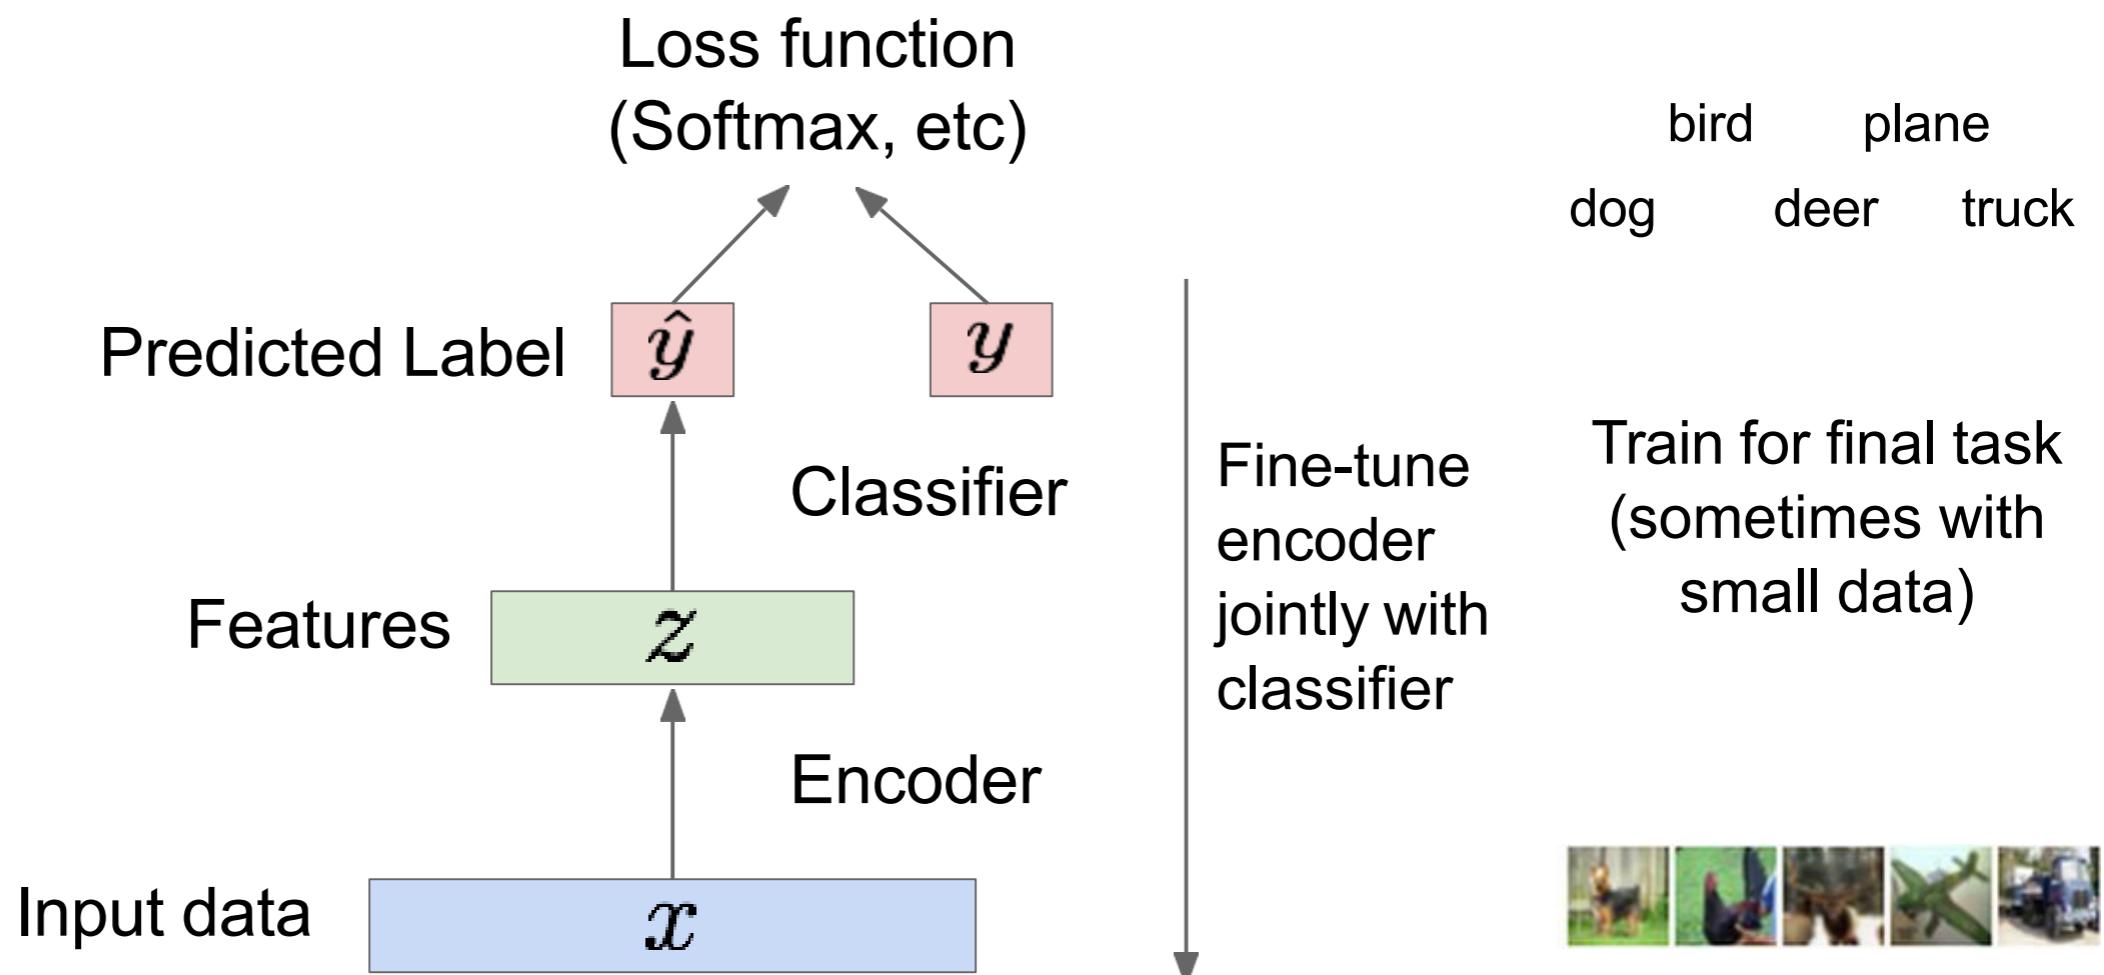


# Some background first: Autoencoders

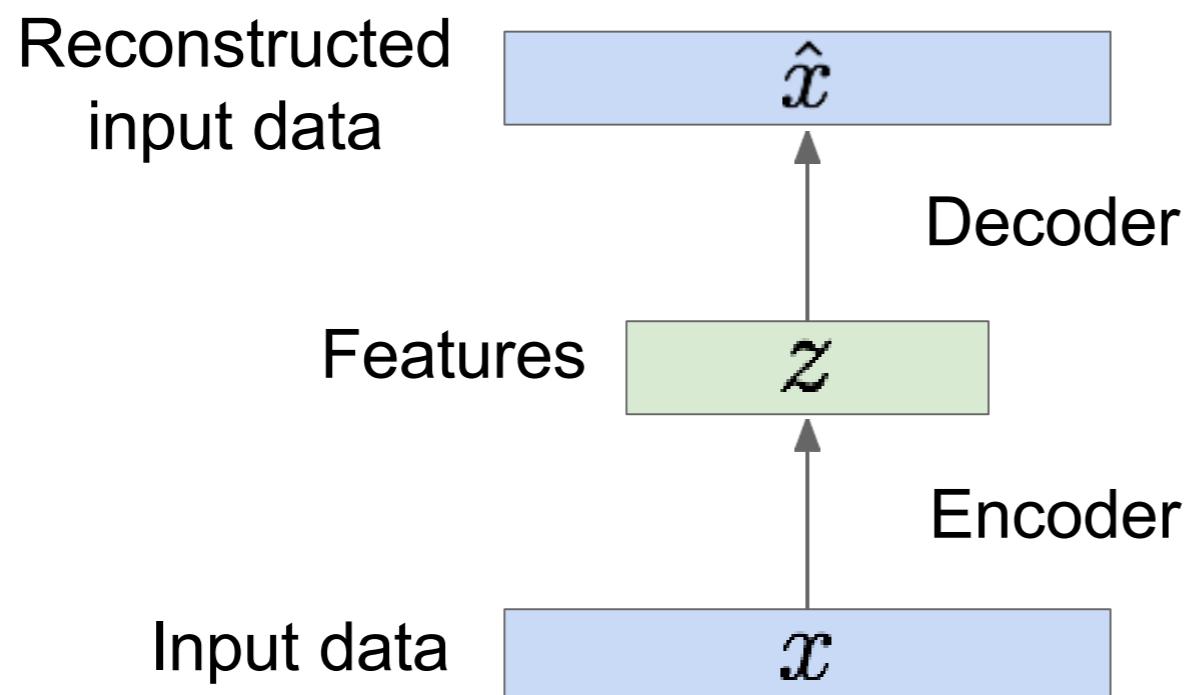


# Some background first: Autoencoders

Encoder can be used to initialize a **supervised** model



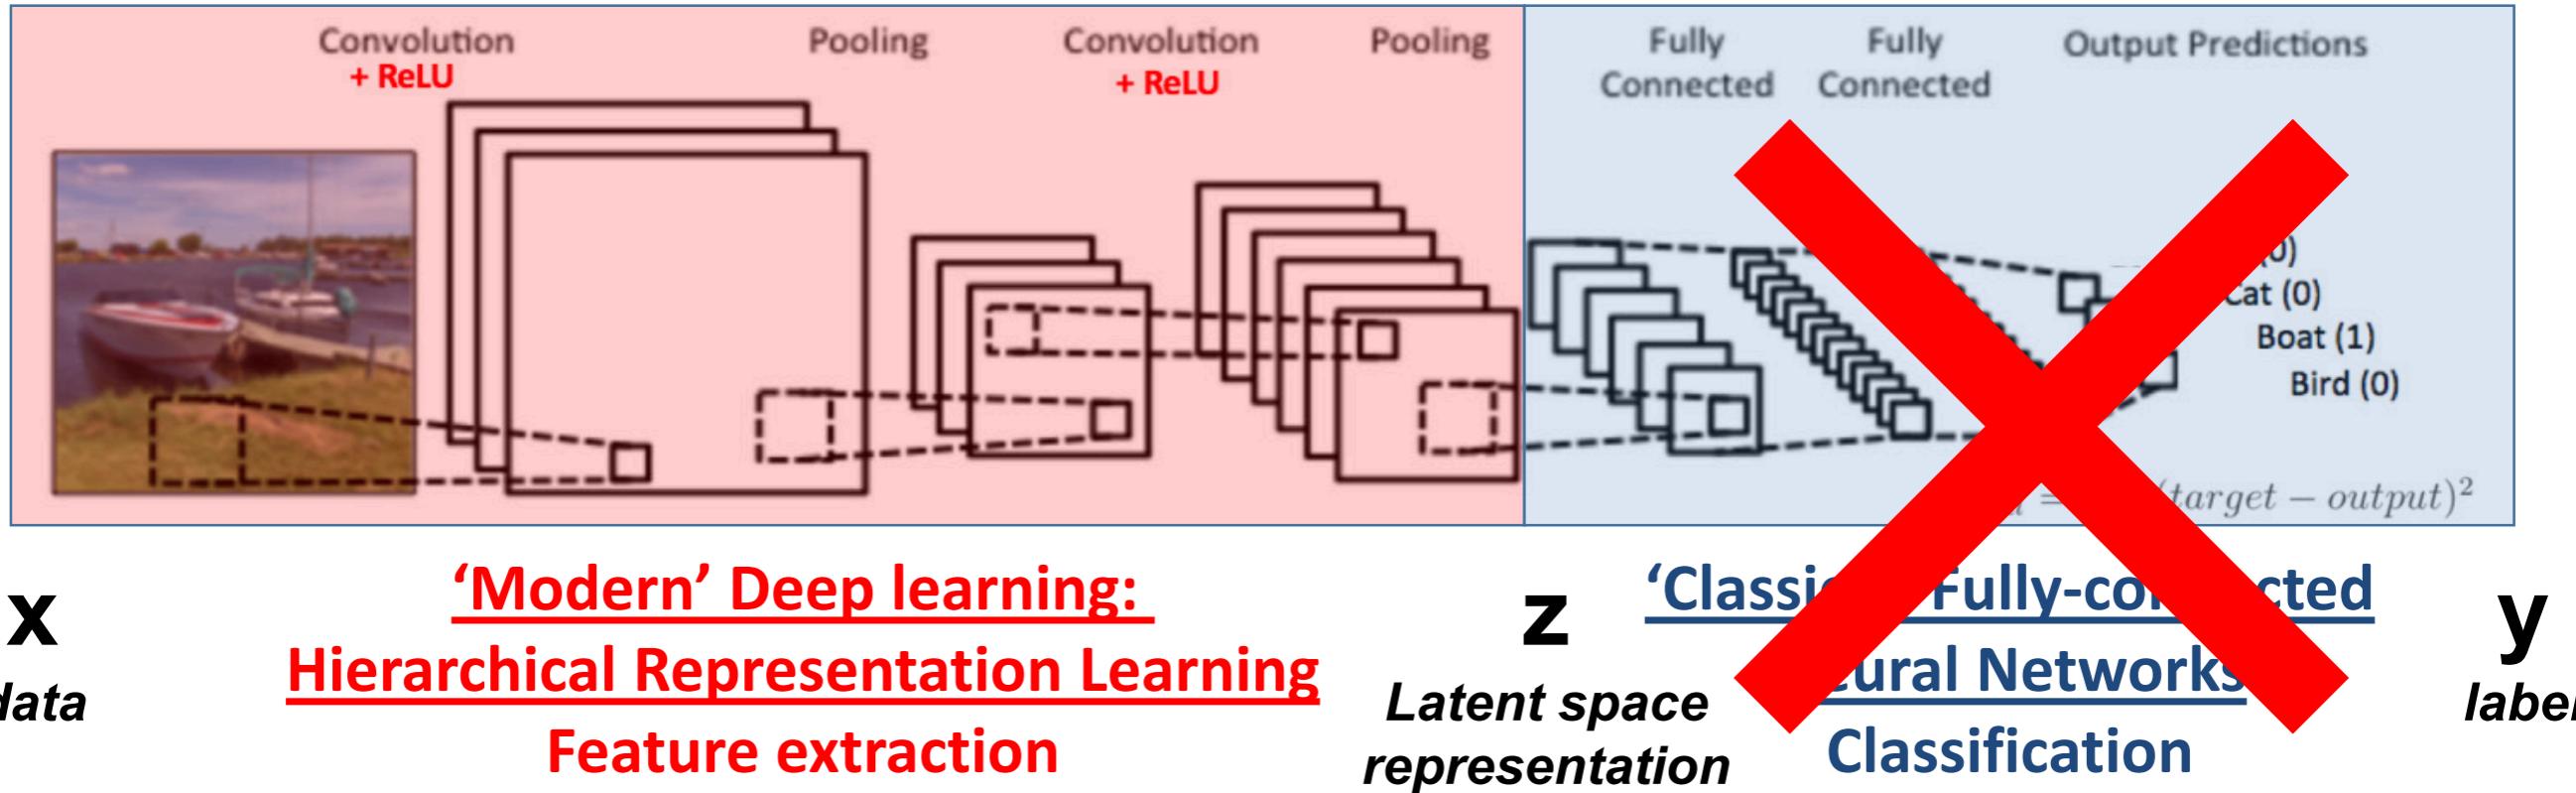
# Some background first: Autoencoders



Autoencoders can reconstruct data, and can learn features to initialize a supervised model

Features capture factors of variation in training data. Can we generate new images from an autoencoder?

# Representation learning without annotations?



**Many ideas are possible (and yours could be even better!):**

1. Predict the future: RNNs, Video
2. Compression: Autoencoder (predict self, through clamp), representation  $z$
3. Pretext tasks: predict self, before/after, missing patch, correct rotation, colorization, up-sampling, multimodal
4. Capture parameter distribution (variance): Variational Auto-Encoders
5. Make latent space parameters  $z$  meaningful, orthogonal, explicit, tuneable
6. Train using a second network: GANs - Improve quality of output images
7. The Sky is the Limit

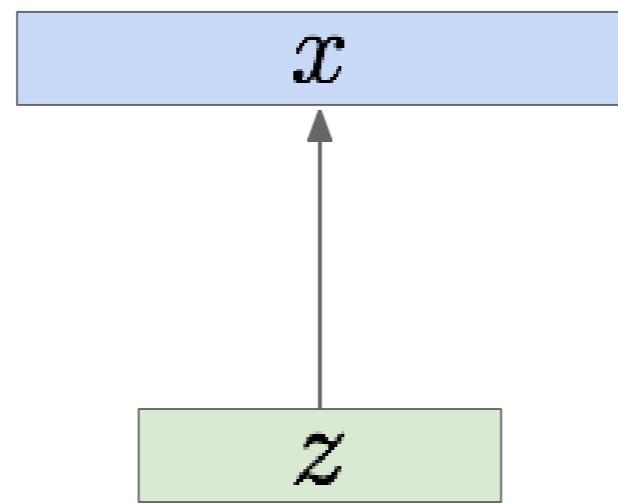
# Variational AutoEncoders (VAEs)

# Variational Autoencoders

Probabilistic spin on autoencoders - will let us sample from the model to generate data!

Assume training data  $\{x^{(i)}\}_{i=1}^N$  is generated from underlying unobserved (latent) representation  $\mathbf{z}$

Sample from  
true conditional  
 $p_{\theta^*}(x \mid z^{(i)})$



Sample from  
true prior  
 $p_{\theta^*}(z)$

**Intuition** (remember from autoencoders!):  
 $\mathbf{x}$  is an image,  $\mathbf{z}$  is latent factors used to  
generate  $\mathbf{x}$ : attributes, orientation, etc.

Kingma and Welling, "Auto-Encoding Variational Bayes", ICLR 2014

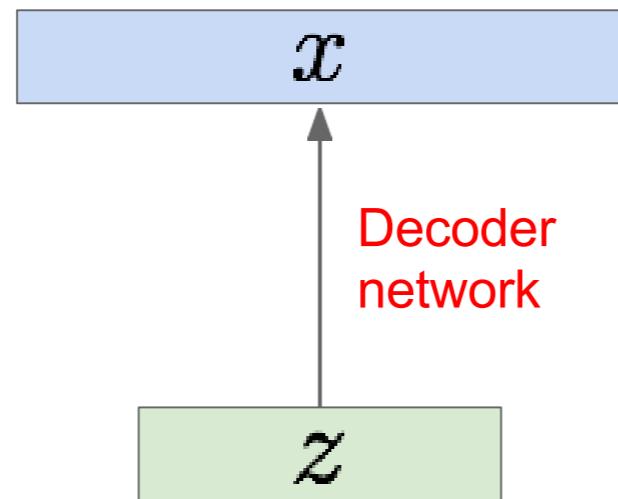
# Variational Autoencoders

Sample from  
true conditional

$$p_{\theta^*}(x \mid z^{(i)})$$

Sample from  
true prior

$$p_{\theta^*}(z)$$



We want to estimate the true parameters  $\theta^*$  of this generative model.

How should we represent this model?

Choose prior  $p(z)$  to be simple, e.g.  
Gaussian.

Conditional  $p(x|z)$  is complex (generates  
image) => represent with neural network

Kingma and Welling, “Auto-Encoding Variational Bayes”, ICLR 2014

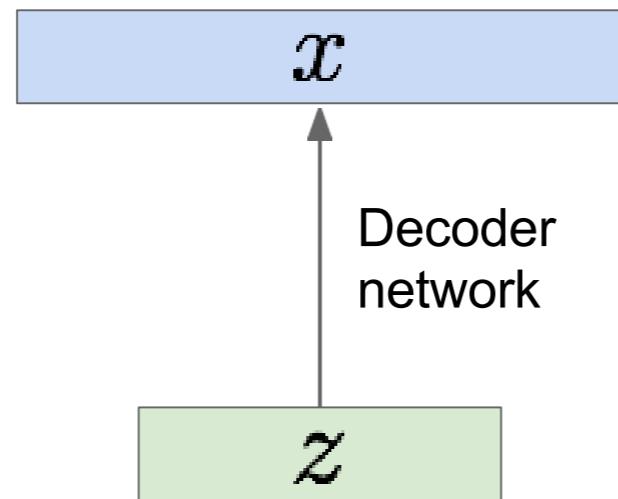
# Variational Autoencoders

Sample from  
true conditional

$$p_{\theta^*}(x \mid z^{(i)})$$

Sample from  
true prior

$$p_{\theta^*}(z)$$



We want to estimate the true parameters  $\theta^*$  of this generative model.

How to train the model?

Remember strategy for training generative models from FVBMs. Learn model parameters to maximize likelihood of training data

$$p_{\theta}(x) = \int p_{\theta}(z)p_{\theta}(x|z)dz$$

Now with latent  $z$

Kingma and Welling, “Auto-Encoding Variational Bayes”, ICLR 2014

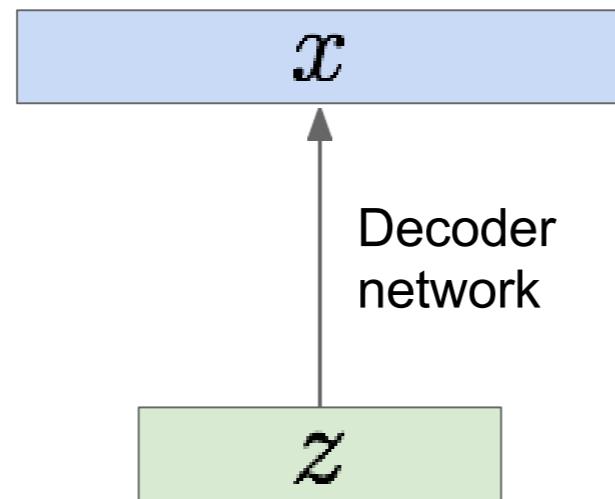
# Variational Autoencoders

Sample from  
true conditional

$$p_{\theta^*}(x \mid z^{(i)})$$

Sample from  
true prior

$$p_{\theta^*}(z)$$



We want to estimate the true parameters  $\theta^*$  of this generative model.

How to train the model?

Remember strategy for training generative models from FVBMs. Learn model parameters to maximize likelihood of training data

$$p_{\theta}(x) = \int p_{\theta}(z)p_{\theta}(x|z)dz$$

Q: What is the problem with this?

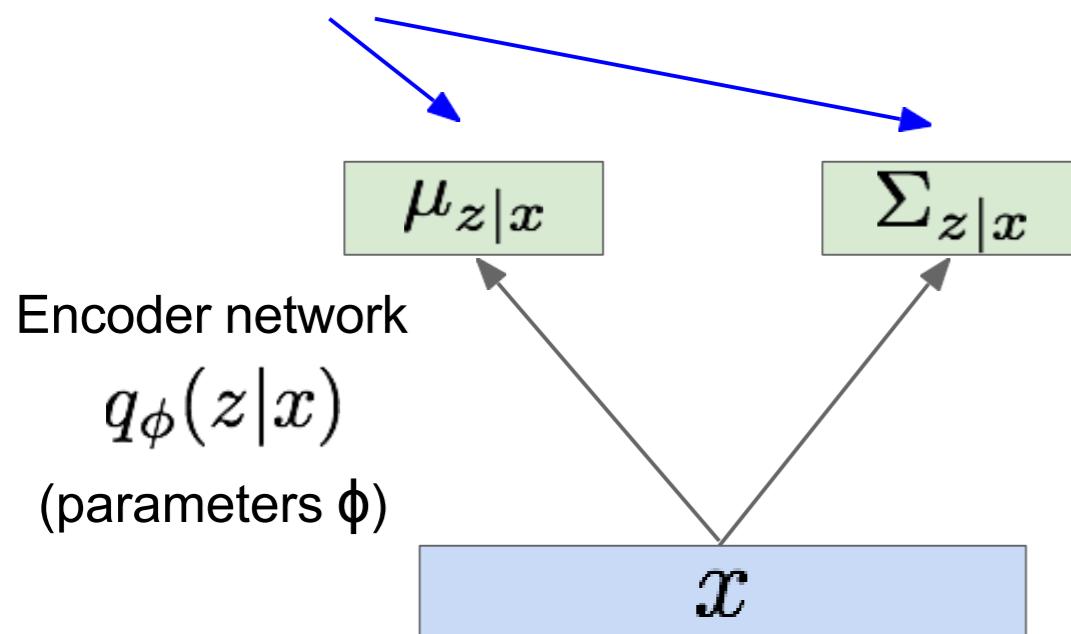
Intractable!

Kingma and Welling, “Auto-Encoding Variational Bayes”, ICLR 2014

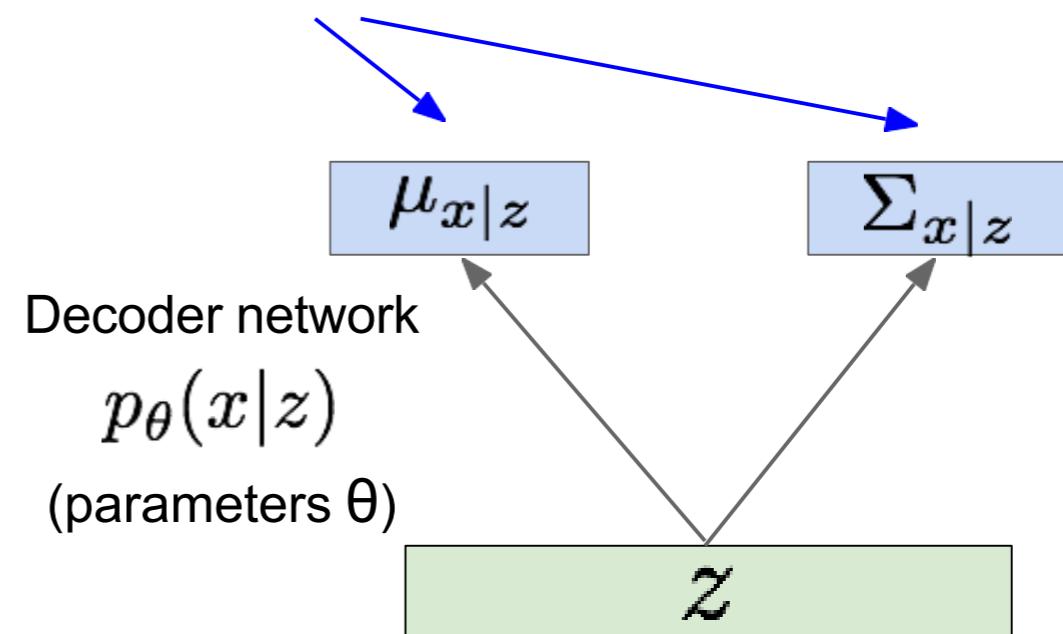
# Variational Autoencoders

Since we're modeling probabilistic generation of data, encoder and decoder networks are probabilistic Mean and

(diagonal) covariance of  $\mathbf{z} | \mathbf{x}$



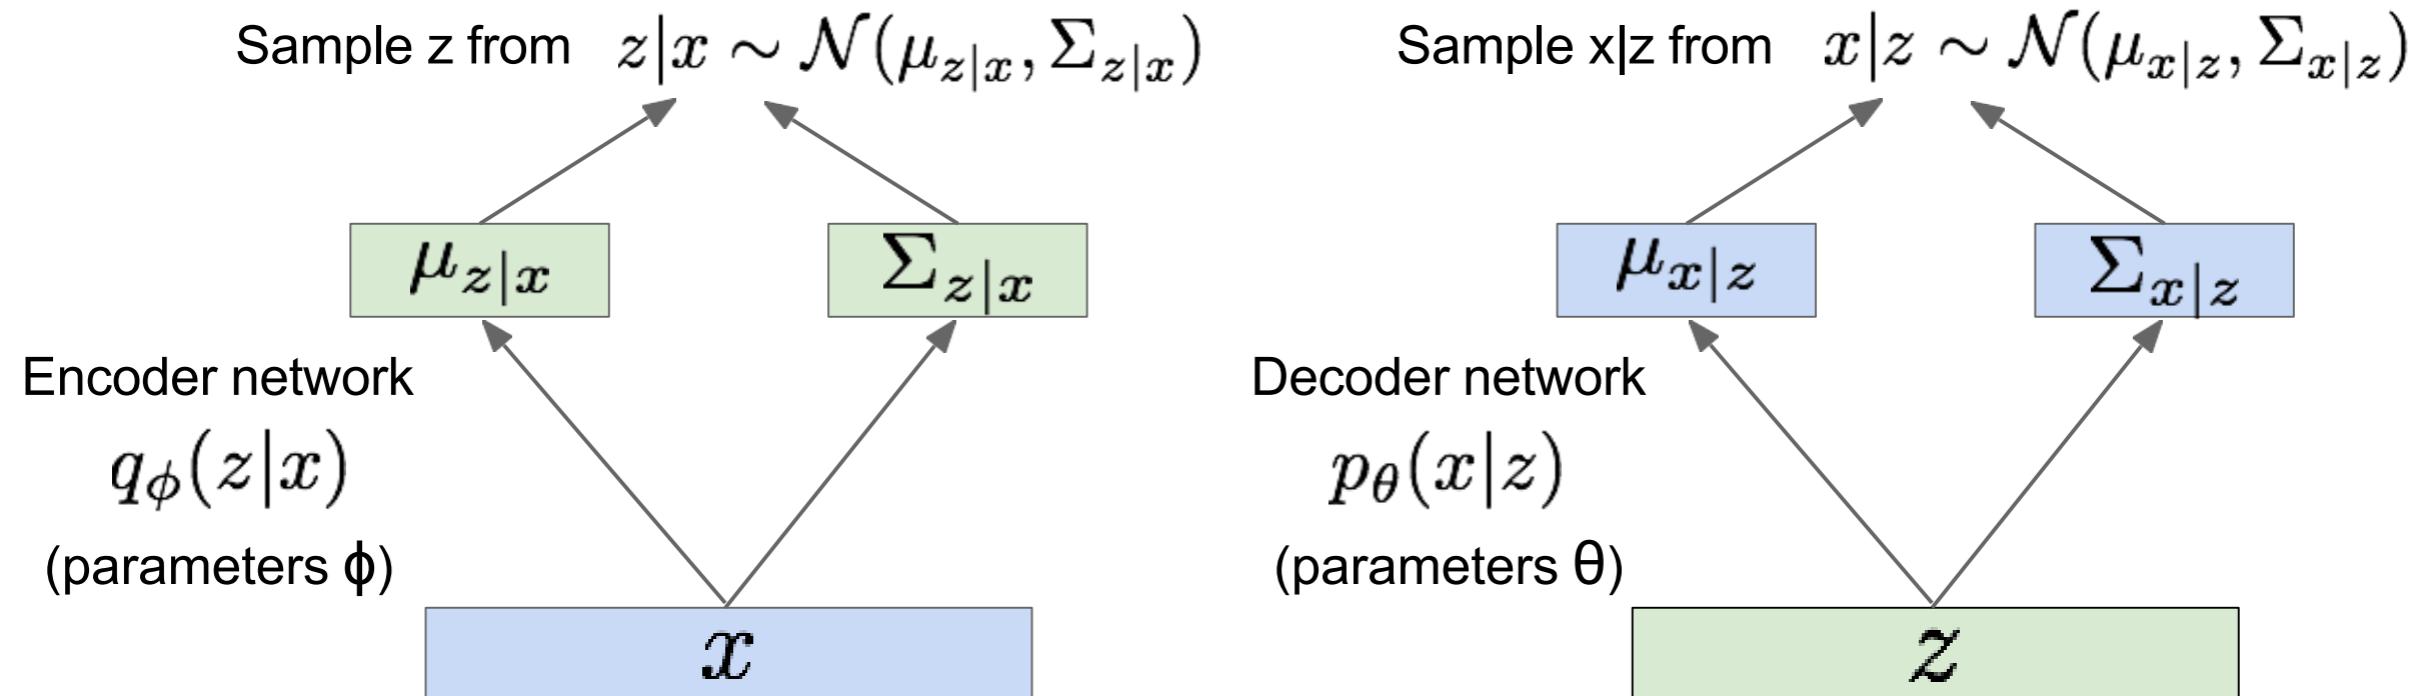
Mean and (diagonal) covariance of  $\mathbf{x} | \mathbf{z}$



Kingma and Welling, "Auto-Encoding Variational Bayes", ICLR 2014

# Variational Autoencoders

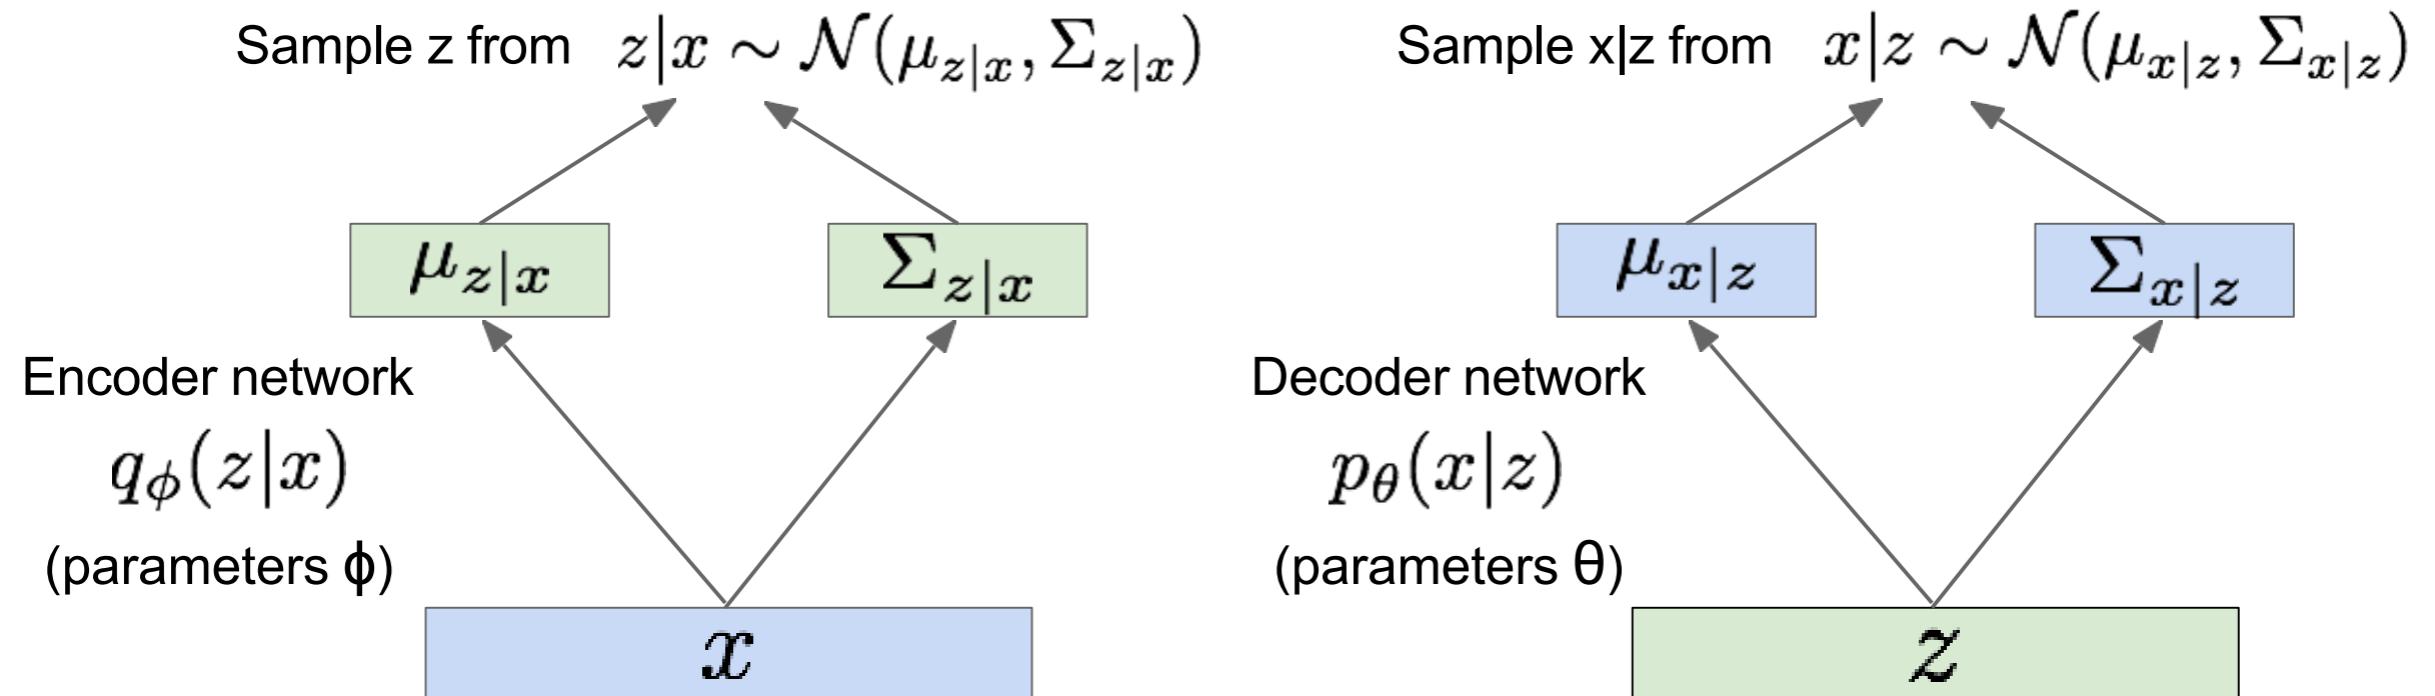
Since we're modeling probabilistic generation of data, encoder and decoder networks are probabilistic



Kingma and Welling, “Auto-Encoding Variational Bayes”, ICLR 2014

# Variational Autoencoders

Since we're modeling probabilistic generation of data, encoder and decoder networks are probabilistic



Encoder and decoder networks also called  
“recognition”/“inference” and “generation” networks

Kingma and Welling, “Auto-Encoding Variational Bayes”, ICLR 2014

# Variational Autoencoders

Now equipped with our encoder and decoder networks, let's work out the (log) data likelihood:

$$\begin{aligned}\log p_\theta(x^{(i)}) &= \mathbf{E}_{z \sim q_\phi(z | x^{(i)})} [\log p_\theta(x^{(i)})] \quad (p_\theta(x^{(i)}) \text{ Does not depend on } z) \\ &= \mathbf{E}_z \left[ \log \frac{p_\theta(x^{(i)} | z)p_\theta(z)}{p_\theta(z | x^{(i)})} \right] \quad (\text{Bayes' Rule}) \\ &= \mathbf{E}_z \left[ \log \frac{p_\theta(x^{(i)} | z)p_\theta(z)}{p_\theta(z | x^{(i)})} \frac{q_\phi(z | x^{(i)})}{q_\phi(z | x^{(i)})} \right] \quad (\text{Multiply by constant}) \\ &= \mathbf{E}_z [\log p_\theta(x^{(i)} | z)] - \mathbf{E}_z \left[ \log \frac{q_\phi(z | x^{(i)})}{p_\theta(z)} \right] + \mathbf{E}_z \left[ \log \frac{q_\phi(z | x^{(i)})}{p_\theta(z | x^{(i)})} \right] \quad (\text{Logarithms}) \\ &= \mathbf{E}_z [\log p_\theta(x^{(i)} | z)] - D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z)) + D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z | x^{(i)}))\end{aligned}$$

The expectation wrt.  $z$  (using encoder network) let us write nice KL terms

# Variational Autoencoders

Now equipped with our encoder and decoder networks, let's work out the (log) data likelihood:

$$\log p_\theta(x^{(i)}) = \mathbf{E}_{z \sim q_\phi(z|x^{(i)})} [\log p_\theta(x^{(i)})] \quad (p_\theta(x^{(i)}) \text{ Does not depend on } z)$$

We want to  
maximize the  
data  
likelihood

$$\uparrow$$
$$= \mathbf{E}_z \left[ \log \frac{p_\theta(x^{(i)} | z)p_\theta(z)}{p_\theta(z | x^{(i)})} \right] \quad (\text{Bayes' Rule})$$

$$= \mathbf{E}_z \left[ \log \frac{p_\theta(x^{(i)} | z)p_\theta(z)}{p_\theta(z | x^{(i)})} \frac{q_\phi(z | x^{(i)})}{q_\phi(z | x^{(i)})} \right] \quad (\text{Multiply by constant})$$

$$= \mathbf{E}_z [\log p_\theta(x^{(i)} | z)] - \mathbf{E}_z \left[ \log \frac{q_\phi(z | x^{(i)})}{p_\theta(z)} \right] + \mathbf{E}_z \left[ \log \frac{q_\phi(z | x^{(i)})}{p_\theta(z | x^{(i)})} \right] \quad (\text{Logarithms})$$

$$= \mathbf{E}_z [\log p_\theta(x^{(i)} | z)] - D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z)) + D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z | x^{(i)}))$$

Decoder network gives  $p_\theta(x|z)$ , can  
compute estimate of this term through  
sampling. (Sampling differentiable  
through reparam. trick, see paper.)

This KL term (between  
Gaussians for encoder and z  
prior) has nice closed-form  
solution!

$p_\theta(z|x)$  intractable (saw  
earlier), can't compute this KL  
term :( But we know KL  
divergence always  $\geq 0$ .

# Variational Autoencoders

Now equipped with our encoder and decoder networks, let's work out the (log) data likelihood:

We want to  
maximize the  
data  
likelihood

$$\begin{aligned} \log p_\theta(x^{(i)}) &= \mathbf{E}_{z \sim q_\phi(z|x^{(i)})} [\log p_\theta(x^{(i)})] \quad (p_\theta(x^{(i)}) \text{ Does not depend on } z) \\ &= \mathbf{E}_z \left[ \log \frac{p_\theta(x^{(i)} | z)p_\theta(z)}{p_\theta(z | x^{(i)})} \right] \quad (\text{Bayes' Rule}) \\ &= \mathbf{E}_z \left[ \log \frac{p_\theta(x^{(i)} | z)p_\theta(z)}{p_\theta(z | x^{(i)})} \frac{q_\phi(z | x^{(i)})}{q_\phi(z | x^{(i)})} \right] \quad (\text{Multiply by constant}) \\ &= \mathbf{E}_z [\log p_\theta(x^{(i)} | z)] - \mathbf{E}_z \left[ \log \frac{q_\phi(z | x^{(i)})}{p_\theta(z)} \right] + \mathbf{E}_z \left[ \log \frac{q_\phi(z | x^{(i)})}{p_\theta(z | x^{(i)})} \right] \quad (\text{Logarithms}) \\ &= \underbrace{\mathbf{E}_z [\log p_\theta(x^{(i)} | z)] - D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z))}_{\mathcal{L}(x^{(i)}, \theta, \phi)} + \underbrace{D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z | x^{(i)}))}_{\geq 0} \end{aligned}$$

**Tractable lower bound** which we can take  
gradient of and optimize! ( $p_\theta(x|z)$  differentiable,  
KL term differentiable)

# Variational Autoencoders

Now equipped with our encoder and decoder networks, let's work out the (log) data likelihood:

$$\log p_\theta(x^{(i)}) = \mathbf{E}_{z \sim q_\phi(z | x^{(i)})} [\log p_\theta(x^{(i)})] \quad (p_\theta(x^{(i)}) \text{ Does not depend on } z)$$

$$= \mathbf{E}_z \left[ \log \frac{p_\theta(x^{(i)} | z)p_\theta(z)}{p_\theta(z | x^{(i)})} \right] \quad (\text{Bayes' Rule})$$

**Reconstruct  
the input data**

$$= \mathbf{E}_z \left[ \log \frac{p_\theta(x^{(i)} | z)p_\theta(z)}{p_\theta(z | x^{(i)})} \frac{q_\phi(z | x^{(i)})}{q_\phi(z | x^{(i)})} \right] \quad (\text{Multiply by constant})$$

Make approximate posterior distribution close to prior

$$= \mathbf{E}_z [\log p_\theta(x^{(i)} | z)] - \mathbf{E}_z \left[ \log \frac{q_\phi(z | x^{(i)})}{p_\theta(z)} \right] + \mathbf{E}_z \left[ \log \frac{q_\phi(z | x^{(i)})}{p_\theta(z | x^{(i)})} \right] \quad (\text{Logarithms})$$

$$= \underbrace{\mathbf{E}_z [\log p_\theta(x^{(i)} | z)]}_{\mathcal{L}(x^{(i)}, \theta, \phi)} - D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z)) + \underbrace{D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z | x^{(i)}))}_{> 0}$$

$$\log p_\theta(x^{(i)}) \geq \mathcal{L}(x^{(i)}, \theta, \phi)$$

Variational lower bound ("ELBO")

$$\theta^*, \phi^* = \arg \max_{\theta, \phi} \sum_{i=1}^N \mathcal{L}(x^{(i)}, \theta, \phi)$$

Training: Maximize lower bound

# Variational Autoencoders

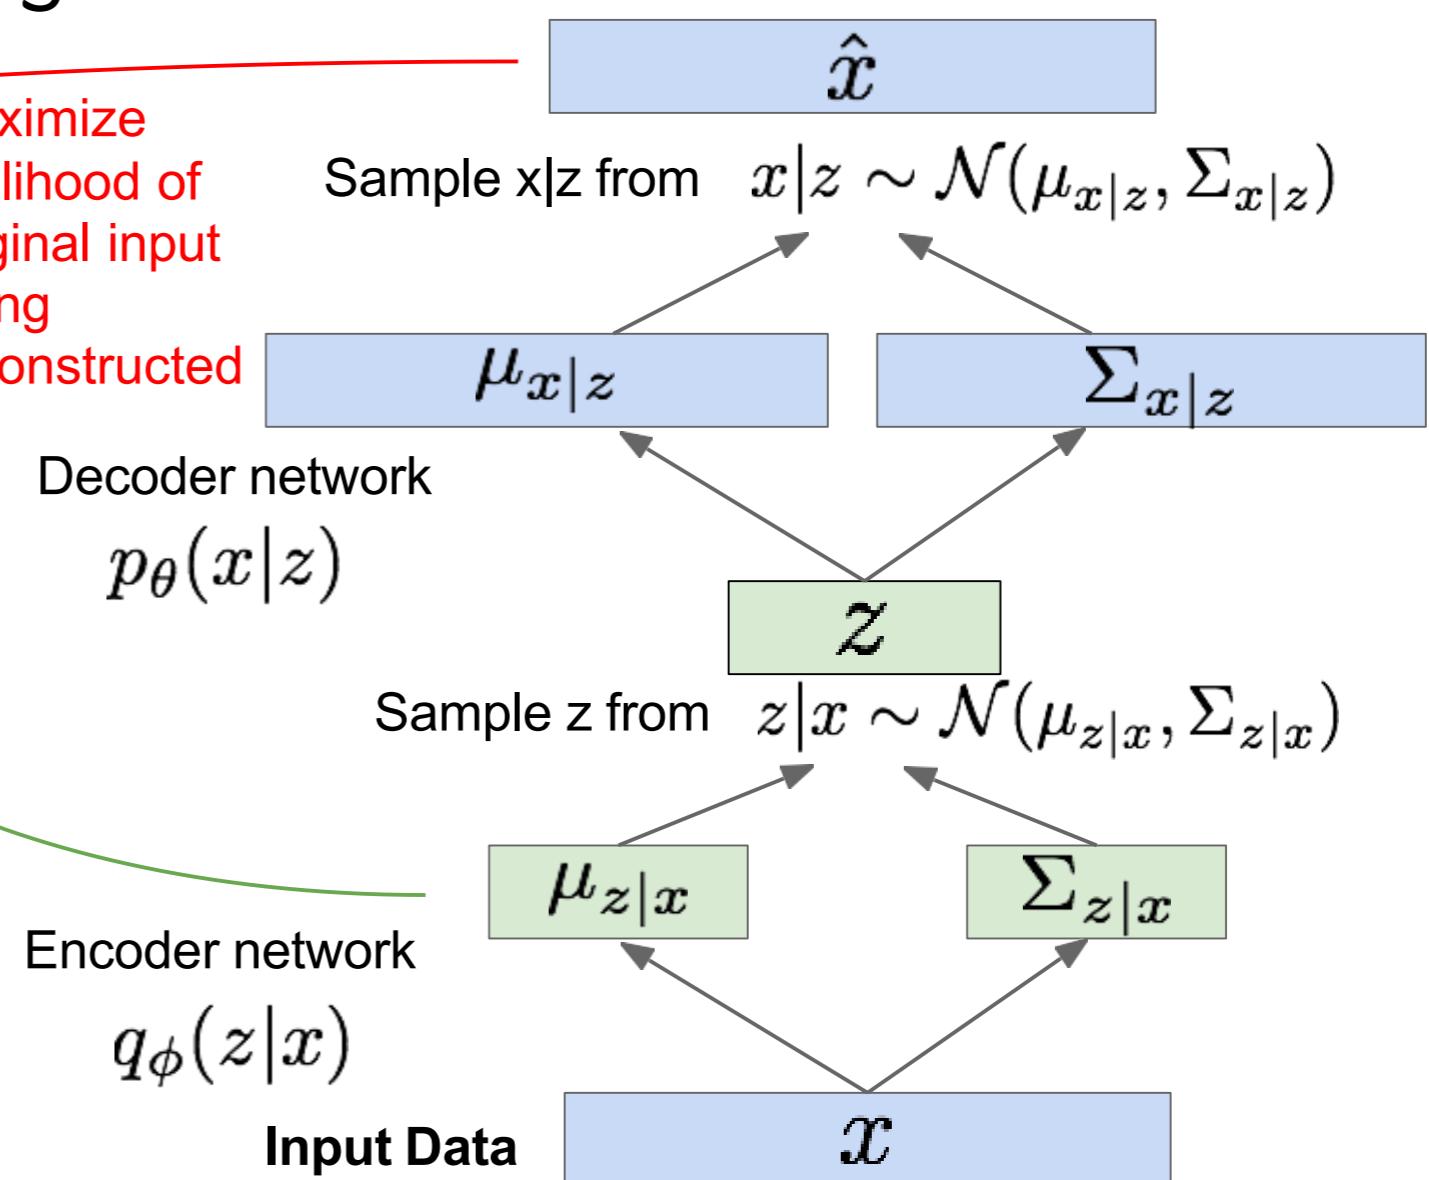
Putting it all together: maximizing the likelihood lower bound

$$\underbrace{\mathbb{E}_z \left[ \log p_\theta(x^{(i)} | z) \right] - D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z))}_{\mathcal{L}(x^{(i)}, \theta, \phi)}$$

Maximize likelihood of original input being reconstructed

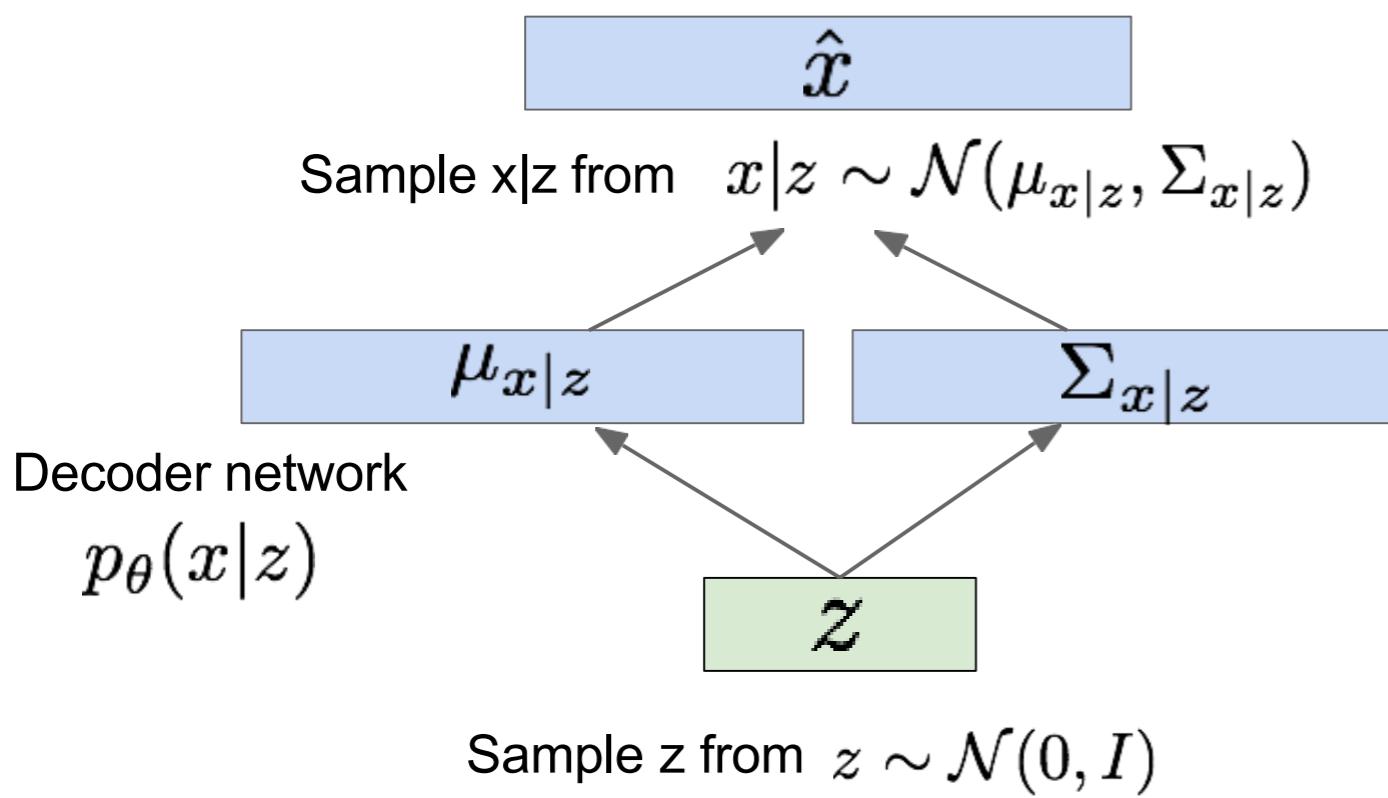
Make approximate posterior distribution close to prior

For every minibatch of input data: compute this forward pass, and then backprop!

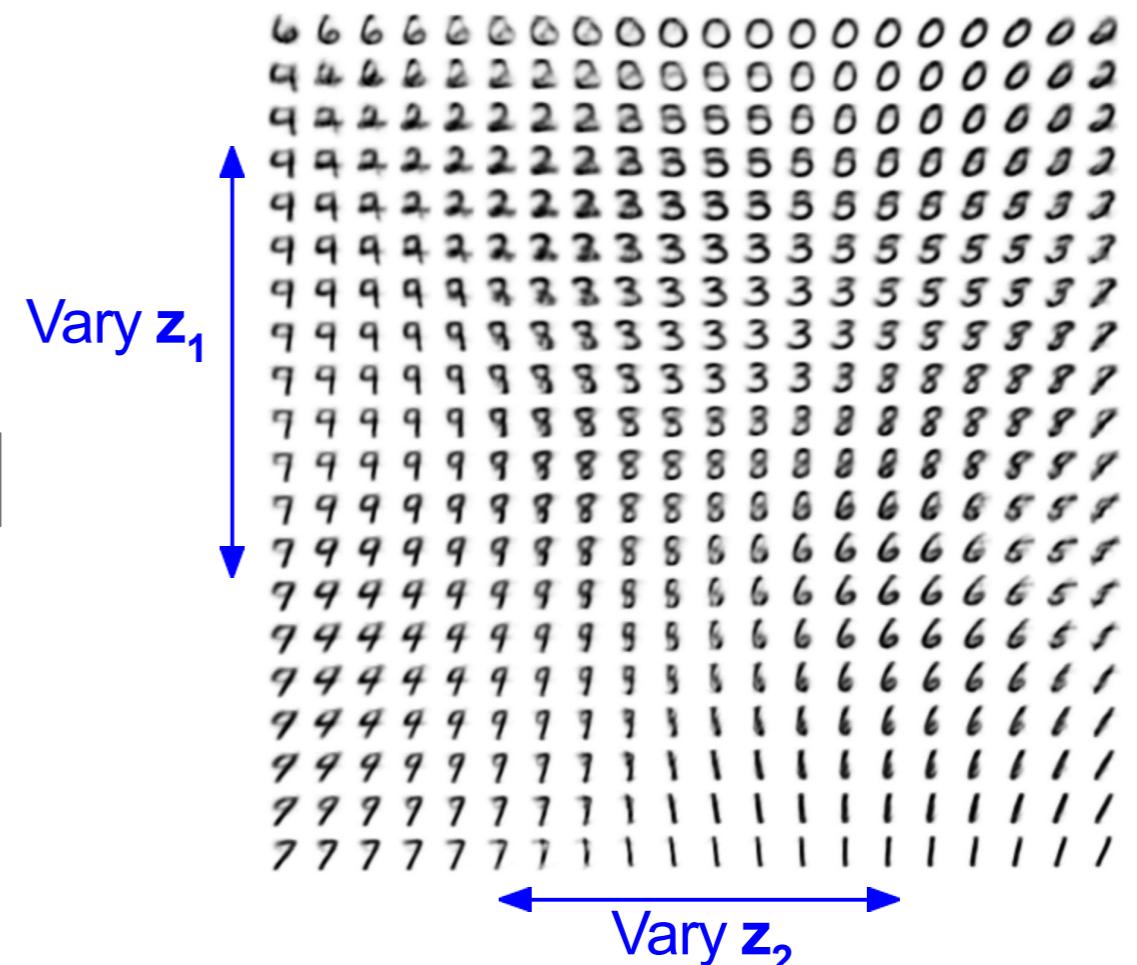


# Variational Autoencoders: Generating Data!

Use decoder network. Now sample z from prior!



Data manifold for 2-d  $z$



Kingma and Welling, "Auto-Encoding Variational Bayes", ICLR 2014

# Variational Autoencoders: Generating Data!

Diagonal prior on  $\mathbf{z}$   
=> independent  
latent variables

Different  
dimensions of  $\mathbf{z}$   
encode  
interpretable factors  
of variation

Also good feature representation that  
can be computed using  $q_\phi(z|x)$ !

Kingma and Welling, "Auto-Encoding Variational Bayes", ICLR 2014

Degree of smile  
Vary  $\mathbf{z}_1$

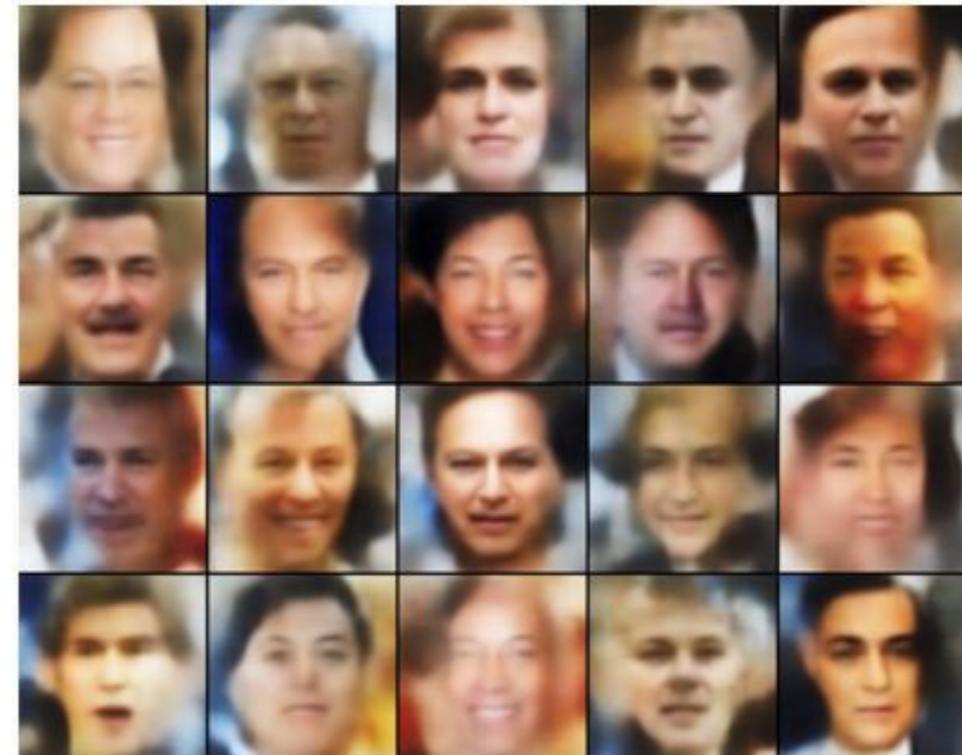


Head pose

# Variational Autoencoders: Generating Data!



32x32 CIFAR-10



Labeled Faces in the Wild

Figures copyright (L) Dirk Kingma et al. 2016; (R) Anders Larsen et al. 2017. Reproduced with permission.

# Variational Autoencoders

Probabilistic spin to traditional autoencoders => allows generating data

Defines an intractable density => derive and optimize a (variational) lower bound

## Pros:

- Principled approach to generative models
- Allows inference of  $q(z|x)$ , can be useful feature representation for other tasks

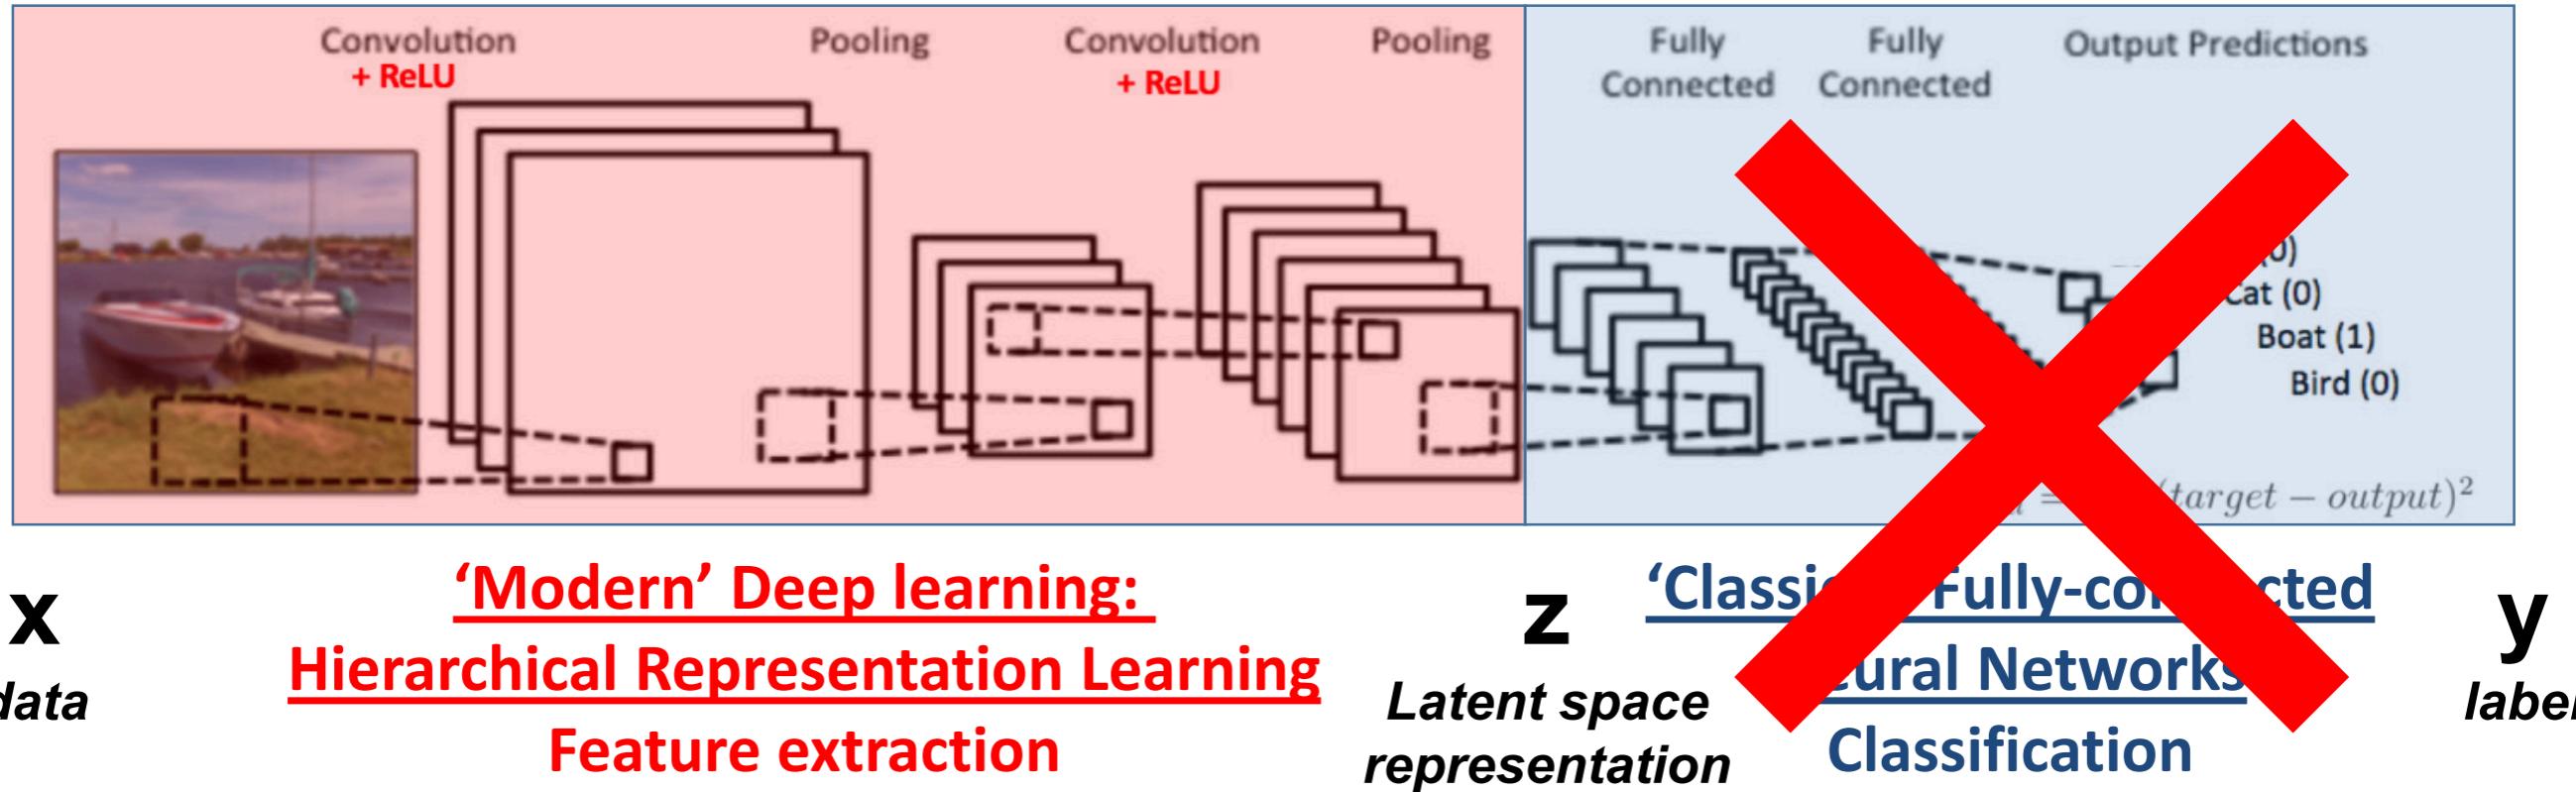
## Cons:

- Maximizes lower bound of likelihood: okay, but not as good evaluation as PixelRNN/PixelCNN
- Samples blurrier and lower quality compared to state-of-the-art (GANs)

## Active areas of research:

- More flexible approximations, e.g. richer approximate posterior instead of diagonal Gaussian, e.g., Gaussian Mixture Models (GMMs)
- Incorporating structure in latent variables, e.g., Categorical Distributions

# Representation learning without annotations?



**Many ideas are possible (and yours could be even better!):**

1. Predict the future: RNNs, Video
2. Compression: Autoencoder (predict self, through clamp), representation  $z$
3. Pretext tasks: predict self, before/after, missing patch, correct rotation, colorization, up-sampling, multimodal
4. Capture parameter distribution (variance): Variational Auto-Encoders
5. Make latent space parameters  $z$  meaningful, orthogonal, explicit, tuneable
6. Train using a second network: GANs - Improve quality of output images
7. The Sky is the Limit

# **GANs:**

# **Generative Adversarial Networks**

# Generative Adversarial Networks

Ian Goodfellow et al., “Generative Adversarial Nets”, NIPS 2014

Problem: Want to sample from complex, high-dimensional training distribution. No direct way to do this!

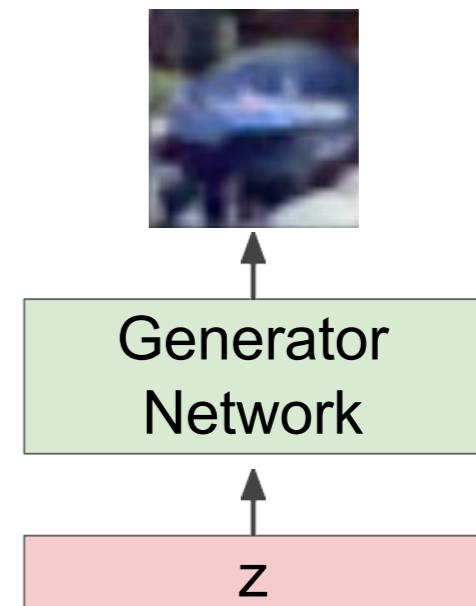
Solution: Sample from a simple distribution, e.g. random noise. Learn transformation to training distribution.

Q: What can we use to represent this complex transformation?

A: A neural network!

Output: Sample from training distribution

Input: Random noise

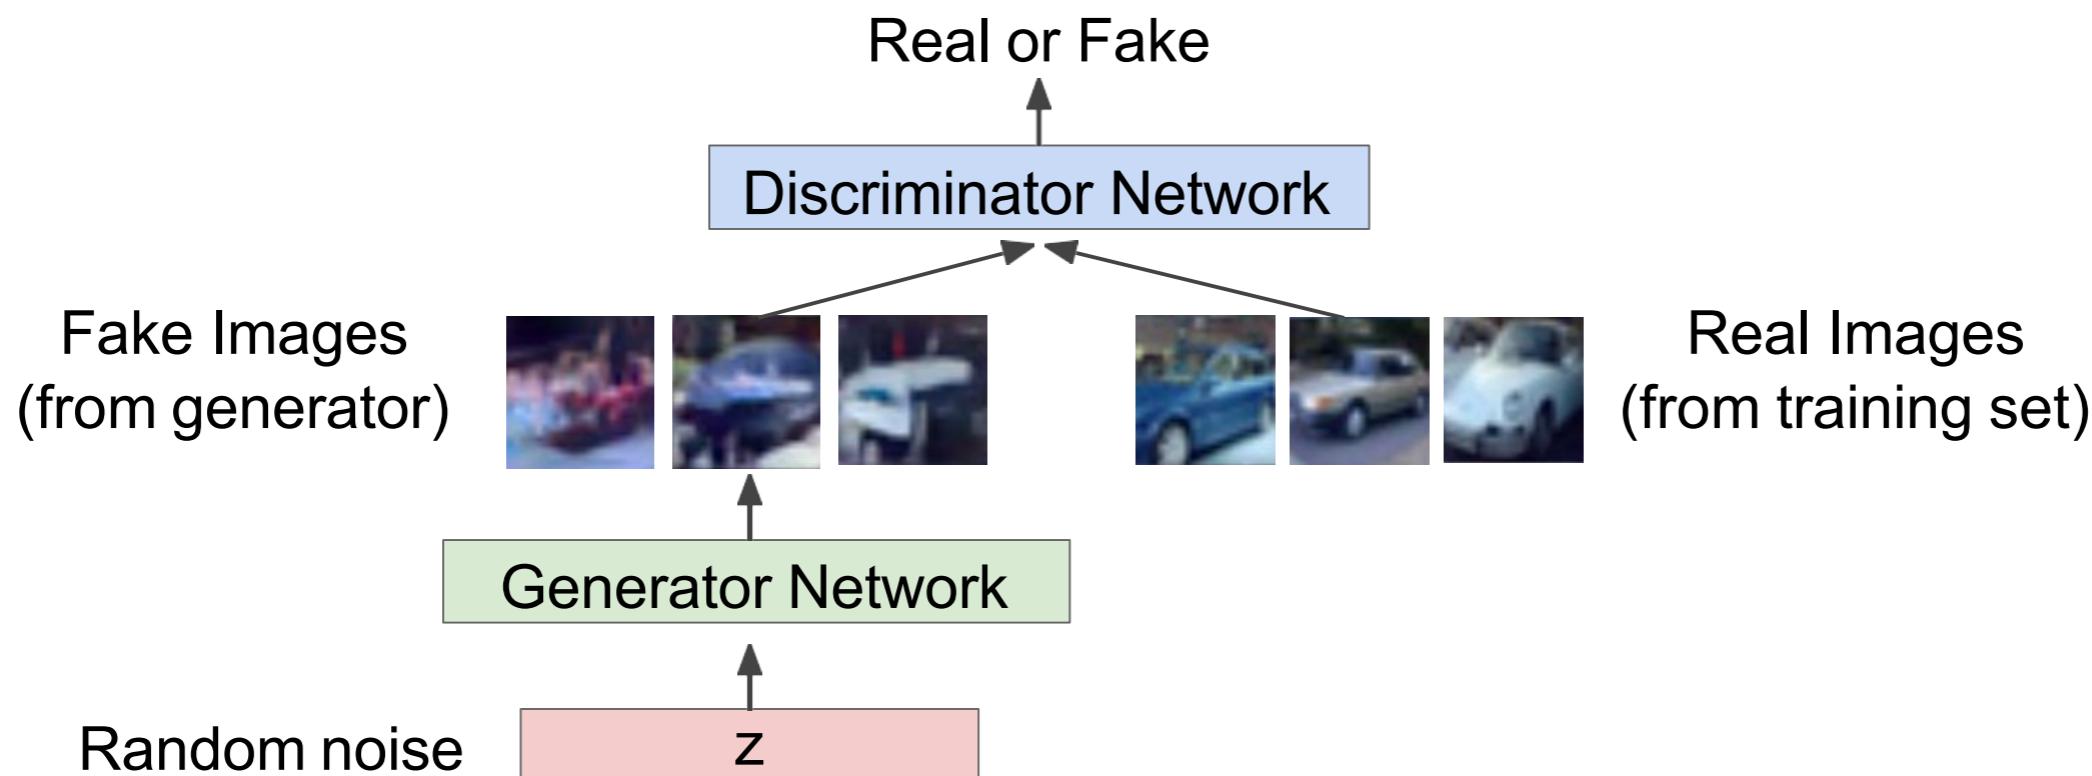


# Training GANs: Two-player game

Ian Goodfellow et al., “Generative Adversarial Nets”, NIPS 2014

**Generator network:** try to fool the discriminator by generating real-looking images

**Discriminator network:** try to distinguish between real and fake images



Fake and real images copyright Emily Denton et al. 2015. Reproduced with permission.

# Training GANs: Two-player game

Ian Goodfellow et al., “Generative Adversarial Nets”, NIPS 2014

**Generator network:** try to fool the discriminator by generating real-looking images

**Discriminator network:** try to distinguish between real and fake images

Train jointly in **minimax game**

Discriminator outputs likelihood in (0,1) of real image

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log \underbrace{D_{\theta_d}(x)}_{\text{Discriminator output for real data } x} + \mathbb{E}_{z \sim p(z)} \log(1 - \underbrace{D_{\theta_d}(G_{\theta_g}(z))}_{\text{Discriminator output for generated fake data } G(z)}) \right]$$

- Discriminator ( $\theta_d$ ) wants to **maximize objective** such that  $D(x)$  is close to 1 (real) and  $D(G(z))$  is close to 0 (fake)
- Generator ( $\theta_g$ ) wants to **minimize objective** such that  $D(G(z))$  is close to 1 (discriminator is fooled into thinking generated  $G(z)$  is real)

# Training GANs: Two-player game

Ian Goodfellow et al., “Generative Adversarial Nets”, NIPS 2014

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Alternate between:

1. **Gradient ascent** on discriminator

$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

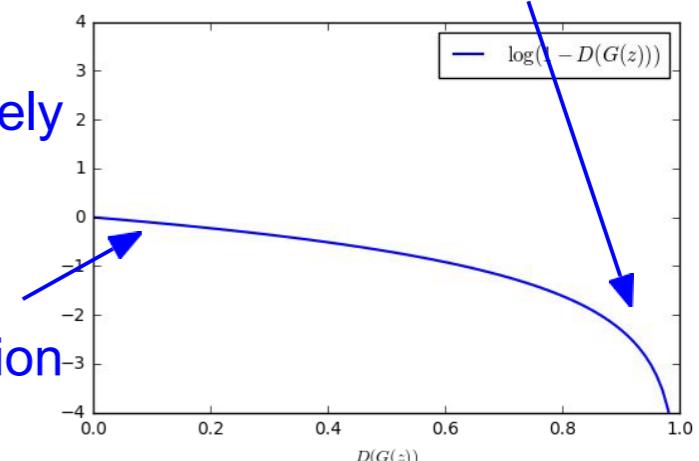
Gradient signal dominated by region where sample is already good

2. **Gradient descent** on generator

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

In practice, optimizing this generator objective does not work well!

When sample is likely fake, want to learn from it to improve generator. But gradient in this region is relatively flat!



# Training GANs: Two-player game

Ian Goodfellow et al., “Generative Adversarial Nets”, NIPS 2014

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Alternate between:

1. **Gradient ascent** on discriminator

$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

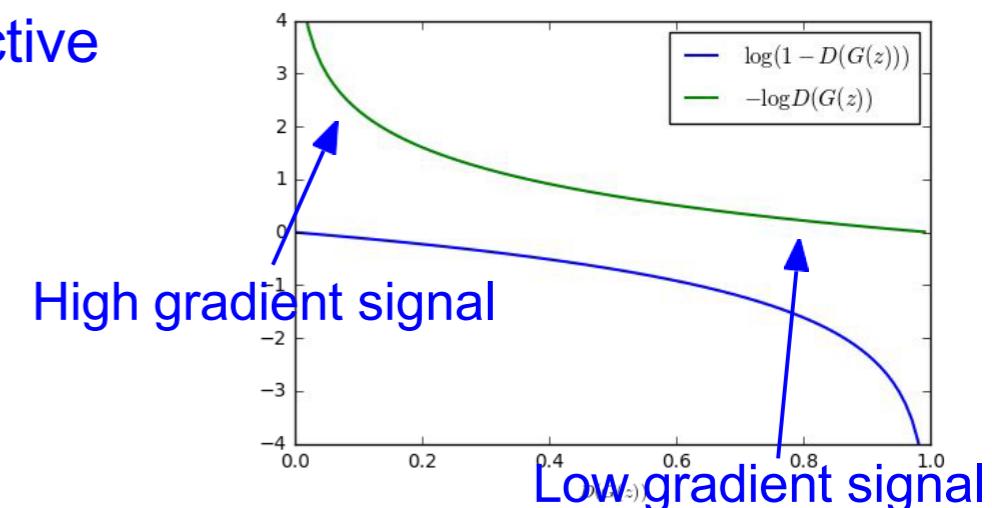
2. **Instead: Gradient ascent** on generator, **different objective**

$$\max_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))$$

Instead of minimizing likelihood of discriminator being correct, now maximize likelihood of discriminator being wrong.

Same objective of fooling discriminator, but now higher gradient signal for bad samples => works much better! Standard in practice.

Aside: Jointly training two networks is challenging, can be unstable. Choosing objectives with better loss landscapes helps training, is an active area of research.



# Training GANs: Two-player game

Ian Goodfellow et al., “Generative Adversarial Nets”, NIPS 2014

## Putting it together: GAN training algorithm

```
for number of training iterations do
    for k steps do
        • Sample minibatch of m noise samples { $\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}$ } from noise prior  $p_g(\mathbf{z})$ .
        • Sample minibatch of m examples { $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}$ } from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
        • Update the discriminator by ascending its stochastic gradient:
            
$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D_{\theta_d}(\mathbf{x}^{(i)}) + \log(1 - D_{\theta_d}(G_{\theta_g}(\mathbf{z}^{(i)}))) \right]$$

    end for
    • Sample minibatch of m noise samples { $\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}$ } from noise prior  $p_g(\mathbf{z})$ .
    • Update the generator by ascending its stochastic gradient (improved objective):
        
$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log(D_{\theta_d}(G_{\theta_g}(\mathbf{z}^{(i)})))$$

end for
```

Some find  $k=1$  more stable, others use  $k > 1$ , no best rule.

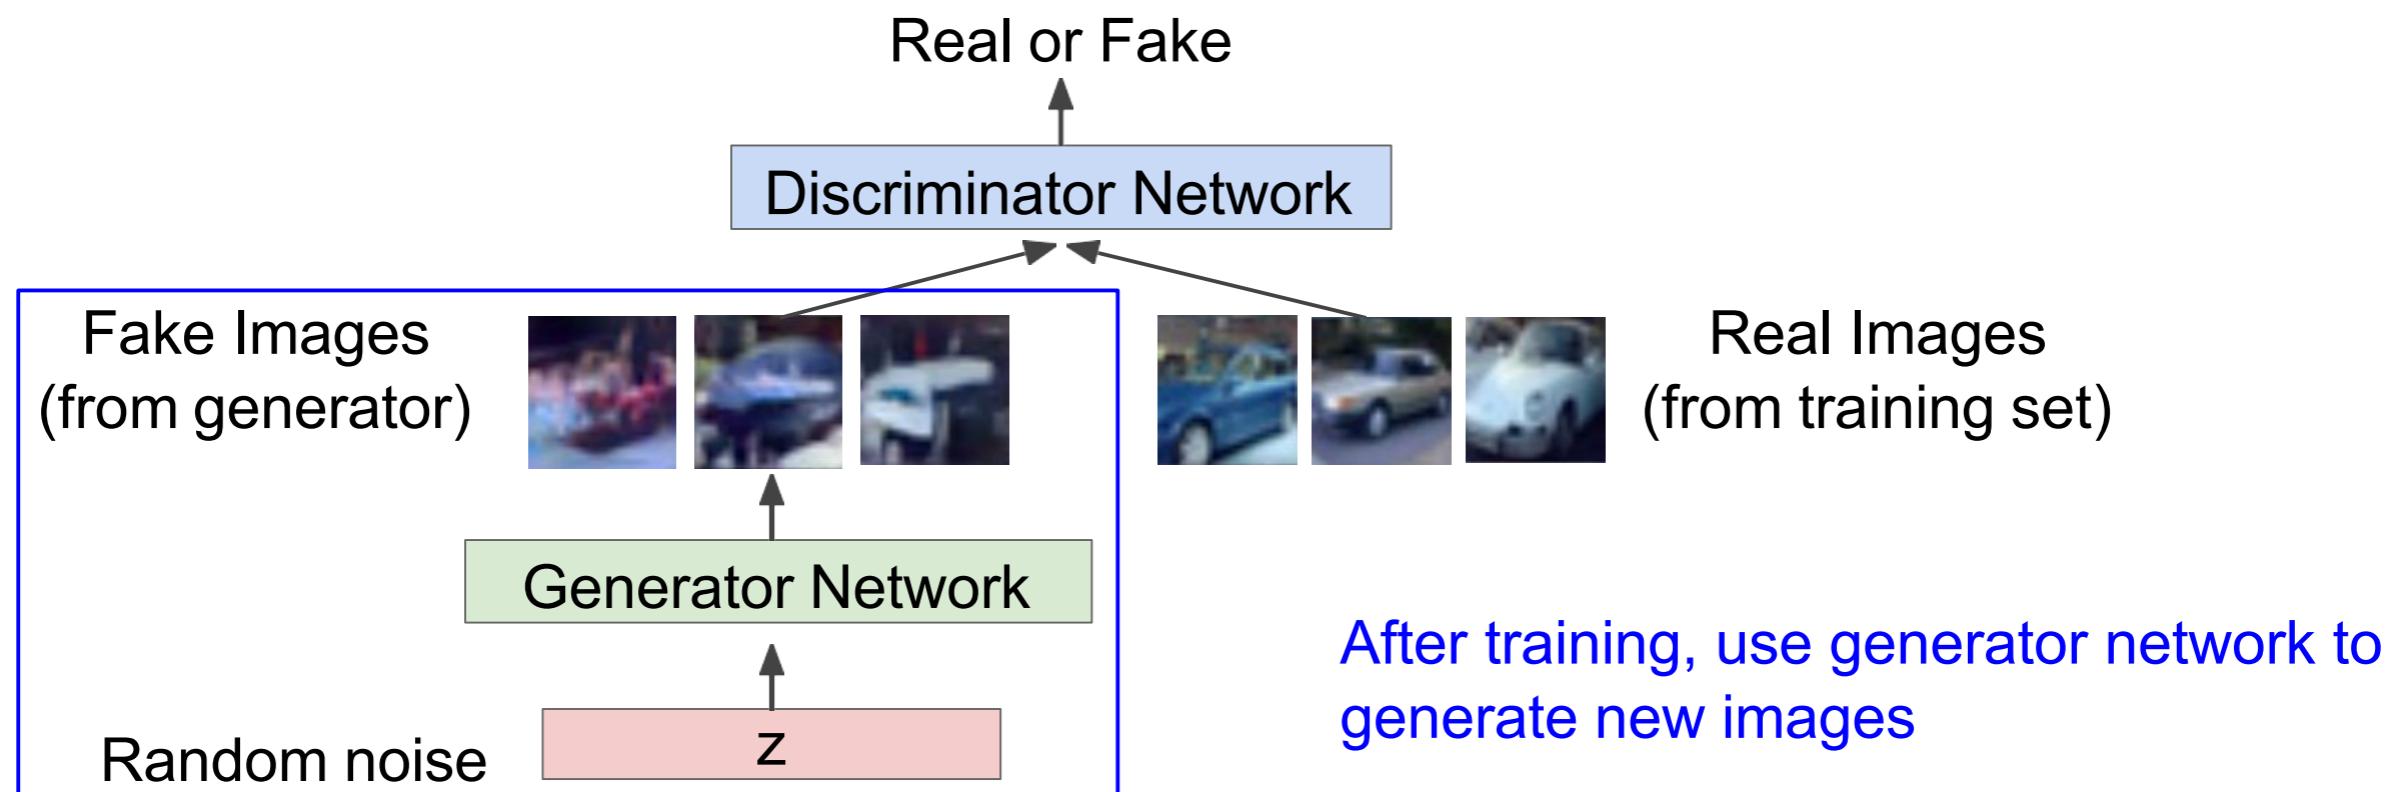
Recent work (e.g. Wasserstein GAN) alleviates this problem, better stability!

# Training GANs: Two-player game

Ian Goodfellow et al., “Generative Adversarial Nets”, NIPS 2014

**Generator network:** try to fool the discriminator by generating real-looking images

**Discriminator network:** try to distinguish between real and fake images



# Mean Squared Error Can Ignore Small but Task-Relevant Features

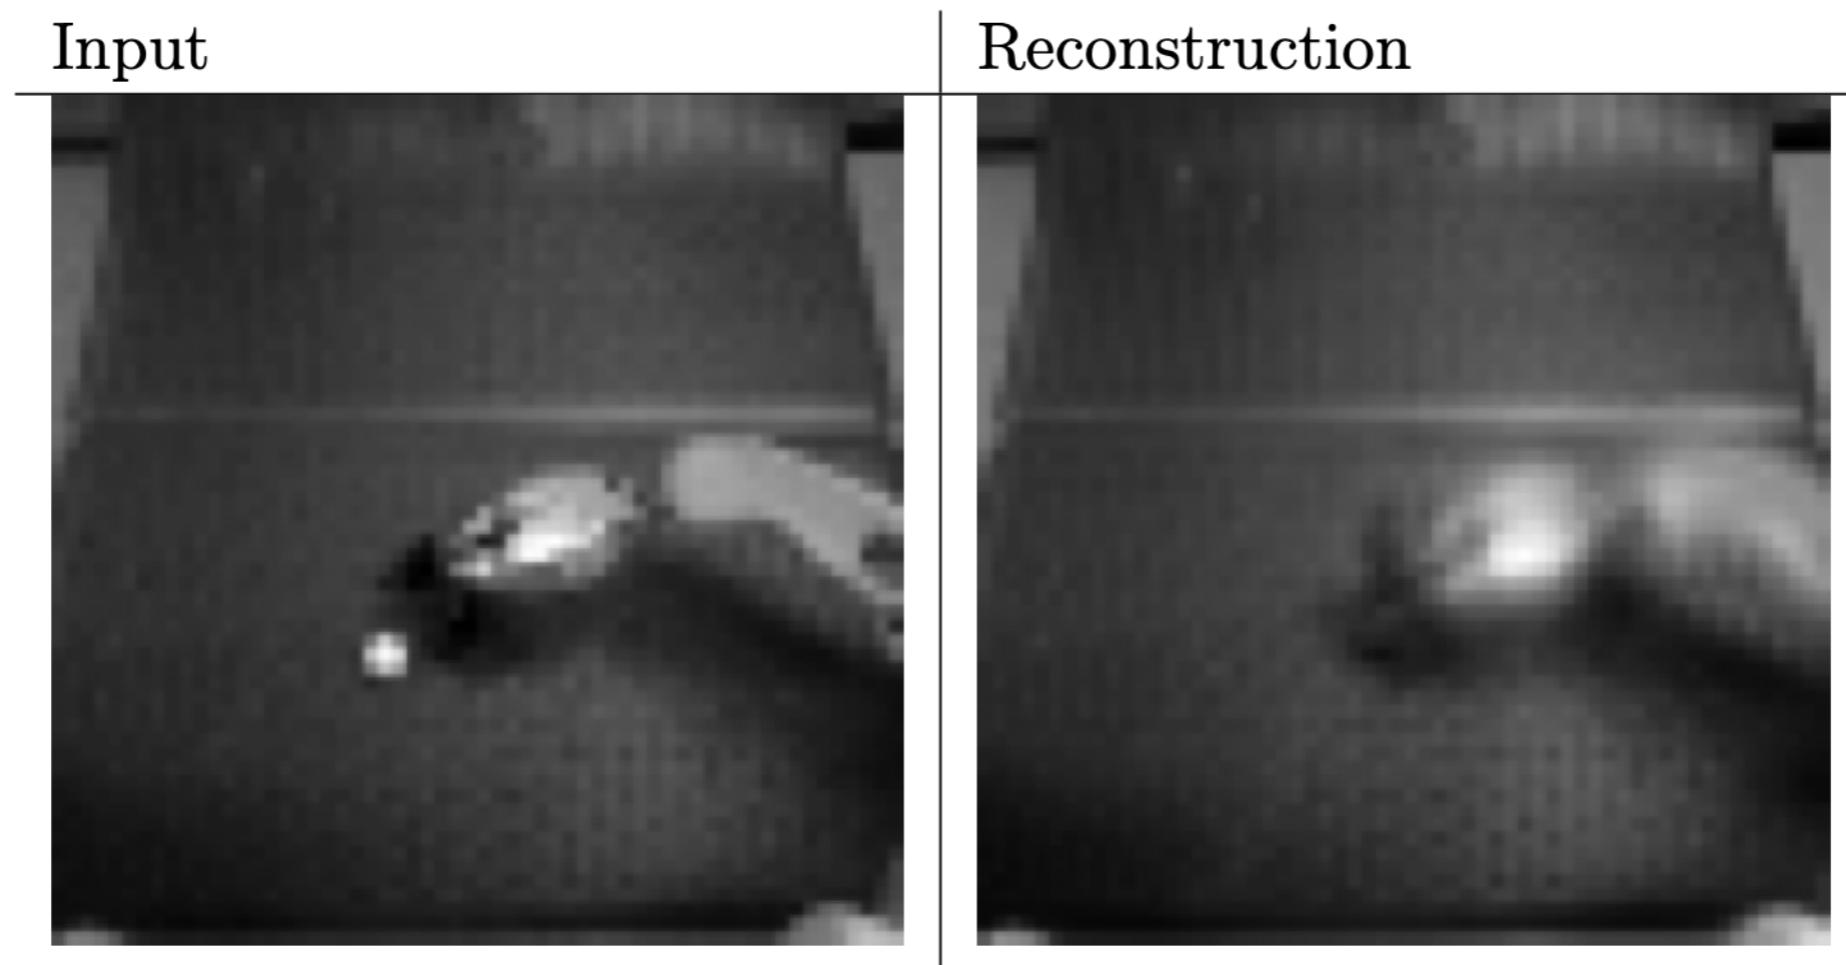


Figure 15.5

The ping pong ball vanishes because it is not large enough to significantly affect the mean squared error

# Adversarial Losses Preserve Any Features with Highly Structured Patterns

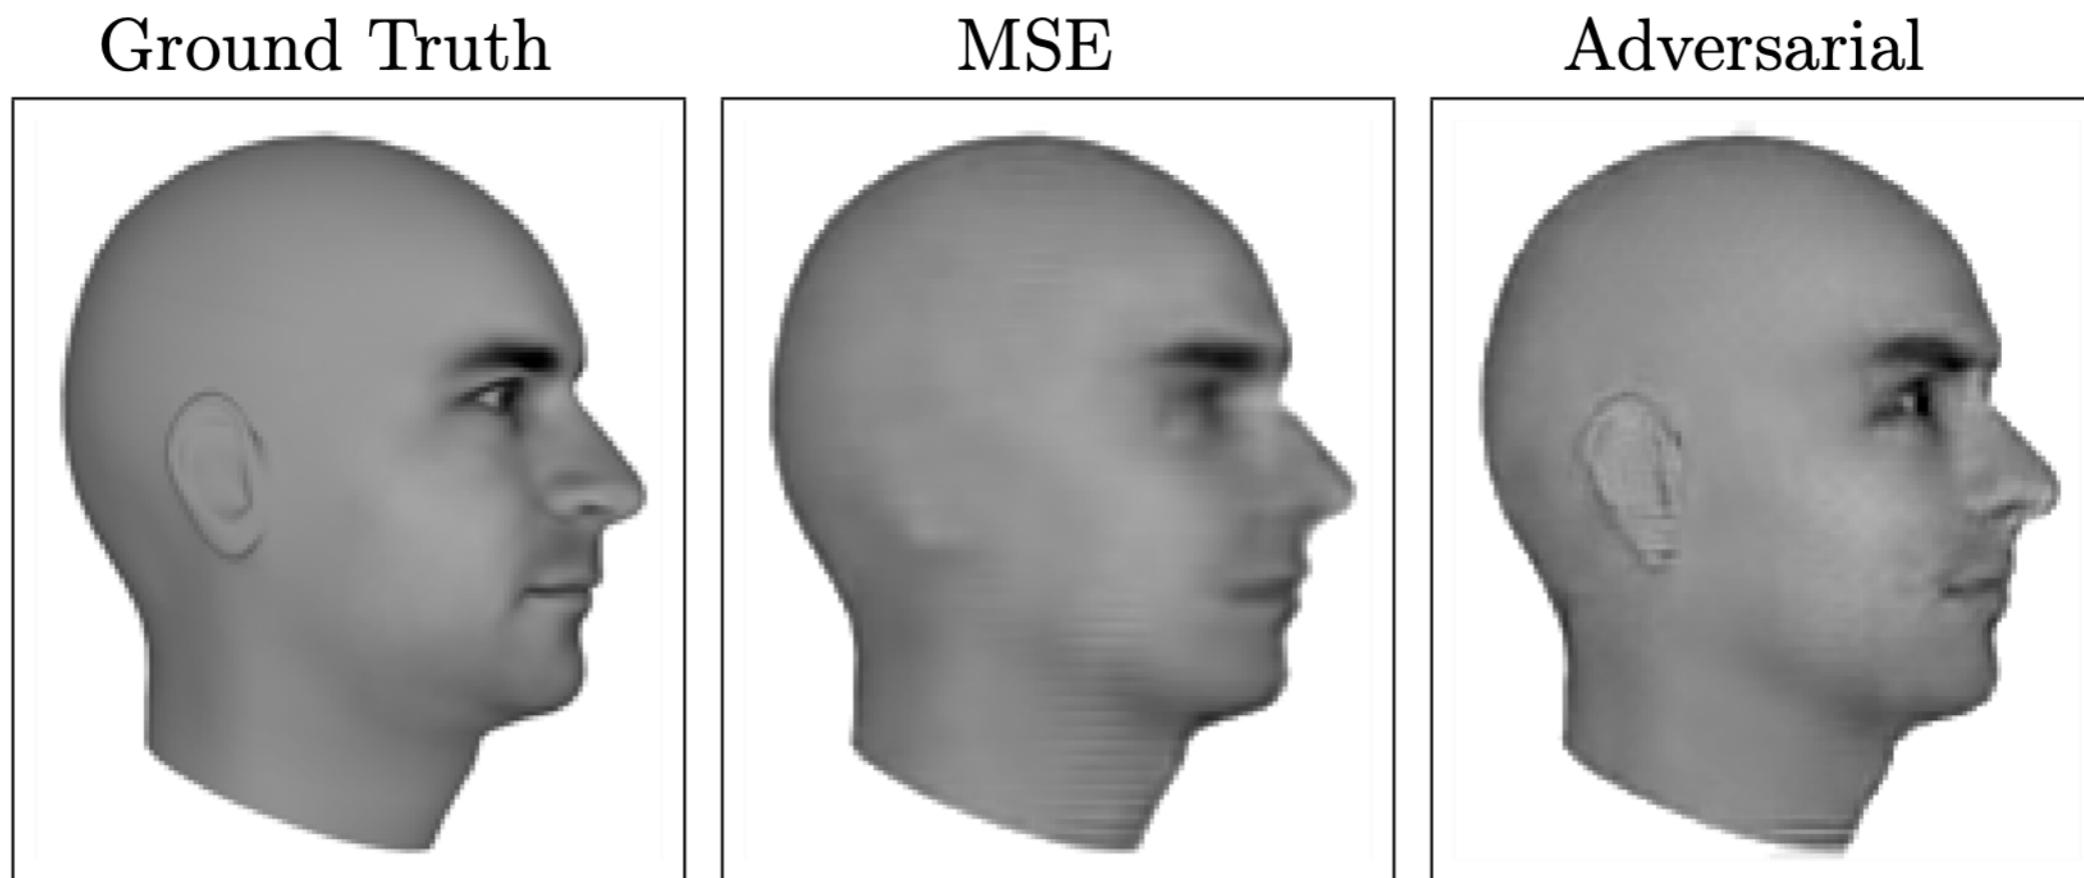


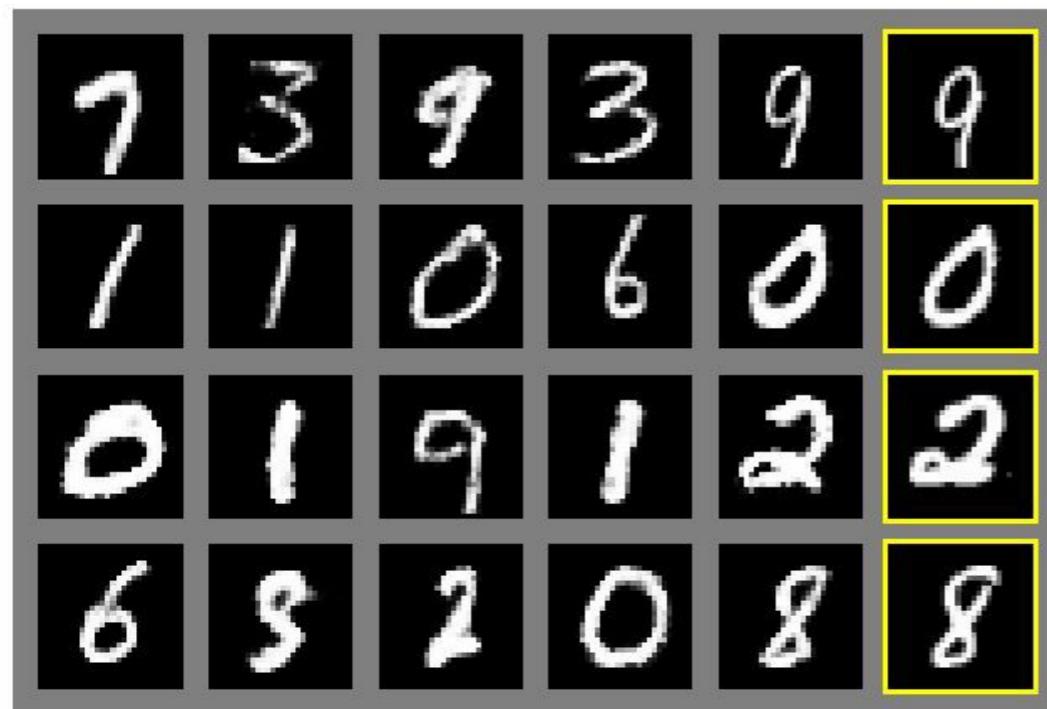
Figure 15.6

Mean squared error loses the ear because it causes a small change in few pixels. Adversarial loss preserves the ear because it is easy to notice its absence.

Ian Goodfellow et al., “Generative Adversarial Nets”, NIPS 2014

# Generative Adversarial Nets

Generated samples



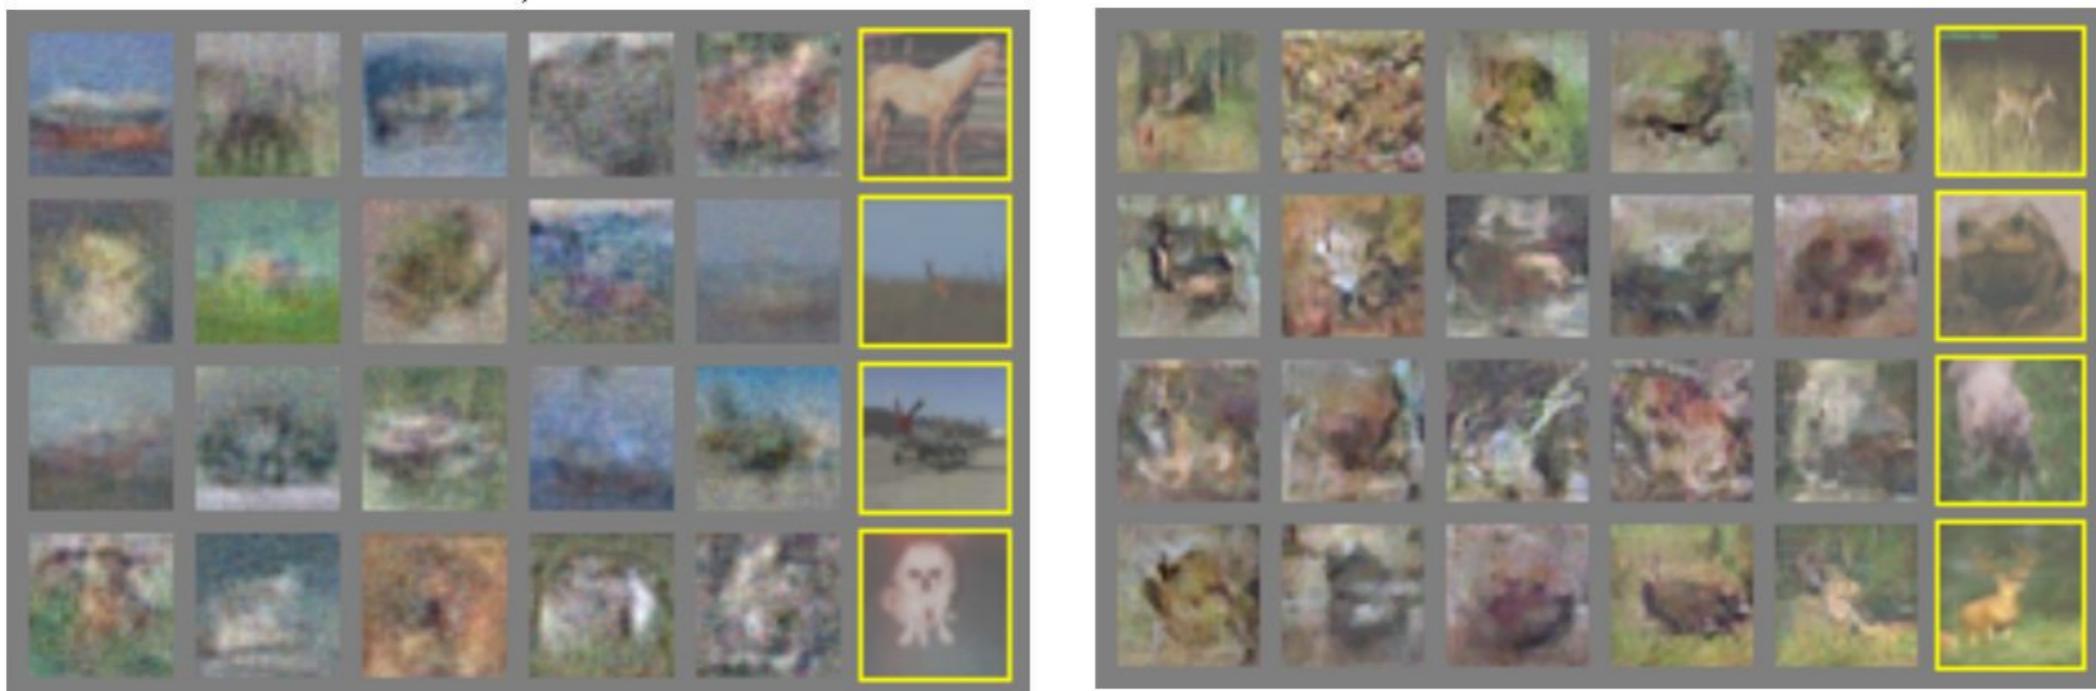
Nearest neighbor from training set

Figures copyright Ian Goodfellow et al., 2014. Reproduced with permission.

Ian Goodfellow et al., “Generative Adversarial Nets”, NIPS 2014

# Generative Adversarial Nets

Generated samples (CIFAR-10)



Nearest neighbor from training set

Figures copyright Ian Goodfellow et al., 2014. Reproduced with permission.

# **GANs + CNNs: Convolutional Architectures for Generative Adversarial Networks**

# Generative Adversarial Nets: Convolutional Architectures

Generator is an upsampling network with fractionally-strided convolutions

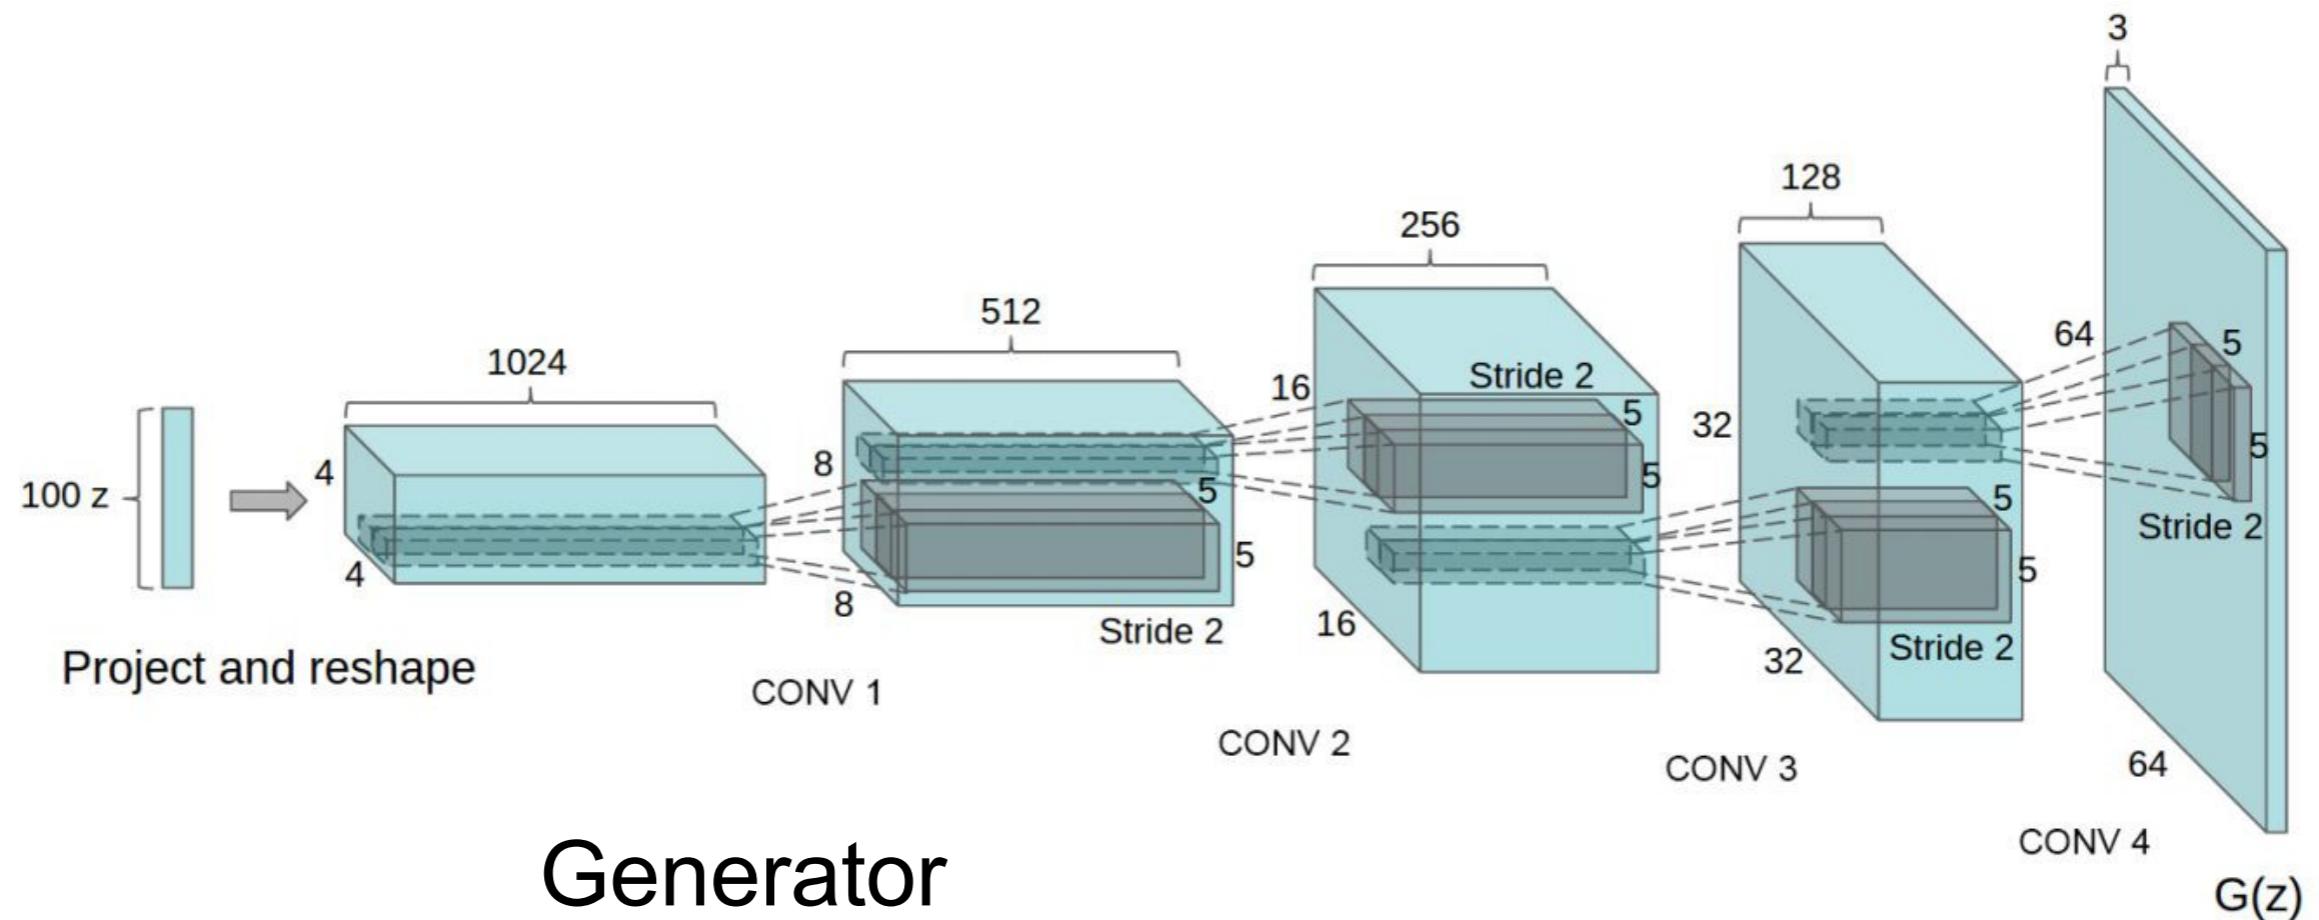
Discriminator is a convolutional network

## Architecture guidelines for stable Deep Convolutional GANs

- Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- Use batchnorm in both the generator and the discriminator.
- Remove fully connected hidden layers for deeper architectures.
- Use ReLU activation in generator for all layers except for the output, which uses Tanh.
- Use LeakyReLU activation in the discriminator for all layers.

Radford et al, “Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks”, ICLR 2016

# Generative Adversarial Nets: Convolutional Architectures



Radford et al, “Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks”, ICLR 2016

# Generative Adversarial Nets: Convolutional Architectures

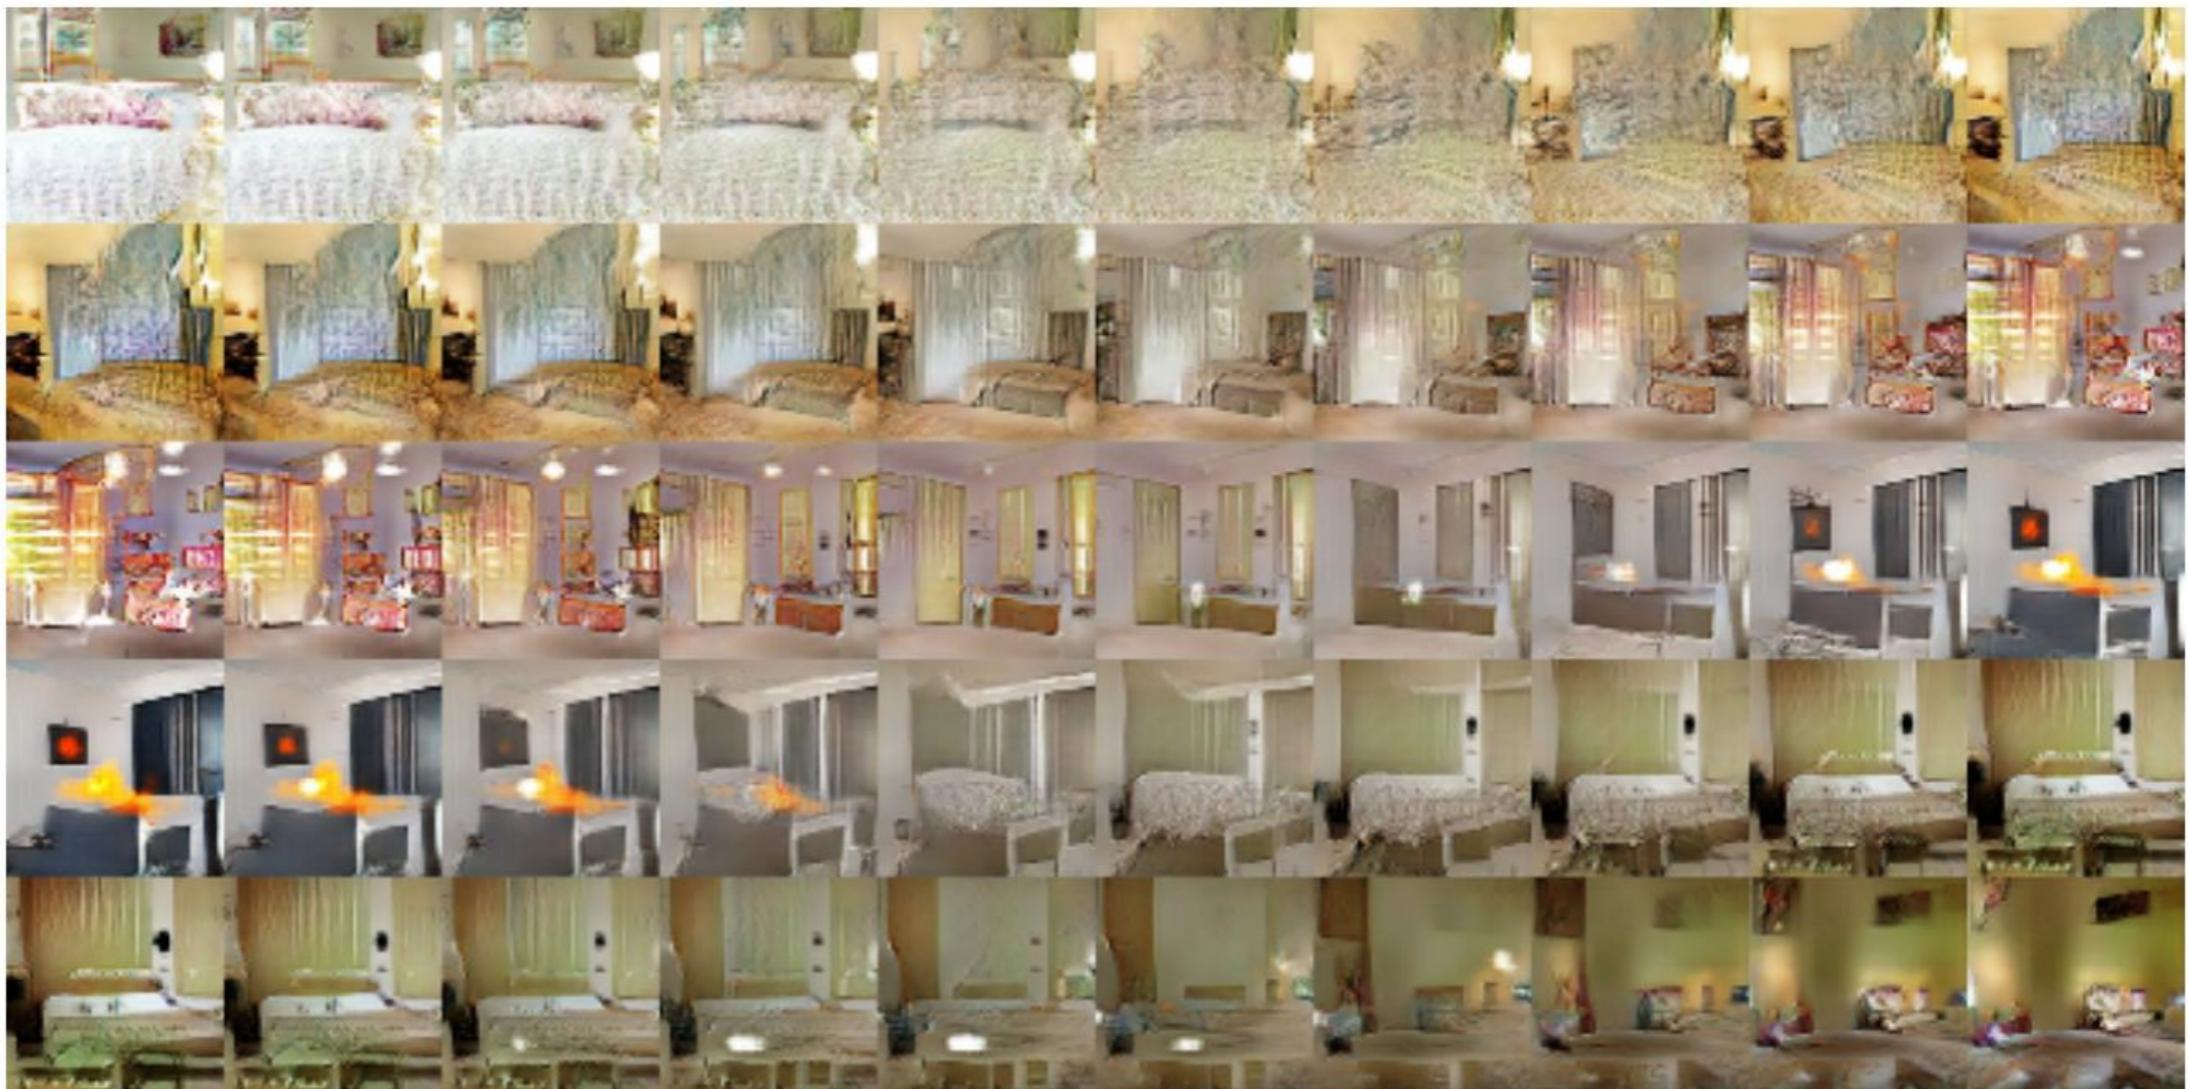
Samples  
from the  
model look  
much  
better!



Radford et al,  
ICLR 2016

# Generative Adversarial Nets: Convolutional Architectures

Interpolating  
between  
random  
points in latent  
space



Radford et al,  
ICLR 2016

# Generative Adversarial Nets: Interpretable Vector Math

Smiling woman   Neutral woman   Neutral man

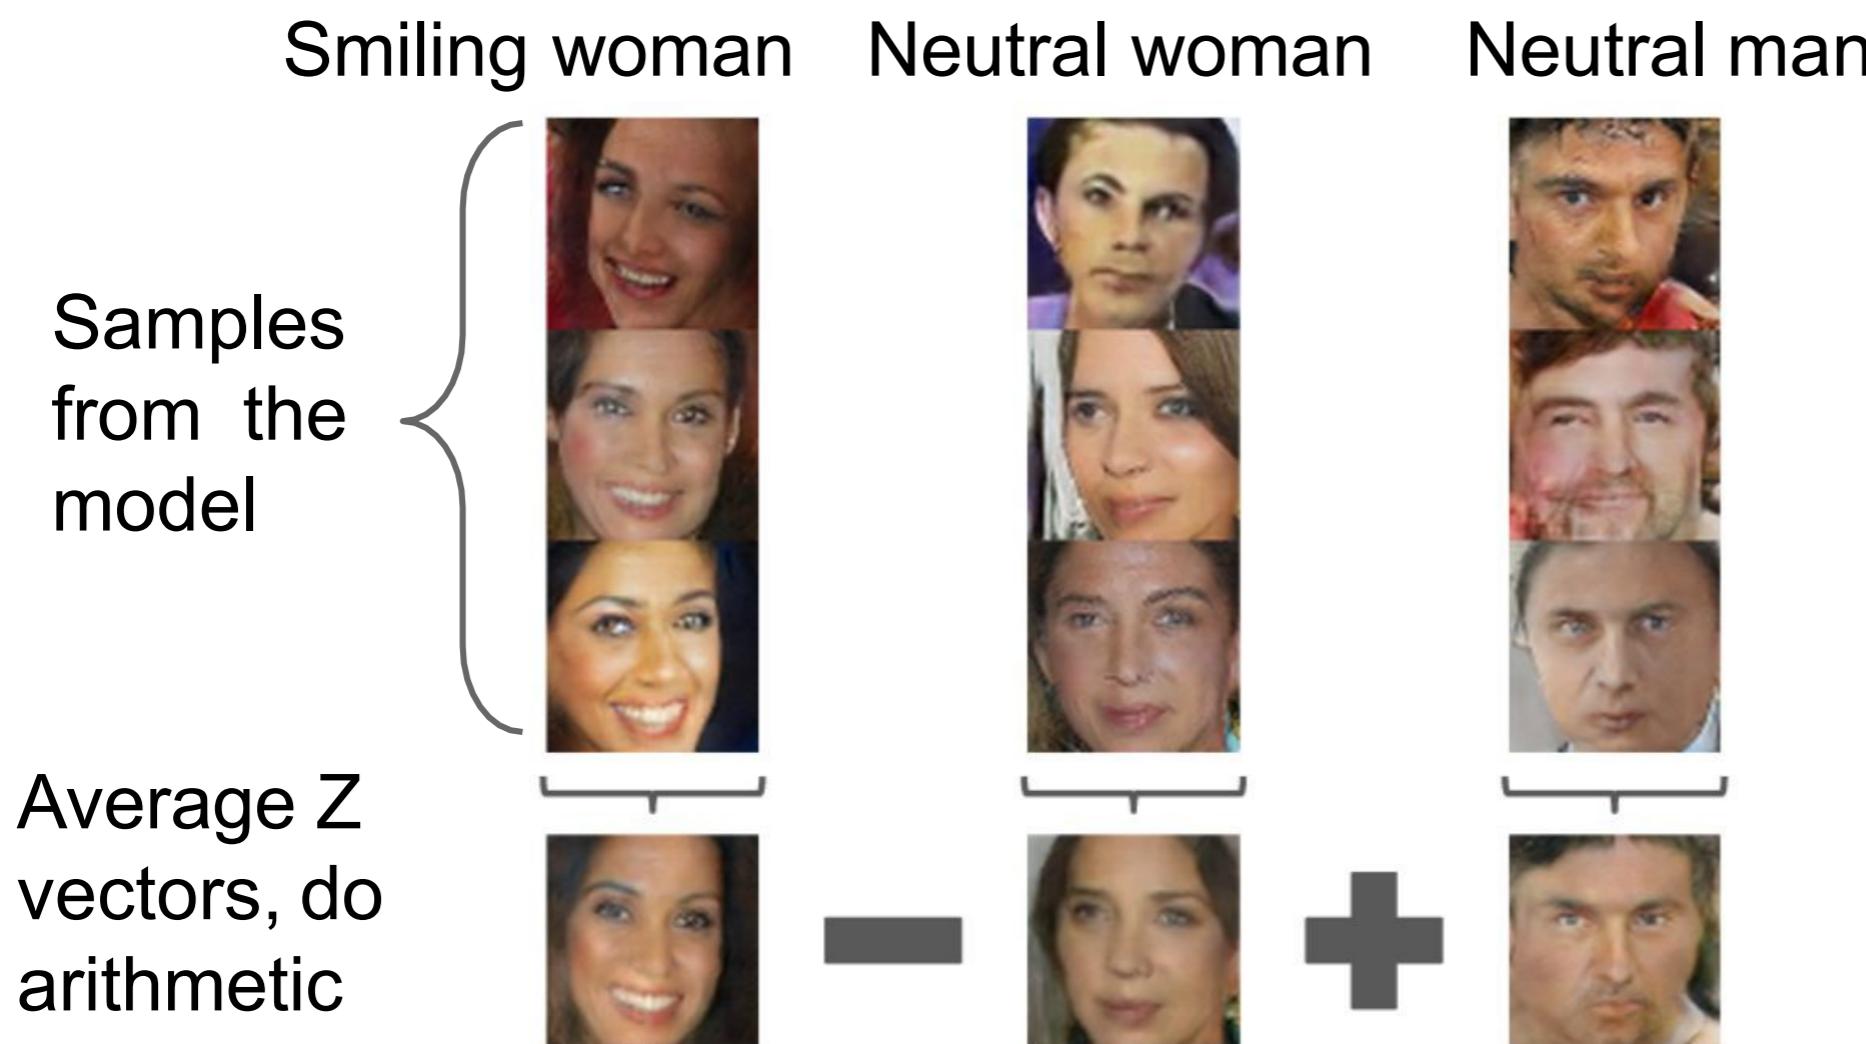
Samples  
from the  
model



Radford et al, ICLR 2016

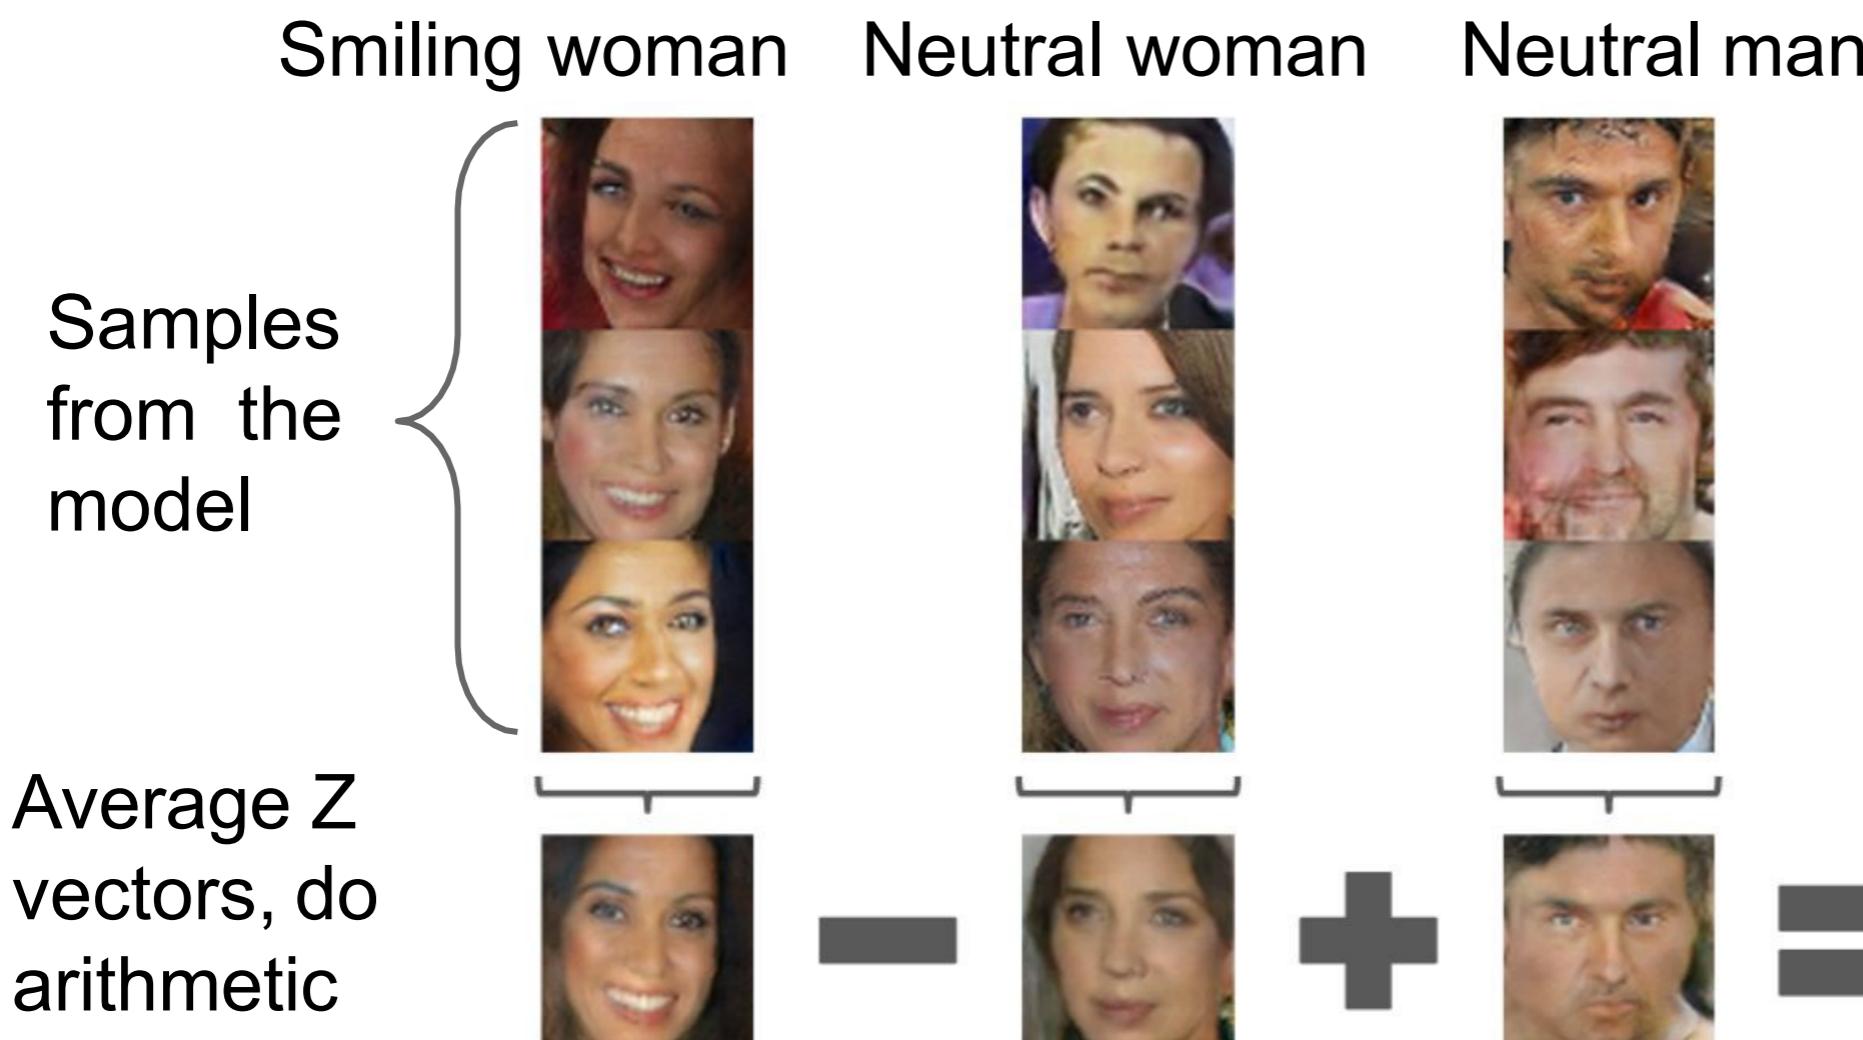
# Generative Adversarial Nets: Interpretable Vector Math

Radford et al, ICLR 2016



# Generative Adversarial Nets: Interpretable Vector Math

Radford et al, ICLR 2016

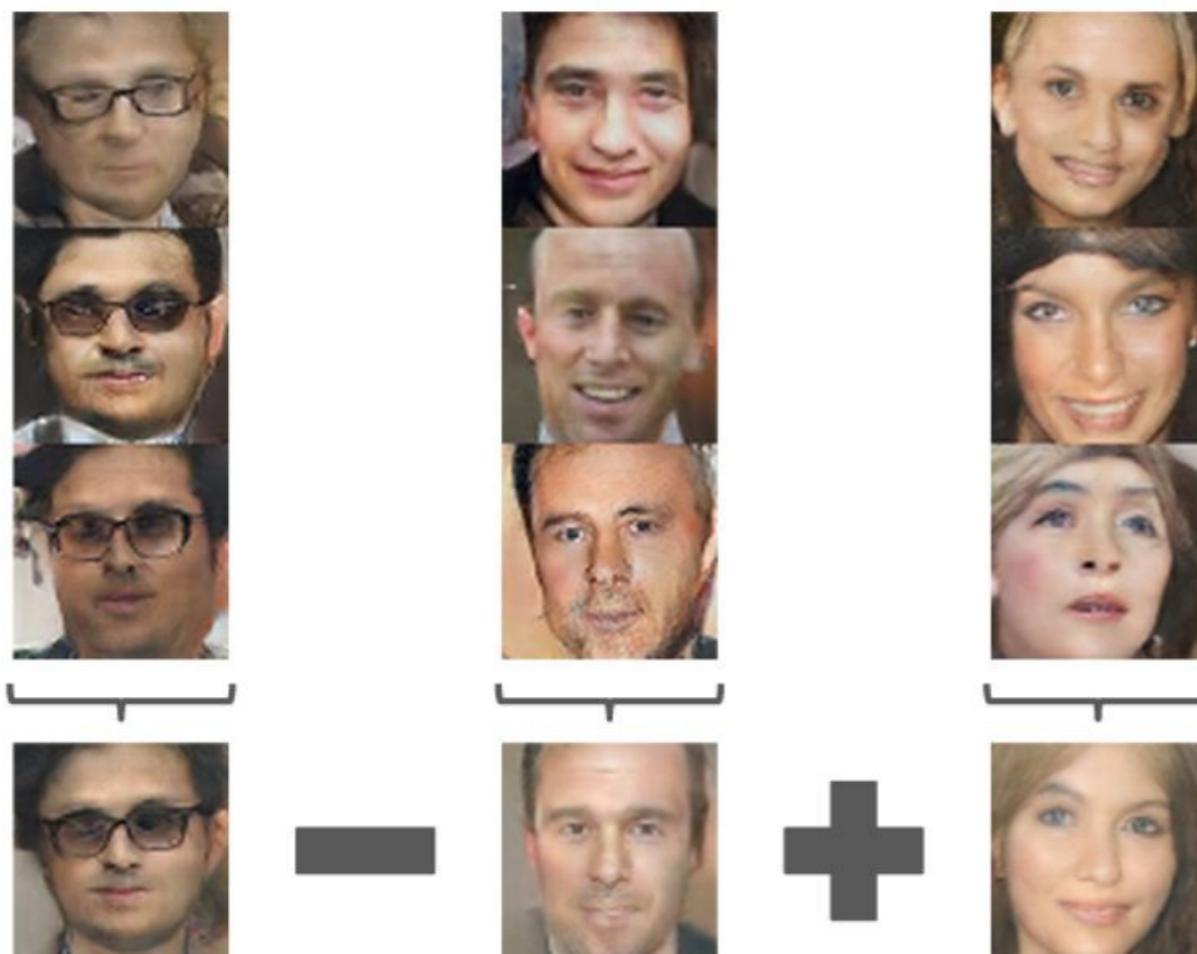


Smiling Man



# Generative Adversarial Nets: Interpretable Vector Math

Glasses man    No glasses man    No glasses woman



Radford et al,  
ICLR 2016

# Generative Adversarial Nets: Interpretable Vector Math

Glasses man



No glasses man

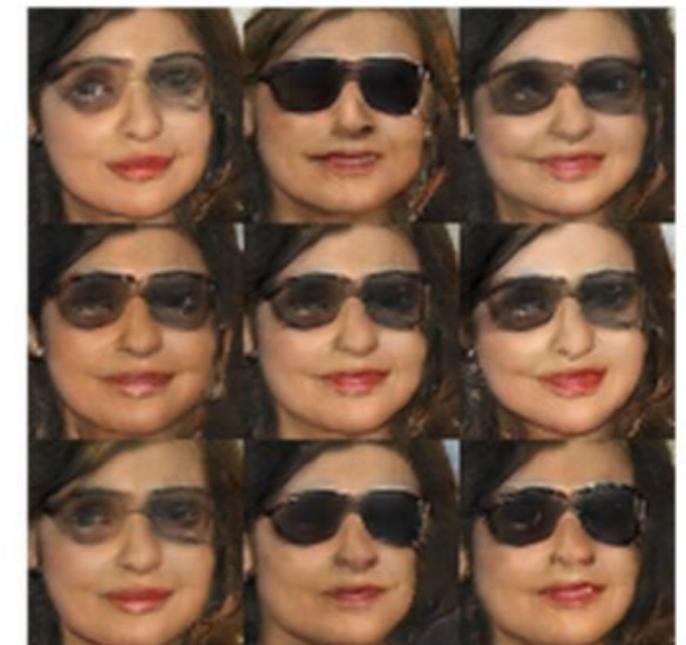


No glasses woman



Radford et al,  
ICLR 2016

Woman with glasses



# Next-Generation GANs

# 2017: Explosion of GANs

## “The GAN Zoo”

- GAN - Generative Adversarial Networks
- 3D-GAN - Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling
- acGAN - Face Aging With Conditional Generative Adversarial Networks
- AC-GAN - Conditional Image Synthesis With Auxiliary Classifier GANs
- AdaGAN - AdaGAN: Boosting Generative Models
- AEGAN - Learning Inverse Mapping by Autoencoder based Generative Adversarial Nets
- AffGAN - Amortised MAP Inference for Image Super-resolution
- AL-CGAN - Learning to Generate Images of Outdoor Scenes from Attributes and Semantic Layouts
- ALI - Adversarially Learned Inference
- AM-GAN - Generative Adversarial Nets with Labeled Data by Activation Maximization
- AnoGAN - Unsupervised Anomaly Detection with Generative Adversarial Networks to Guide Marker Discovery
- ArtGAN - ArtGAN: Artwork Synthesis with Conditional Categorical GANs
- b-GAN - b-GAN: Unified Framework of Generative Adversarial Networks
- Bayesian GAN - Deep and Hierarchical Implicit Models
- BEGAN - BEGAN: Boundary Equilibrium Generative Adversarial Networks
- BiGAN - Adversarial Feature Learning
- BS-GAN - Boundary-Seeking Generative Adversarial Networks
- CGAN - Conditional Generative Adversarial Nets
- CaloGAN - CaloGAN: Simulating 3D High Energy Particle Showers in Multi-Layer Electromagnetic Calorimeters with Generative Adversarial Networks
- CCGAN - Semi-Supervised Learning with Context-Conditional Generative Adversarial Networks
- CatGAN - Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks
- CoGAN - Coupled Generative Adversarial Networks
- Context-RNN-GAN - Contextual RNN-GANs for Abstract Reasoning Diagram Generation
- C-RNN-GAN - C-RNN-GAN: Continuous recurrent neural networks with adversarial training
- CS-GAN - Improving Neural Machine Translation with Conditional Sequence Generative Adversarial Nets
- CVAE-GAN - CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training
- CycleGAN - Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks
- DTN - Unsupervised Cross-Domain Image Generation
- DCGAN - Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks
- DiscoGAN - Learning to Discover Cross-Domain Relations with Generative Adversarial Networks
- DR-GAN - Disentangled Representation Learning GAN for Pose-Invariant Face Recognition
- DualGAN - DualGAN: Unsupervised Dual Learning for Image-to-Image Translation
- EBGAN - Energy-based Generative Adversarial Network
- f-GAN - f-GAN: Training Generative Neural Samplers using Variational Divergence Minimization
- FF-GAN - Towards Large-Pose Face Frontalization in the Wild
- GAWWN - Learning What and Where to Draw
- GeneGAN - GeneGAN: Learning Object Transfiguration and Attribute Subspace from Unpaired Data
- Geometric GAN - Geometric GAN
- GoGAN - Gang of GANs: Generative Adversarial Networks with Maximum Margin Ranking
- GP-GAN - GP-GAN: Towards Realistic High-Resolution Image Blending
- IAN - Neural Photo Editing with Introspective Adversarial Networks
- iGAN - Generative Visual Manipulation on the Natural Image Manifold
- IcGAN - Invertible Conditional GANs for image editing
- ID-CGAN - Image De-raining Using a Conditional Generative Adversarial Network
- Improved GAN - Improved Techniques for Training GANs
- InfoGAN - InfoGAN: Interpretable Representation Learning by Information Maximizing Generative Adversarial Nets
- LAGAN - Learning Particle Physics by Example: Location-Aware Generative Adversarial Networks for Physics Synthesis
- LAPGAN - Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks

<https://github.com/hindupuravinash/the-gan-zoo>

# 2017: Explosion of GANs

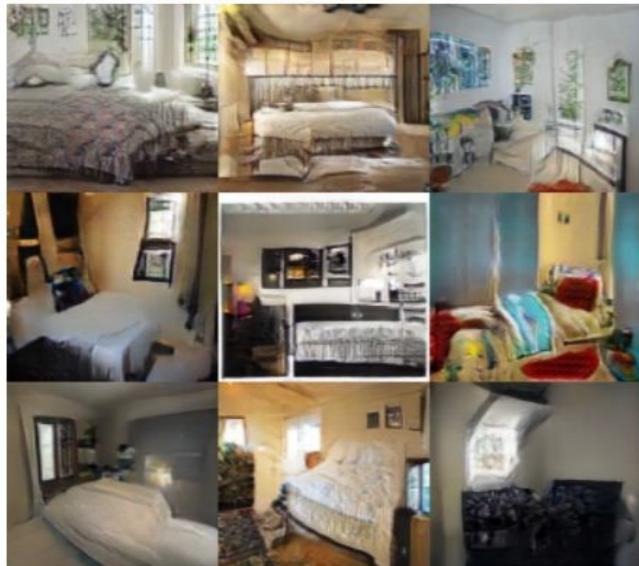
See also: <https://github.com/soumith/ganhacks> for tips  
“The GAN Zoo”

- GAN - Generative Adversarial Networks
- 3D-GAN - Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling
- acGAN - Face Aging With Conditional Generative Adversarial Networks
- AC-GAN - Conditional Image Synthesis With Auxiliary Classifier GANs
- AdaGAN - AdaGAN: Boosting Generative Models
- AEGAN - Learning Inverse Mapping by Autoencoder based Generative Adversarial Nets
- AffGAN - Amortised MAP Inference for Image Super-resolution
- AL-CGAN - Learning to Generate Images of Outdoor Scenes from Attributes and Semantic Layouts
- ALI - Adversarially Learned Inference
- AM-GAN - Generative Adversarial Nets with Labeled Data by Activation Maximization
- AnoGAN - Unsupervised Anomaly Detection with Generative Adversarial Networks to Guide Marker Discovery
- ArtGAN - ArtGAN: Artwork Synthesis with Conditional Categorical GANs
- b-GAN - b-GAN: Unified Framework of Generative Adversarial Networks
- Bayesian GAN - Deep and Hierarchical Implicit Models
- BEGAN - BEGAN: Boundary Equilibrium Generative Adversarial Networks
- BiGAN - Adversarial Feature Learning
- BS-GAN - Boundary-Seeking Generative Adversarial Networks
- CGAN - Conditional Generative Adversarial Nets
- CaloGAN - CaloGAN: Simulating 3D High Energy Particle Showers in Multi-Layer Electromagnetic Calorimeters with Generative Adversarial Networks
- CCGAN - Semi-Supervised Learning with Context-Conditional Generative Adversarial Networks
- CatGAN - Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks
- CoGAN - Coupled Generative Adversarial Networks
- Context-RNN-GAN - Contextual RNN-GANs for Abstract Reasoning Diagram Generation
- C-RNN-GAN - C-RNN-GAN: Continuous recurrent neural networks with adversarial training
- CS-GAN - Improving Neural Machine Translation with Conditional Sequence Generative Adversarial Nets
- CVAE-GAN - CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training
- CycleGAN - Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks
- DTN - Unsupervised Cross-Domain Image Generation
- DCGAN - Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks
- DiscoGAN - Learning to Discover Cross-Domain Relations with Generative Adversarial Networks
- DR-GAN - Disentangled Representation Learning GAN for Pose-Invariant Face Recognition
- DualGAN - DualGAN: Unsupervised Dual Learning for Image-to-Image Translation
- EBGAN - Energy-based Generative Adversarial Network
- f-GAN - f-GAN: Training Generative Neural Samplers using Variational Divergence Minimization
- FF-GAN - Towards Large-Pose Face Frontalization in the Wild
- GAWWN - Learning What and Where to Draw
- GeneGAN - GeneGAN: Learning Object Transfiguration and Attribute Subspace from Unpaired Data
- Geometric GAN - Geometric GAN
- GoGAN - Gang of GANs: Generative Adversarial Networks with Maximum Margin Ranking
- GP-GAN - GP-GAN: Towards Realistic High-Resolution Image Blending
- IAN - Neural Photo Editing with Introspective Adversarial Networks
- iGAN - Generative Visual Manipulation on the Natural Image Manifold
- IcGAN - Invertible Conditional GANs for image editing
- ID-CGAN - Image De-raining Using a Conditional Generative Adversarial Network
- Improved GAN - Improved Techniques for Training GANs
- InfoGAN - InfoGAN: Interpretable Representation Learning by Information Maximizing Generative Adversarial Nets
- LAGAN - Learning Particle Physics by Example: Location-Aware Generative Adversarial Networks for Physics Synthesis
- LAPGAN - Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks

<https://github.com/hindupuravinash/the-gan-zoo>

# 2017: Explosion of GANs

Better training and generation



LSGAN, Zhu 2017.



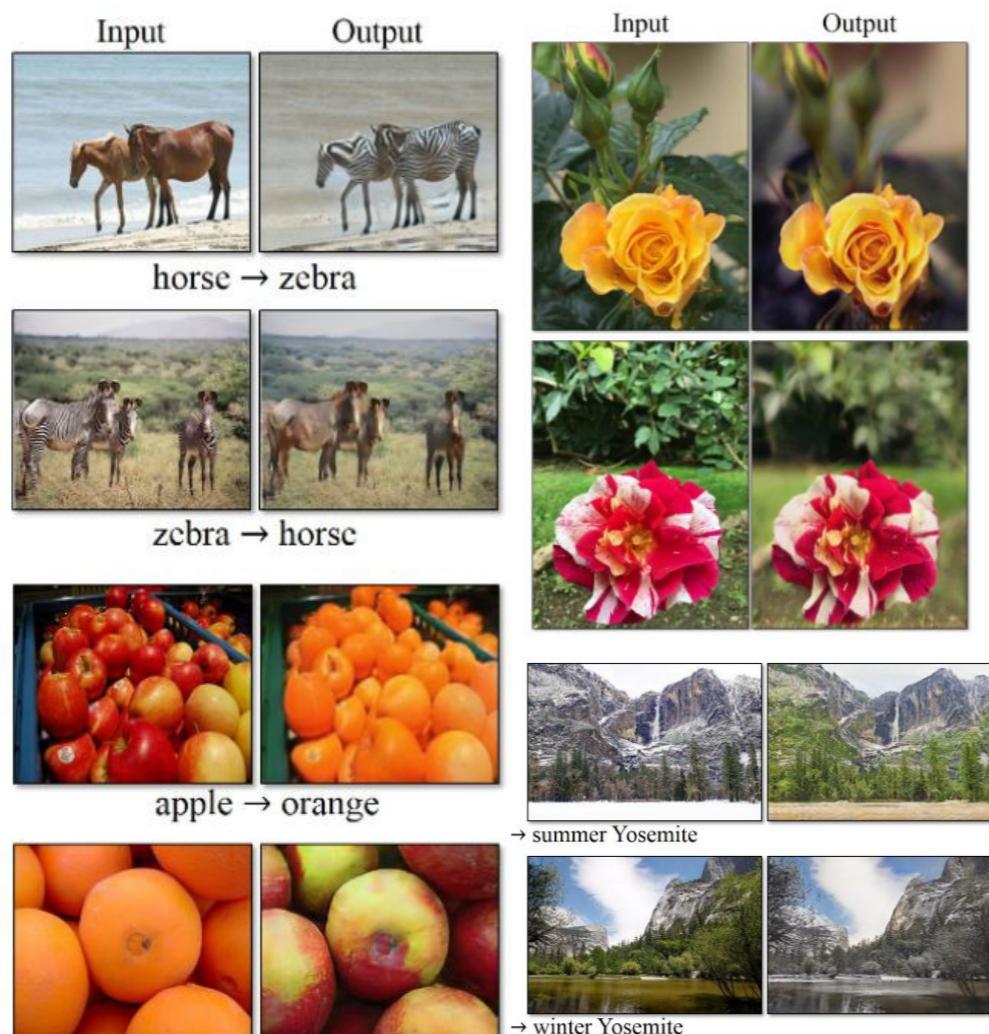
Wasserstein GAN,  
Arjovsky 2017.  
Improved Wasserstein  
GAN, Gulrajani 2017.



Progressive GAN, Karras 2018.

# 2017: Explosion of GANs

Source->Target domain transfer



CycleGAN. Zhu et al. 2017.

## Text -> Image Synthesis

this small bird has a pink breast and crown, and black primaries and secondaries.

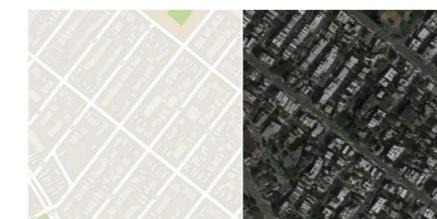
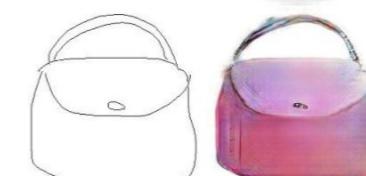


this magnificent fellow is almost all black with a red crest, and white cheek patch.



Reed et al. 2017.

## Many GAN applications



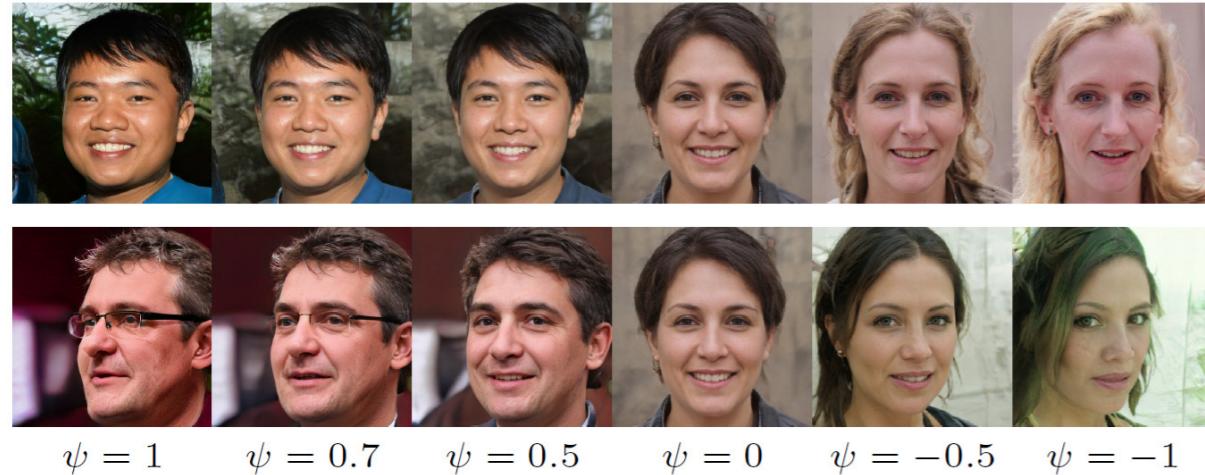
Pix2pix. Isola 2017. Many examples at <https://phillipi.github.io/pix2pix/>

# 2019: BigGAN



Brock et al., 2019

# Style GANs



$$\psi = 1 \quad \psi = 0.7 \quad \psi = 0.5 \quad \psi = 0 \quad \psi = -0.5 \quad \psi = -1$$

Figure 8. The effect of truncation trick as a function of style scale  $\psi$ . When we fade  $\psi \rightarrow 0$ , all faces converge to the “mean” face of FFHQ. This face is similar for all trained networks, and the interpolation towards it never seems to cause artifacts. By applying negative scaling to styles, we get the corresponding opposite or “anti-face”. It is interesting that various high-level attributes often flip between the opposites, including viewpoint, glasses, age, coloring, hair length, and often gender.

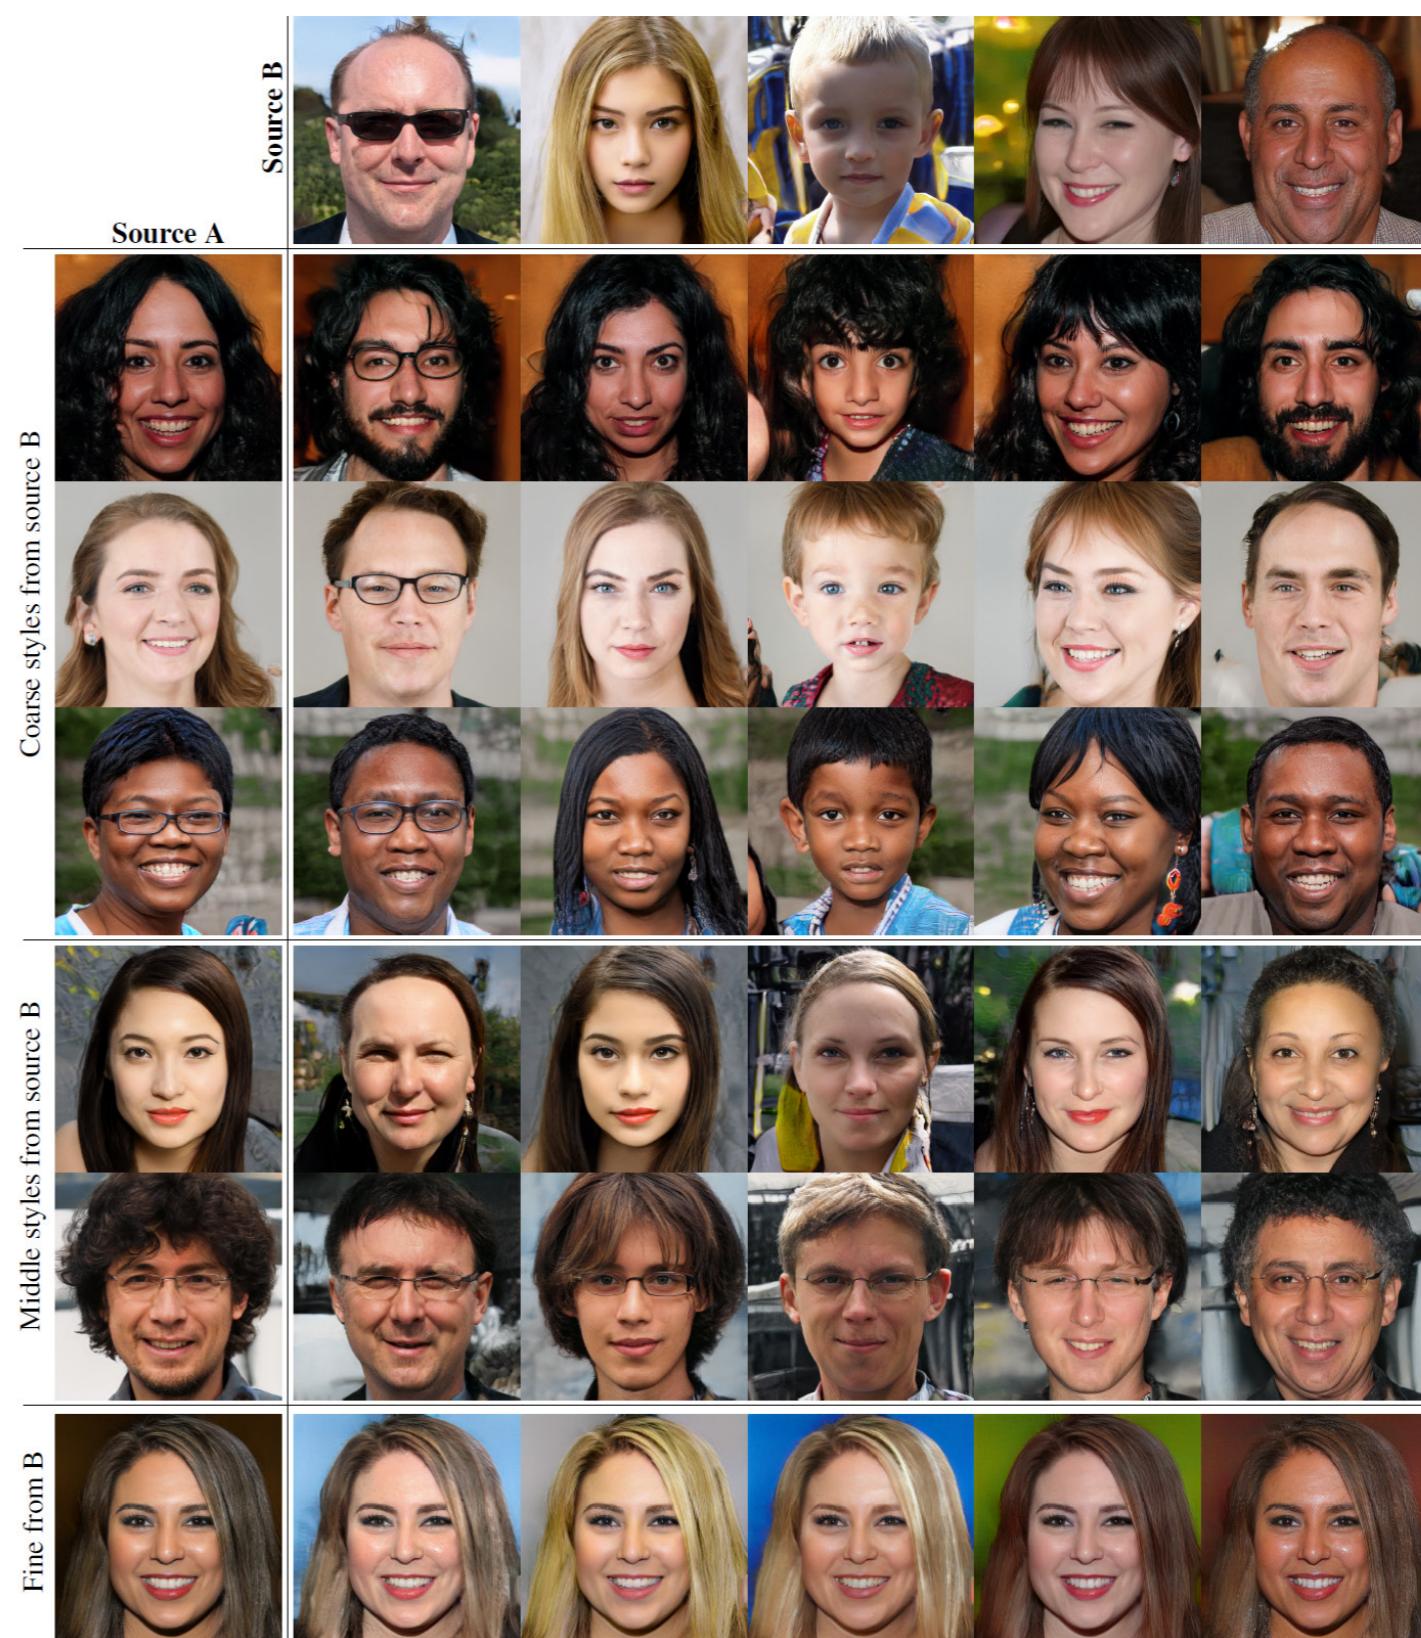
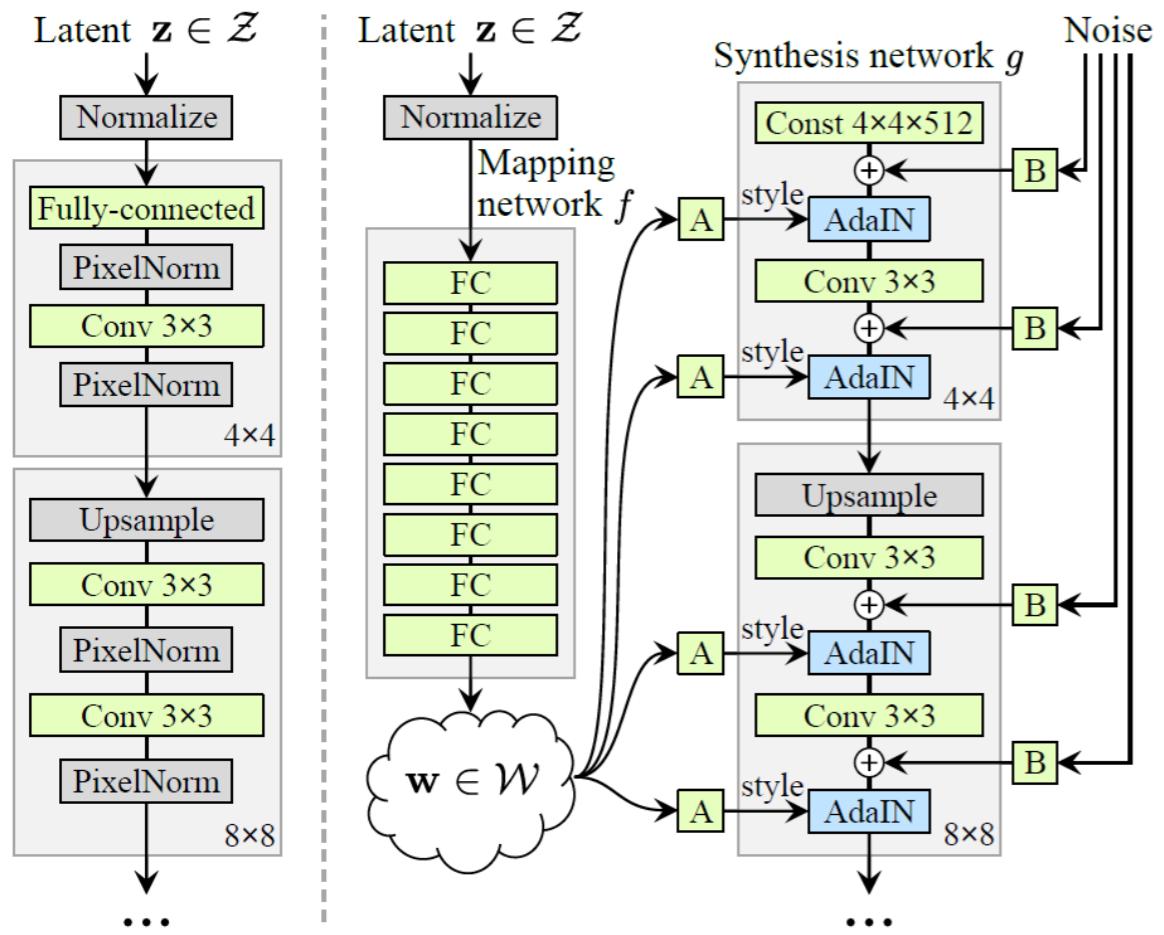
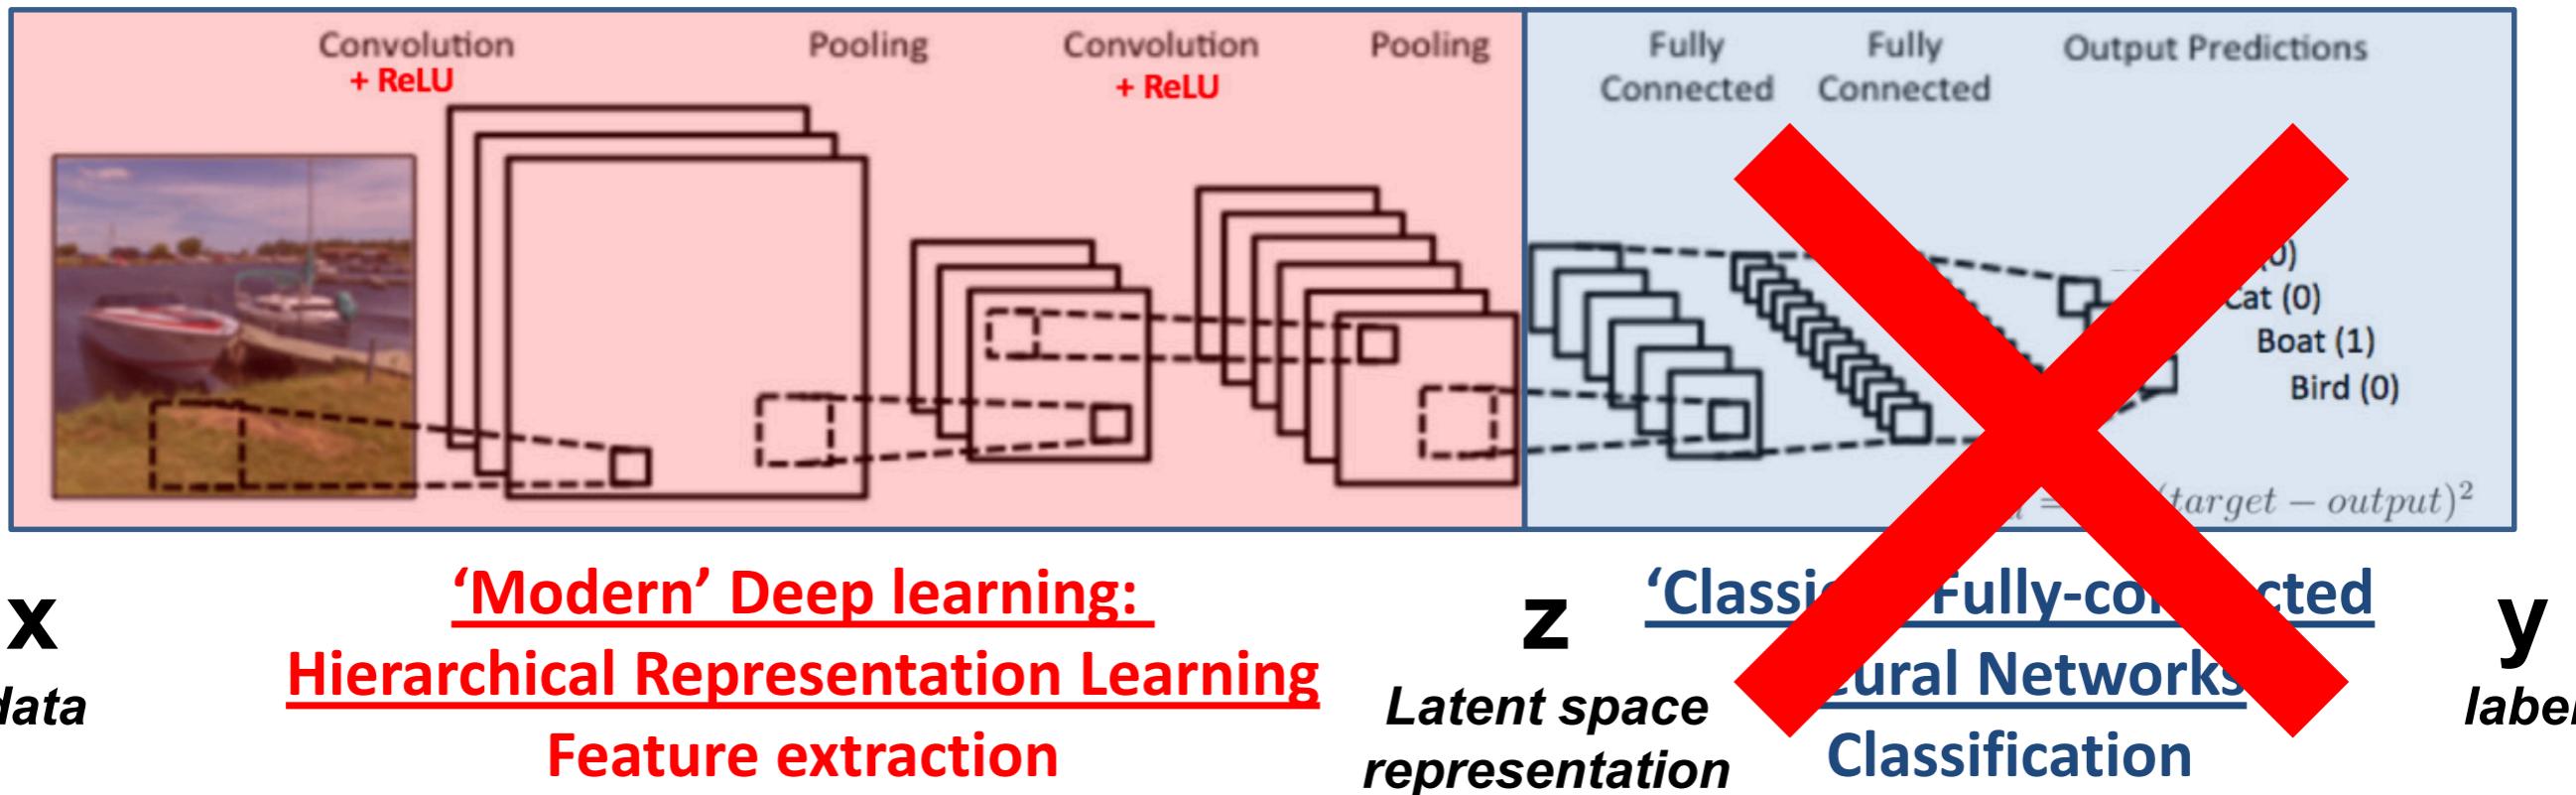


Figure 3. Two sets of images were generated from their respective latent codes (sources A and B); the rest of the images were generated by copying a specified subset of styles from source B and taking the rest from source A. Copying the styles corresponding to coarse spatial resolutions ( $4^2 - 8^2$ ) brings high-level aspects such as pose, general hair style, face shape, and eyeglasses from source B, while all colors (eyes, hair, lighting) and finer facial features resemble A. If we instead copy the styles of middle resolutions ( $16^2 - 32^2$ ) from B, we inherit smaller scale facial features, hair style, eyes open/closed from B, while the pose, general face shape, and eyeglasses from A are preserved. Finally, copying the fine styles ( $64^2 - 1024^2$ ) from B brings mainly the color scheme and microstructure.

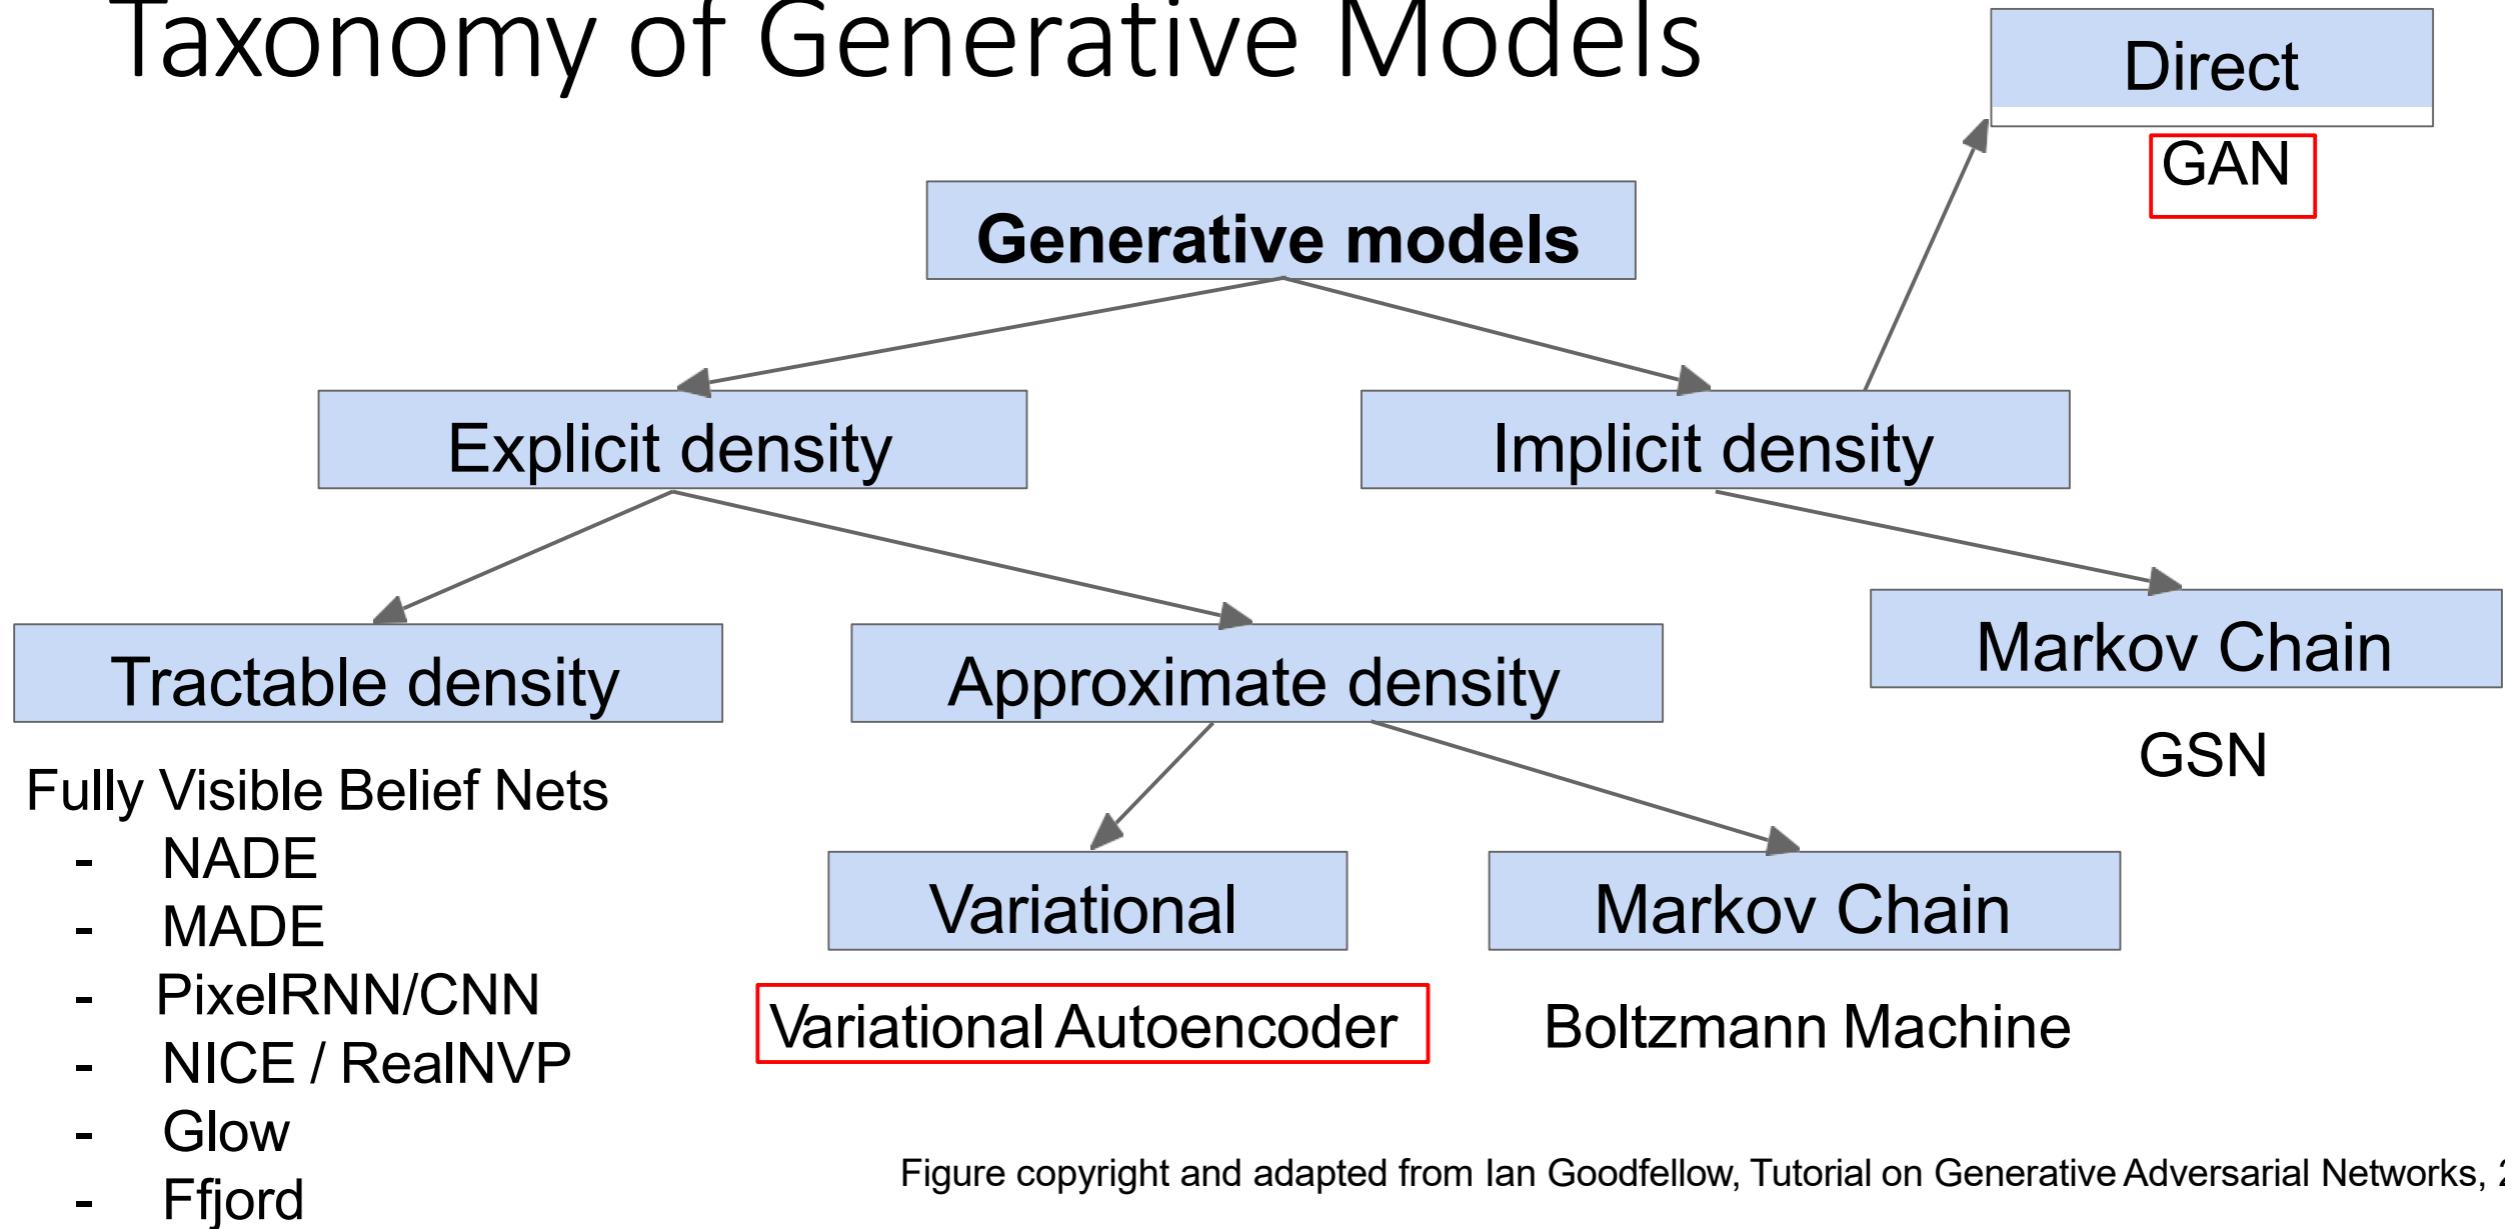
# Representation learning without annotations?



**Many ideas are possible (and yours could be even better!):**

1. Predict the future: RNNs, Video
2. Compression: Autoencoder (predict self, through clamp), representation **z**
3. Pretext tasks: predict self, before/after, missing patch, correct rotation, colorization, up-sampling, multimodal
4. Capture parameter distribution (variance): Variational Auto-Encoders
5. Make latent space parameters **z** meaningful, orthogonal, explicit, tuneable
6. Train using a second network: GANs - Improve quality of output images
7. The Sky is the Limit

# Taxonomy of Generative Models



**GANs  $\leftrightarrow$  VAEs**

Let's look at VAE and GAN more closely...

From Eric Xing's slides for CMU 10-708  
Based on "[On Unifying Deep Generative Models](#)"



# Variational Autoencoders (VAEs)

- [Kingma & Welling, 2014]
- Use variational inference with an inference model
  - Enjoy similar applicability with wake-sleep algorithm
- Generative model  $p_\theta(\mathbf{x}|\mathbf{z})$ , and prior  $p(\mathbf{z})$ 
  - Joint distribution  $p_\theta(\mathbf{x}, \mathbf{z}) = p_\theta(\mathbf{x}|\mathbf{z})p(\mathbf{z})$
- Inference model  $q_\phi(\mathbf{z}|\mathbf{x})$

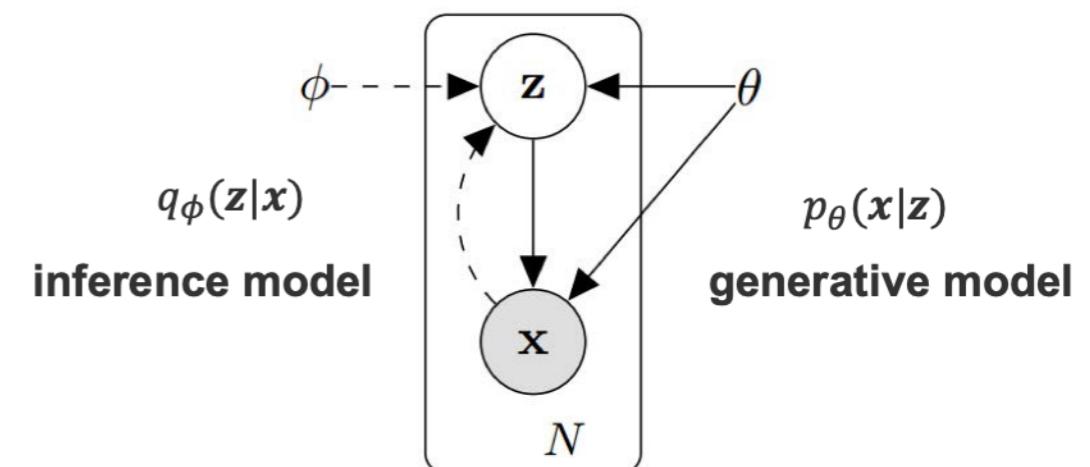


Figure courtesy: Kingma & Welling, 2014





# Generative Adversarial Nets (GANs)

- [Goodfellow et al., 2014]
- Generative model  $\mathbf{x} = G_{\theta}(\mathbf{z})$ ,  $\mathbf{z} \sim p(\mathbf{z})$ 
  - Map noise variable  $\mathbf{z}$  to data space  $\mathbf{x}$
  - Define an **implicit distribution** over  $\mathbf{x}$ :  $p_{g_{\theta}}(\mathbf{x})$ 
    - a stochastic process to simulate data  $\mathbf{x}$
    - Intractable to evaluate likelihood
- Discriminator  $D_{\phi}(\mathbf{x})$ 
  - Output the probability that  $\mathbf{x}$  came from the data rather than the generator
- No explicit inference model
- No obvious connection to previous models with inference networks like VAEs
  - We will build formal connections between GANs and VAEs later





# A unified view of deep generative models

- ❑ Literatures have viewed these DGM approaches as distinct model training paradigms
  - ❑ GANs: achieve an equilibrium between generator and discriminator
  - ❑ VAEs: maximize lower bound of the data likelihood
- ❑ Let's study a new formulation for DGMs
  - ❑ Connects GANs, VAEs, and other variants, under a unified view
  - ❑ Links them back to inference and learning of Graphical Models, and the wake-sleep heuristic that approximates this
  - ❑ Provides a tool to analyze many GAN-VAE-based algorithms
  - ❑ Encourages mutual exchange of ideas from each individual class of models





## Generative Adversarial Nets (GANs):

- Implicit distribution over  $\mathbf{x} \sim p_{\theta}(\mathbf{x}|y)$

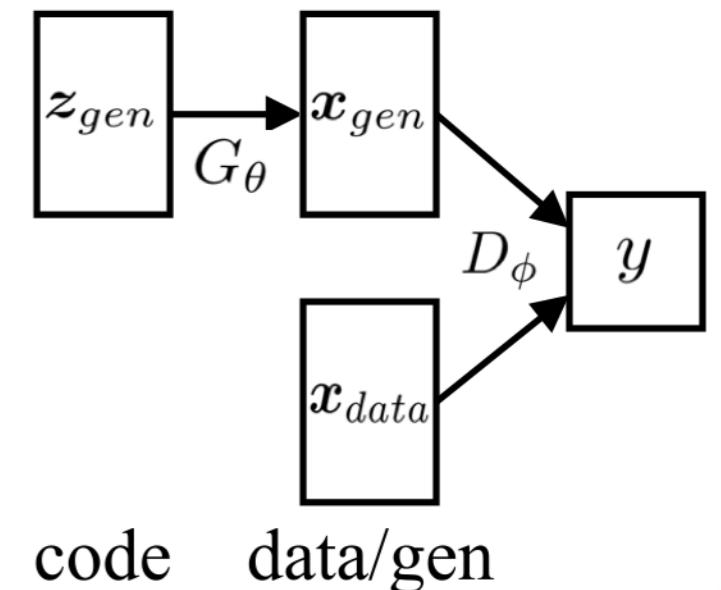
$$p_{\theta}(\mathbf{x}|y) = \begin{cases} p_{g_{\theta}}(\mathbf{x}) & y = 0 \\ p_{data}(\mathbf{x}) & y = 1. \end{cases}$$

(distribution of generated images)  
(distribution of real images)

- $\mathbf{x} \sim p_{g_{\theta}}(\mathbf{x}) \Leftrightarrow \mathbf{x} = G_{\theta}(\mathbf{z}), \mathbf{z} \sim p(\mathbf{z}|y = 0)$

- $\mathbf{x} \sim p_{data}(\mathbf{x})$

- the code space of  $\mathbf{z}$  is degenerated
- sample directly from data





## A new formulation

- Rewrite GAN objectives in the "variational-EM" format
- Recap: conventional formulation:

$$\max_{\phi} \mathcal{L}_{\phi} = \mathbb{E}_{\mathbf{x}=G_{\theta}(\mathbf{z}), \mathbf{z} \sim p(\mathbf{z}|y=0)} [\log(1 - D_{\phi}(\mathbf{x}))] + \mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\log D_{\phi}(\mathbf{x})]$$

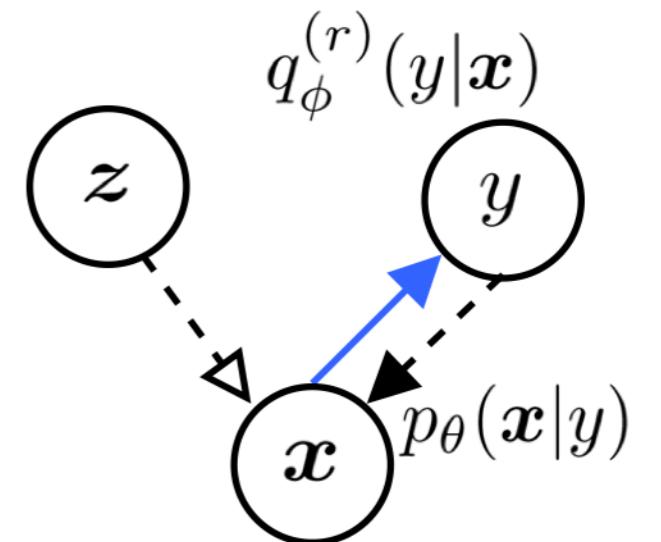
$$\begin{aligned} \max_{\theta} \mathcal{L}_{\theta} &= \mathbb{E}_{\mathbf{x}=G_{\theta}(\mathbf{z}), \mathbf{z} \sim p(\mathbf{z}|y=0)} [\log D_{\phi}(\mathbf{x})] + \mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\log(1 - D_{\phi}(\mathbf{x}))] \\ &= \mathbb{E}_{\mathbf{x}=G_{\theta}(\mathbf{z}), \mathbf{z} \sim p(\mathbf{z}|y=0)} [\log D_{\phi}(\mathbf{x})] \end{aligned}$$

- Rewrite in the new form
  - Implicit distribution over  $\mathbf{x} \sim p_{\theta}(\mathbf{x}|y)$   
 $\mathbf{x} = G_{\theta}(\mathbf{z}), \mathbf{z} \sim p(\mathbf{z}|y)$
  - Discriminator distribution  $q_{\phi}(y|\mathbf{x})$   
 $q_{\phi}^r(y|\mathbf{x}) = q_{\phi}(1 - y|\mathbf{x})$  (reverse)

$$\max_{\phi} \mathcal{L}_{\phi} = \mathbb{E}_{p_{\theta}(\mathbf{x}|y)p(y)} [\log q_{\phi}(y|\mathbf{x})]$$

$$\max_{\theta} \mathcal{L}_{\theta} = \mathbb{E}_{p_{\theta}(\mathbf{x}|y)p(y)} [\log q_{\phi}^r(y|\mathbf{x})]$$

123





# GANs vs. Variational EM

## Variational EM

- ❑ Objectives

$$\max_{\phi} \mathcal{L}_{\phi, \theta} = \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x|z)] + KL(q_{\phi}(z|x) || p(z))$$

$$\max_{\theta} \mathcal{L}_{\phi, \theta} = \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x|z)] + KL(q_{\phi}(z|x) || p(z))$$

- ❑ Single objective for both  $\theta$  and  $\phi$
- ❑ Extra prior regularization by  $p(z)$
- ❑ The reconstruction term: maximize the conditional log-likelihood of  $x$  with the generative distribution  $p_{\theta}(x|z)$  conditioning on the latent code  $z$  inferred by  $q_{\phi}(z|x)$



- ❑  $p_{\theta}(x|z)$  is the generative model
- ❑  $q_{\phi}(z|x)$  is the inference model

## GAN

- ❑ Objectives

$$\max_{\phi} \mathcal{L}_{\phi} = \mathbb{E}_{p_{\theta}(\mathbf{x}|y)p(y)} [\log q_{\phi}(y|\mathbf{x})]$$

$$\max_{\theta} \mathcal{L}_{\theta} = \mathbb{E}_{p_{\theta}(\mathbf{x}|y)p(y)} [\log q_{\phi}^r(y|\mathbf{x})]$$

- ❑ Two objectives
- ❑ Have global optimal state in the game theoretic view
- ❑ The objectives: maximize the conditional log-likelihood of  $y$  (or  $1 - y$ ) with the distribution  $q_{\phi}(y|\mathbf{x})$  conditioning on data/generation  $x$  inferred by  $p_{\theta}(x|y)$



- ❑ Interpret  $q_{\phi}(y|\mathbf{x})$  as the generative model
- ❑ Interpret  $p_{\theta}(x|y)$  as the inference model





# GANs vs. Variational EM

## Variational EM

- Objectives

$$\max_{\phi} \mathcal{L}_{\phi, \theta} = \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x|z)] + KL(q_{\phi}(z|x) || p(z))$$

$$\max_{\theta} \mathcal{L}_{\phi, \theta} = \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x|z)] + KL(q_{\phi}(z|x) || p(z))$$

- Single objective for both  $\theta$  and  $\phi$
- Extra prior regularization by  $p(z)$
- The reconstruction term: maximize the conditional log-likelihood of  $x$  with the generative distribution  $p_{\theta}(x|z)$  conditioning on the latent code  $z$  inferred by  $q_{\phi}(z|x)$



- $p_{\theta}(x|z)$  is the generative model
- $q_{\phi}(z|x)$  is the inference model

- Interpret  $x$  as latent variables
- Interpret generation of  $x$  as performing inference over latent

In VEM, we minimize the following:

$$F(\theta, \phi; x) = -\log p(x) + KL(q_{\phi}(z|x) || p_{\theta}(z|x))$$

$\Rightarrow$  KL (inference model | posterior)

## GAN

- Objectives

$$\max_{\phi} \mathcal{L}_{\phi} = \mathbb{E}_{p_{\theta}(x|y)p(y)} [\log q_{\phi}(y|x)]$$

$$\max_{\theta} \mathcal{L}_{\theta} = \mathbb{E}_{p_{\theta}(x|y)p(y)} [\log q_{\phi}^r(y|x)]$$

- Two objectives
- Have global optimal state in the game theoretic view
- The objectives: maximize the conditional log-likelihood of  $y$  (or  $1 - y$ ) with the distribution  $q_{\phi}(y|x)$  conditioning on data/generation  $x$  inferred by  $p_{\theta}(x|y)$



- Interpret  $q_{\phi}(y|x)$  as the generative model
- Interpret  $p_{\theta}(x|y)$  as the inference model

125





## GANs: minimizing KLD

- As in Variational EM, we can further rewrite in the form of **minimizing KLD** to reveal more insights into the optimization problem
- For each optimization step of  $p_\theta(\mathbf{x}|y)$  at point  $(\theta = \theta_0, \phi = \phi_0)$ , let
  - $p(y)$ : uniform prior distribution
  - $p_{\theta=\theta_0}(\mathbf{x}) = \mathbb{E}_{p(y)}[p_{\theta=\theta_0}(\mathbf{x}|y)]$
  - $q^r(\mathbf{x}|y) \propto q_{\phi=\phi_0}^r(y|\mathbf{x})p_{\theta=\theta_0}(\mathbf{x})$
- *Lemma 1:* The updates of  $\theta$  at  $\theta_0$  have

$$\nabla_\theta \left[ -\mathbb{E}_{p_\theta(\mathbf{x}|y)p(y)} [\log q_{\phi=\phi_0}^r(y|\mathbf{x})] \right] \Big|_{\theta=\theta_0} =$$
$$\nabla_\theta \left[ \mathbb{E}_{p(y)} [KL(p_\theta(\mathbf{x}|y)\|q^r(\mathbf{x}|y))] - JSD(p_\theta(\mathbf{x}|y=0)\|p_\theta(\mathbf{x}|y=1)) \right] \Big|_{\theta=\theta_0};$$

- KL: KL divergence
- JSD: Jensen-shannon divergence





# GANs: minimizing KLD

- *Lemma 1:* The updates of  $\theta$  at  $\theta_0$  have

$$\begin{aligned} \nabla_{\theta} \left[ -\mathbb{E}_{p_{\theta}(\mathbf{x}|y)p(y)} [\log q_{\phi=\phi_0}^r(y|\mathbf{x})] \right] \Big|_{\theta=\theta_0} = \\ \nabla_{\theta} \left[ \mathbb{E}_{p(y)} [\text{KL}(p_{\theta}(\mathbf{x}|y) \| q^r(\mathbf{x}|y))] - \text{JSD}(p_{\theta}(\mathbf{x}|y=0) \| p_{\theta}(\mathbf{x}|y=1)) \right] \Big|_{\theta=\theta_0} \end{aligned}$$

- Connection to variational inference
  - See  $\mathbf{x}$  as latent variables,  $y$  as visible
  - $p_{\theta=\theta_0}(\mathbf{x})$ : prior distribution
  - $q^r(\mathbf{x}|y) \propto q_{\phi=\phi_0}^r(y|\mathbf{x})p_{\theta=\theta_0}(\mathbf{x})$  : posterior distribution
  - $p_{\theta}(\mathbf{x}|y)$ : variational distribution
    - Amortized inference: updates model parameter  $\theta$
- Suggests relations to VAEs, as we will explore shortly

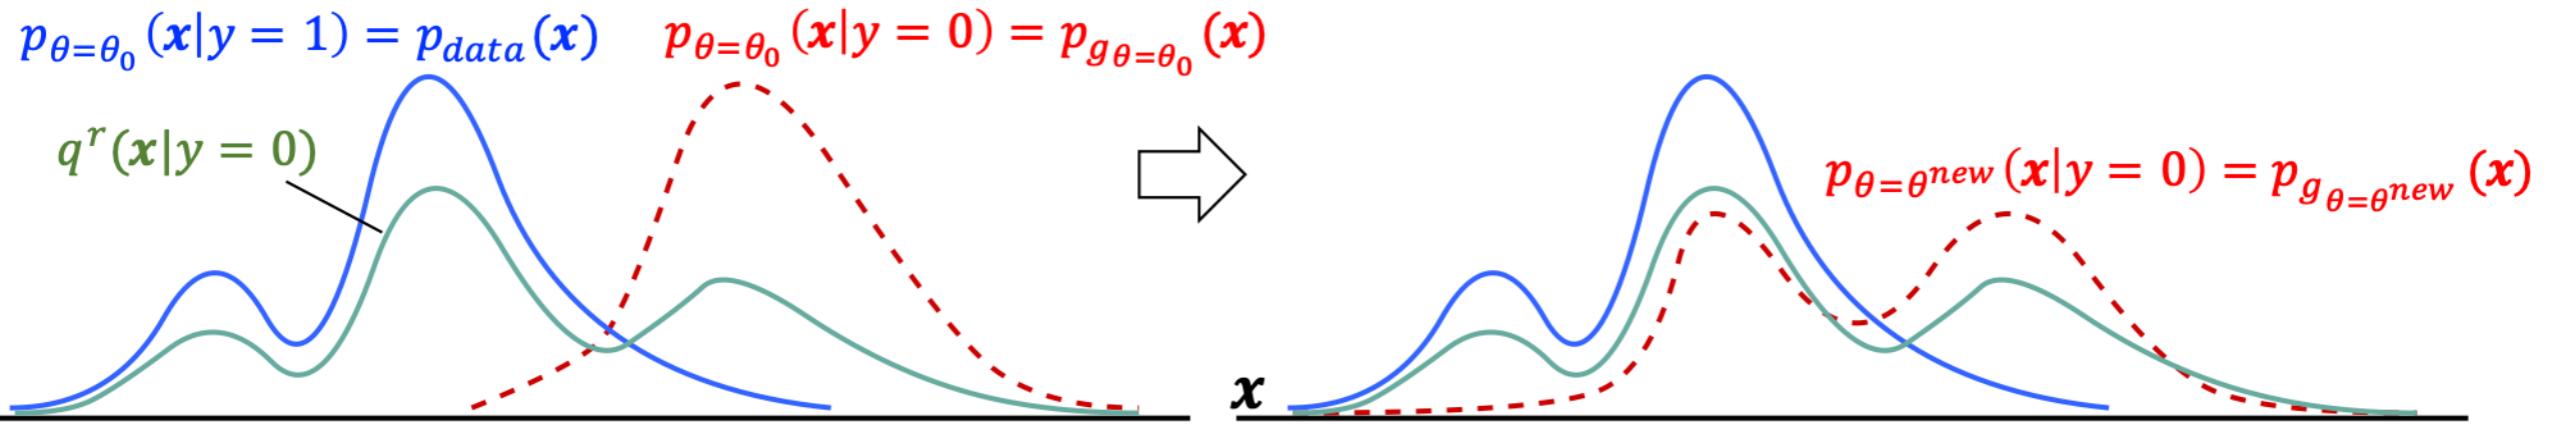
In VEM, we minimize the following:

$$F(\theta, \phi; \mathbf{x}) = -\log p(\mathbf{x}) + KL(q_{\phi}(\mathbf{z}|\mathbf{x}) || p_{\theta}(\mathbf{z}|\mathbf{x}))$$

→ ~~KL (inference model | posterior)~~<sup>127</sup>



# GANs: minimizing KLD



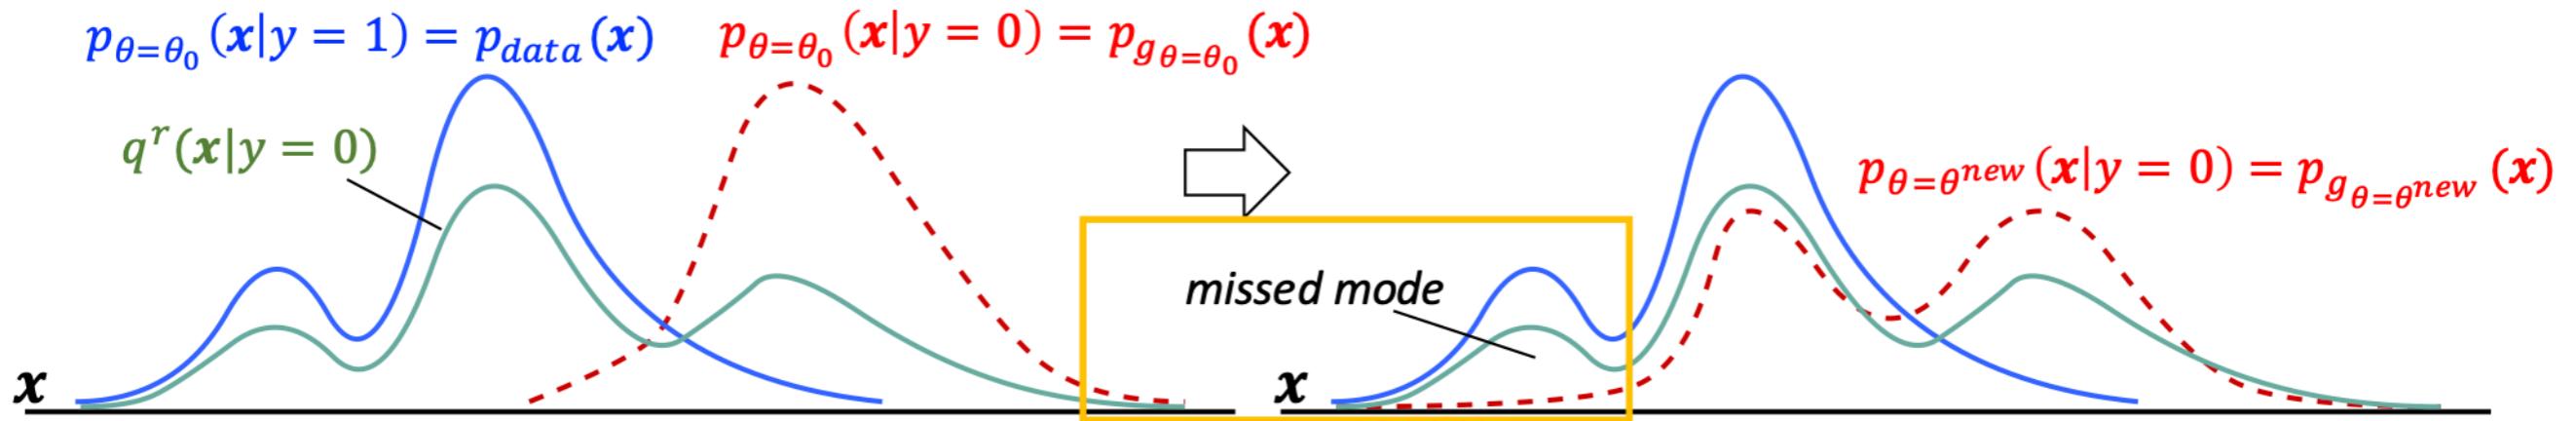
- Minimizing the KLD drives  $p_{g_\theta}(x)$  to  $p_{data}(x)$ 
  - By definition:  $p_{\theta=\theta_0}(x) = E_{p(y)}[p_{\theta=\theta_0}(x|y)] = \left(p_{g_{\theta=\theta_0}}(x) + p_{data}(x)\right)/2$
  - $KL(p_\theta(x|y=1)||q^r(x|y=1)) = KL(p_{data}(x)||q^r(x|y=1))$  : constant, no free parameters
  - $KL(p_\theta(x|y=0)||q^r(x|y=0)) = KL(p_{g_\theta}(x)||q^r(x|y=0))$  : parameter  $\theta$  to optimize
    - $q^r(x|y=0) \propto q_{\phi=\phi_0}^r(y=0|x)p_{\theta=\theta_0}(x)$
    - seen as a mixture of  $p_{g_{\theta=\theta_0}}(x)$  and  $p_{data}(x)$
    - mixing weights induced from  $q_{\phi=\phi_0}^r(y=0|x)$
  - Drives  $p_{g_\theta}(x|y)$  to mixture of  $p_{g_{\theta=\theta_0}}(x)$  and  $p_{data}(x)$   
 $\Rightarrow$  Drives  $p_{g_\theta}(x)$  to  $p_{data}(x)$

128





## GANs: minimizing KLD



- Missing mode phenomena of GANs
  - Asymmetry of KLD
    - Concentrates  $p_\theta(x|y=0)$  to large modes of  $q^r(x|y)$   
 $\Rightarrow p_{g_\theta}(x)$  misses modes of  $p_{data}(x)$
  - Symmetry of JSD
    - Does not affect the behavior of mode missing

$$\begin{aligned} \text{KL}\left(p_{g_\theta}(x)||q^r(x|y=0)\right) \\ = \int p_{g_\theta}(x) \log \frac{p_{g_\theta}(x)}{q^r(x|y=0)} dx \end{aligned}$$

- Large positive contribution to the KLD in the regions of  $x$  space where  $q^r(x|y=0)$  is small, unless  $p_{g_\theta}(x)$  is also small
- $\Rightarrow p_{g_\theta}(x)$  tends to avoid regions where  $q^r(x|y=0)$  is small





## Recap: conventional formulation of VAEs

- Objective:

$$\max_{\theta, \eta} \mathcal{L}_{\theta, \eta}^{\text{vae}} = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} [\mathbb{E}_{\tilde{q}_\eta(\mathbf{z}|\mathbf{x})} [\log \tilde{p}_\theta(\mathbf{x}|\mathbf{z})] - \text{KL}(\tilde{q}_\eta(\mathbf{z}|\mathbf{x})||\tilde{p}(\mathbf{z}))]$$

- $\tilde{p}(\mathbf{z})$ : prior over  $\mathbf{z}$
- $\tilde{p}_\theta(\mathbf{x}|\mathbf{z})$ : generative model
- $\tilde{q}_\eta(\mathbf{z}|\mathbf{x})$ : inference model
- Only uses real examples from  $p_{\text{data}}(\mathbf{x})$ , lacks adversarial mechanism
- To align with GANs, let's introduce the real/fake indicator  $y$  and adversarial discriminator

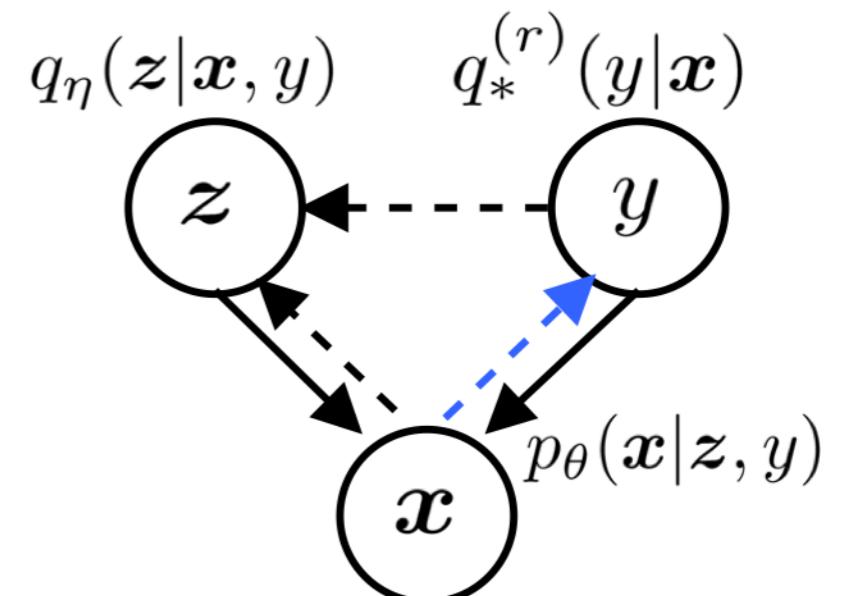




## VAEs: new formulation

- Assume a *perfect* discriminator  $q_*(y|x)$ 
  - $q_*(y=1|x) = 1$  if  $x$  is real examples
  - $q_*(y=0|x) = 1$  if  $x$  is generated samples
  - $q_*^r(y|x) := q_*(1-y|x)$
- Generative distribution

$$p_\theta(x|z, y) = \begin{cases} p_\theta(x|z) & y = 0 \\ p_{data}(x) & y = 1. \end{cases}$$



- Let  $p_\theta(z, y|x) \propto p_\theta(x|z, y)p(z|y)p(y)$
- *Lemma 2*

$$\begin{aligned}\mathcal{L}_{\theta, \eta}^{vae} &= 2 \cdot \mathbb{E}_{p_{\theta_0}(x)} [\mathbb{E}_{q_\eta(z|x,y)q_*^r(y|x)} [\log p_\theta(x|z, y)] - KL(q_\eta(z|x, y)q_*^r(y|x) \| p(z|y)p(y))] \\ &= 2 \cdot \mathbb{E}_{p_{\theta_0}(x)} [-KL(q_\eta(z|x, y)q_*^r(y|x) \| p_\theta(z, y|x))].\end{aligned}$$





## GANs vs VAEs side by side

$$p_{\theta}(\mathbf{z}, y | \mathbf{x}) \propto p_{\theta}(\mathbf{x} | \mathbf{z}, y) p(\mathbf{z} | y) p(y)$$

	GANs (InfoGAN)	VAEs
Generative distribution	$p_{\theta}(\mathbf{x} y) = \begin{cases} p_{g_{\theta}}(\mathbf{x}) & y = 0 \\ p_{data}(\mathbf{x}) & y = 1. \end{cases}$	$p_{\theta}(\mathbf{x} \mathbf{z}, y) = \begin{cases} p_{\theta}(\mathbf{x} \mathbf{z}) & y = 0 \\ p_{data}(\mathbf{x}) & y = 1. \end{cases}$
Discriminator distribution	$q_{\phi}(y \mathbf{x})$	$q_{*}(y \mathbf{x})$ , perfect, degenerated
z-inference model	$q_{\eta}(\mathbf{z} \mathbf{x}, y)$ of InfoGAN	$q_{\eta}(\mathbf{z} \mathbf{x}, y)$
KLD to minimize	$\min_{\theta} \text{KL}(p_{\theta}(\mathbf{x} y)    q^r(\mathbf{x} \mathbf{z}, y))$ $\sim \min_{\theta} \text{KL}(P_{\theta}    Q)$	$\min_{\theta} \text{KL}\left(q_{\eta}(\mathbf{z} \mathbf{x}, y) q_{*}^r(y \mathbf{x})    p_{\theta}(\mathbf{z}, y \mathbf{x})\right)$ $\sim \min_{\theta} \text{KL}(Q    P_{\theta})$

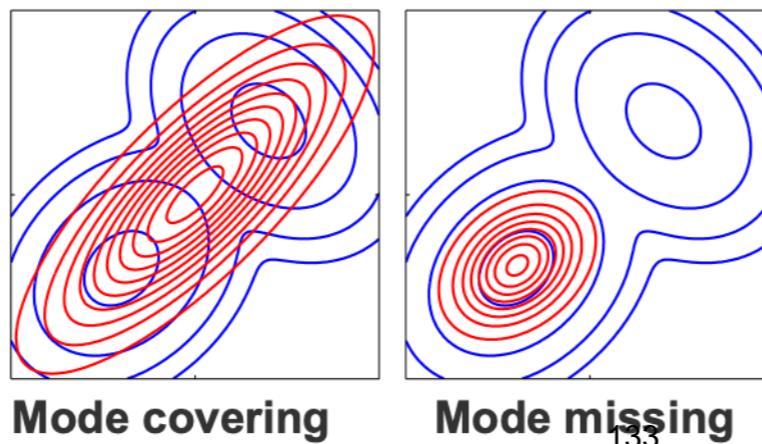




## GANs vs VAEs side by side

	GANs (InfoGAN)	VAEs
KLD to minimize	$\min_{\theta} \text{KL}(p_{\theta}(x y) \parallel q^r(x z, y))$ $\sim \min_{\theta} \text{KL}(P_{\theta} \parallel Q)$	$\min_{\theta} \text{KL}(q_{\eta}(z x, y)q_*^r(y x) \parallel p_{\theta}(z, y x))$ $\sim \min_{\theta} \text{KL}(Q \parallel P_{\theta})$

- Asymmetry of KLDs inspires combination of GANs and VAEs
  - GANs:  $\min_{\theta} \text{KL}(P_{\theta} \parallel Q)$  tends to missing mode
  - VAEs:  $\min_{\theta} \text{KL}(Q \parallel P_{\theta})$  tends to cover regions with small values of  $p_{data}$



[Figure courtesy: PRML]

© Eric Xing @ CMU, 2005-2020

50



# **Discriminative vs. Generative: Blurring the Distinction**

# Generative vs. Discriminative Classifiers

Training classifiers involves estimating  $f: X \rightarrow Y$ , or  $P(Y|X)$

Generative classifiers:

- Assume some functional form for  $P(X|Y)$ ,  $P(X)$
- Estimate parameters of  $P(X|Y)$ ,  $P(X)$  directly from training data
- Use Bayes rule to calculate  $P(Y|X=x_i)$

Discriminative classifiers:

1. Assume some functional form for  $P(Y|X)$
2. Estimate parameters of  $P(Y|X)$  directly from training data

- Consider learning  $f: X \rightarrow Y$ , where
  - $X$  is a vector of real-valued features,  $\langle X_1 \dots X_n \rangle$
  - $Y$  is boolean
  - assume all  $X_i$  are conditionally independent given  $Y$
  - model  $P(X_i | Y = y_k)$  as Gaussian  $N(\mu_{ik}, \sigma)$
  - model  $P(Y)$  as binomial ( $p$ )
- What does that imply about the form of  $P(Y|X)$ ?

$$P(Y = 1 | X = \langle x_1, \dots, x_n \rangle) = \frac{1}{1 + \exp(w_0 + \sum_i w_i x_i)}$$

# Logistic regression

- Logistic regression represents the probability of category  $i$  using a linear function of the input variables:

$$P(Y=i|X=x) = g(w_{i0} + w_{i1}x_1 + \dots + w_{id}x_d)$$

where for  $i < k$

$$g(z_i) = \frac{e^{z_i}}{1 + \sum_{j=1}^{K-1} e^{z_j}}$$

and for  $k$

$$g(z_k) = \frac{1}{1 + \sum_{j=1}^{K-1} e^{z_j}}$$

# Generative-Discriminative Pairs

Example: assume  $Y$  boolean,  $X = \langle X_1, X_2, \dots, X_n \rangle$ , where  $x_i$  are boolean, perhaps dependent on  $Y$ , conditionally independent given  $Y$

Generative model: naïve Bayes:

$$\hat{p}(x_i = 1 | y = b) = \frac{s\{x_i = 1, y = b\} + l}{s\{y = b\} + 2l}$$

$$\hat{p}(y = b) = \frac{s\{y = b\}}{\sum_j s\{y = j\}}$$

$s$  indicates size of set.

$l$  is smoothing parameter

Classify new example  $x$  based on ratio

$$\frac{\hat{p}(y = T|x)}{\hat{p}(y = F|x)} = \frac{\hat{p}(y = T) \prod_{i=1}^n \hat{p}(x_i|y = T)}{\hat{p}(y = F) \prod_{i=1}^n \hat{p}(x_i|y = F)}$$

Equivalently, based on sign of log of this ratio

# Generative-Discriminative Pairs

Example: assume  $Y$  boolean,  $X = \langle x_1, x_2, \dots, x_n \rangle$ , where  $x_i$  are boolean, perhaps dependent on  $Y$ , conditionally independent given  $Y$

Generative model: naïve Bayes:

$$\hat{p}(x_i = 1 | y = b) = \frac{s\{x_i = 1, y = b\} + l}{s\{y = b\} + 2l}$$

$$\hat{p}(y = b) = \frac{s\{y = b\}}{\sum_j s\{y = j\}}$$

Classify new example  $x$  based on ratio

$$\frac{\hat{p}(y = T|x)}{\hat{p}(y = F|x)} = \frac{\hat{p}(y = T) \prod_{i=1}^n \hat{p}(x_i|y = T)}{\hat{p}(y = F) \prod_{i=1}^n \hat{p}(x_i|y = F)}$$

Discriminative model: logistic regression

$$\hat{p}(y = T|x; \beta, \theta) = 1 / (1 + \exp(- \sum_{i=1}^n \beta_i x_i - \theta))$$

Note both learn linear decision surface over  $X$  in this case

# What is the difference asymptotically?

Notation: let  $\epsilon(h_{A,m})$  denote error of hypothesis learned via algorithm A, from  $m$  examples

- If assumed model correct (e.g., naïve Bayes model), and finite number of parameters, then

$$\epsilon(h_{Dis,\infty}) = \epsilon(h_{Gen,\infty})$$

- If assumed model incorrect

$$\epsilon(h_{Dis,\infty}) \leq \epsilon(h_{Gen,\infty})$$

Note assumed discriminative model can be correct even when generative model incorrect, but not vice versa

# Some experiments from UCI data sets

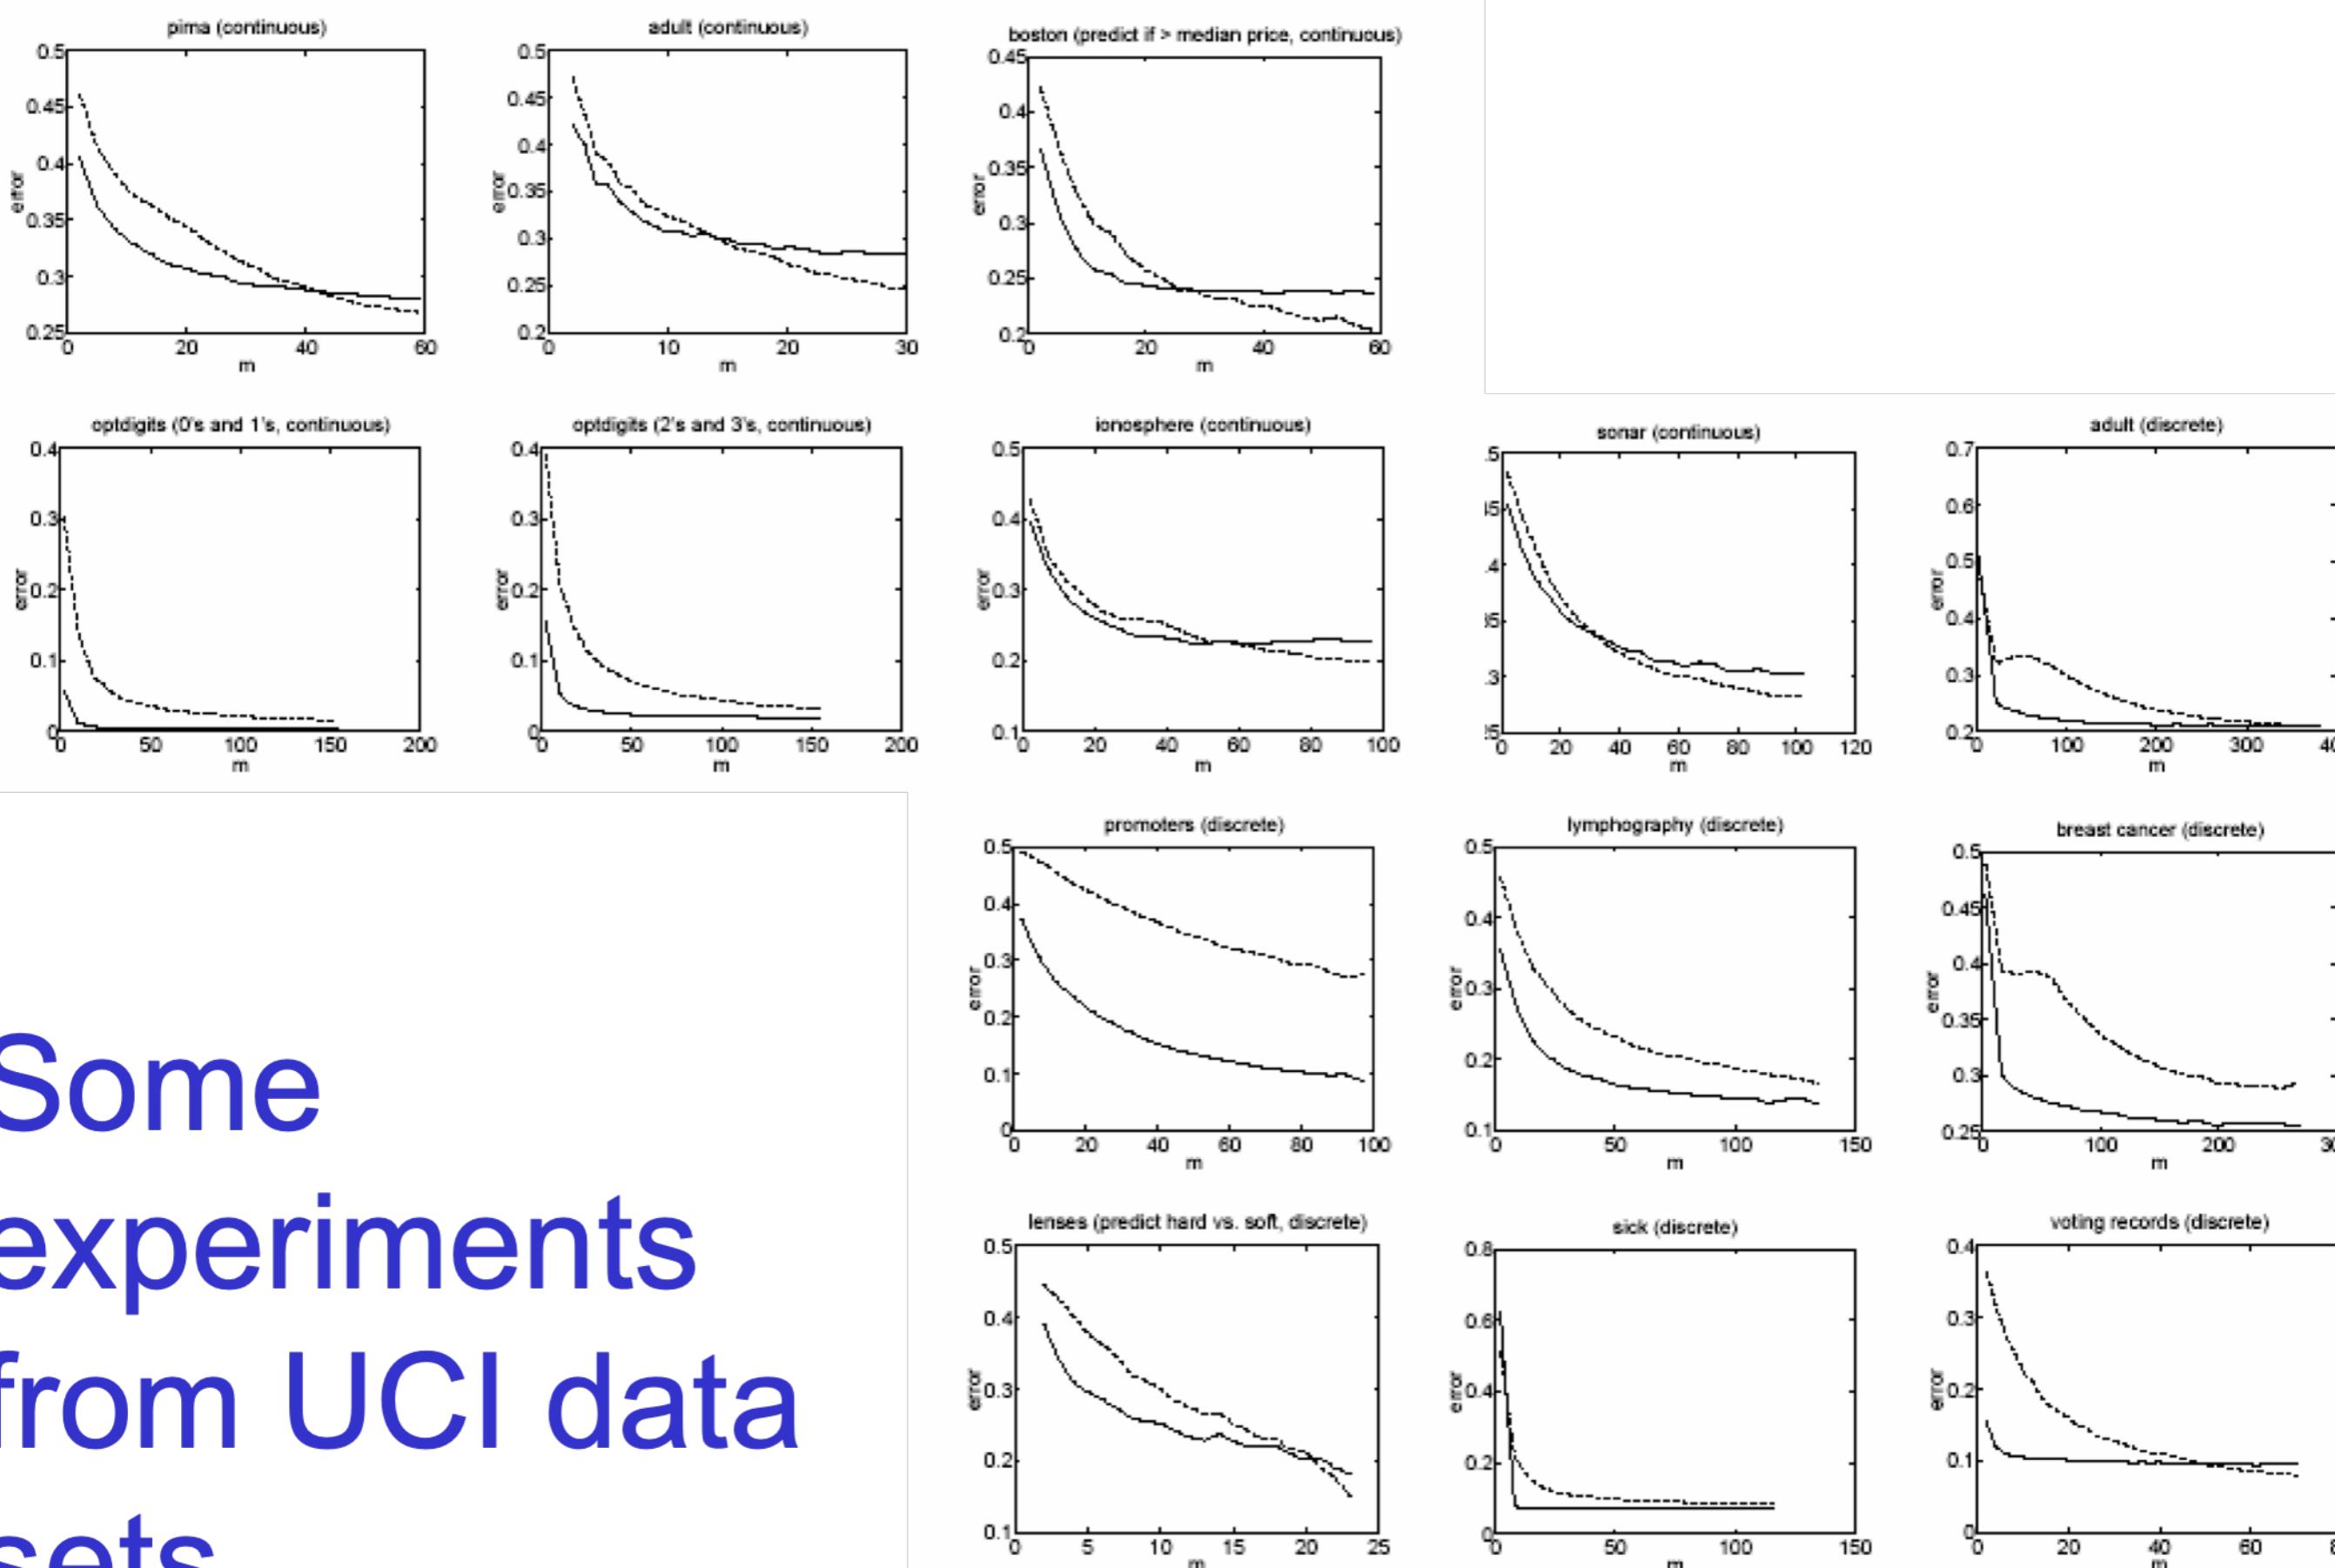


Figure 1: Results of 15 experiments on datasets from the UCI Machine Learning repository. Plots are of generalization error vs.  $m$  (averaged over 1000 random train/test splits). Dashed line is logistic regression; solid line is naive Bayes.